

TEL AVIV UNIVERSITY

The Iby and Aladar Fleischman Faculty of Engineering

Department of Electrical Engineering

**Deep Learning Architectures for
Two-Hop Relay Communication:
A Comparative Study of Classical and
Neural Network Relay Strategies**

A thesis submitted toward the degree of
Master of Science in Electrical and Electronic Engineering
In Tel Aviv University

by

Gil Zukerman

March 2026

TEL AVIV UNIVERSITY

The Iby and Aladar Fleischman Faculty of Engineering

Department of Electrical Engineering

**Deep Learning Architectures for
Two-Hop Relay Communication:
A Comparative Study of Classical and
Neural Network Relay Strategies**

A thesis submitted toward the degree of
Master of Science in Electrical and Electronic Engineering
In Tel Aviv University

by

Gil Zukerman

This research work was carried out at Tel-Aviv University
in the School of Electrical and Electronics Engineering
under the supervision of Dr. Anatoly Khina and Prof. Raja Giryes

March 2026

Contents

List of Figures	v
List of Tables	xi
1 Introduction	1
1.1 System Model	2
1.1.1 Two-Hop Relay Model	3
2 Literature Review	5
2.1 Classical Relay Strategies	5
2.1.1 Amplify-and-Forward (AF).	5
2.1.2 Decode-and-Forward (DF)	5
2.1.3 Other Classical Relay Techniques	6
2.2 Machine Learning in Wireless Communication	6
2.2.1 Theoretical Basis: Universal Approximation and Denoising	7
2.2.2 Prior Work in Deep Learning for Physical-Layer Processing	8
2.2.3 Relay methods comparison	9
2.2.4 Neural Network Relay Processing	10
2.3 Generative Models for Signal Processing	11
2.3.1 Variational Autoencoders	11
2.3.2 Conditional Generative Adversarial Networks	12
2.3.3 Generative vs. Discriminative Paradigms for Relay Processing	14
2.4 Sequence Models: Transformers, State Space Models	14
2.4.1 Transformers and the Attention Mechanism	14
2.4.2 Structured State Space Models	15
2.4.3 Mamba: Selective State Spaces	16
2.4.4 Mamba-2: Structured State Space Duality	18
2.4.5 State Space Models for Signal Processing	20
2.5 Theoretical Foundations: Channel Models	20
2.5.1 AWGN Channel : Theoretical Analysis	21
2.5.1.1 Single-hop BER	21
2.5.1.2 Two-hop AF BER	22
2.5.1.3 Two-hop DF BER	23
2.5.2 Rayleigh Fading Channel : Theoretical Analysis	23

2.5.2.1	Single-hop BER	23
2.5.2.2	Two-hop DF BER	24
2.5.3	Fading Coefficient Distributions	25
2.5.4	Channel Model Validation (Simulative)	26
2.6	Theoretical Foundations: MIMO Equalization	27
3	Research Objectives	28
3.1	Main Objective	28
3.2	Research Gap and Motivation	28
3.3	Research Hypotheses	29
3.4	Specific Objectives	31
3.5	Scope and limitations	32
4	Methods	33
4.1	System Model	33
4.1.1	Modulation.	33
4.1.2	Hop Model	34
4.1.3	Relay Processing	34
4.2	Relay Strategies	34
4.2.1	Selection and Classification of Relay Strategies	35
4.3	Simulation Framework	38
4.3.1	Monte Carlo BER Estimation	38
4.3.2	Statistical Significance Testing	39
4.3.3	Training Protocol	39
4.3.4	Reproducibility and Weight Management	40
4.4	Normalized Parameter Comparison	41
4.5	Extension Methods: QPSK and 16-QAM	42
4.5.1	QPSK : Gray-Coded Quadrature Phase-Shift Keying	42
4.5.2	16-QAM : Gray-Coded Quadrature Amplitude Modulation	43
4.5.3	Constellation Diagram Summary	44
4.5.4	I/Q Splitting for AI Relay Processing of Complex Constellations	45
4.5.5	2D Joint Classification as an Alternative to I/Q Splitting	46
5	Core Experiments: The Canonical Setup	48
5.1	Experiments summary	48
5.2	Channel Model Validation	49
5.2.1	Experiments	50
5.2.2	Results tables	50
5.2.3	Results Figures	51
5.3	Canonical Relay Comparison (SISO, Rayleigh, BPSK)	52
5.3.1	Results Tables	53

5.3.2	Results Figures	53
5.4	Parameter Normalization & Complexity Trade-off	54
5.4.1	Results Tables	54
5.4.2	Results Figures	55
6	Extension: Higher-Order Modulations	56
6.1	Higher-Order Modulation Scalability (Constellation-Aware Training)	56
6.1.1	Results Tables	57
6.1.2	Results Figures	58
6.2	16-Class 2D Classification for QAM16	62
6.2.1	Results Tables	63
6.2.2	Results Figures	63
7	Learned Relaying under Unknown and Mismatched Channels	65
7.1	Unknown Channel: Classical Failure and Learned Mitigation	66
7.1.1	Flat (Memoryless) Unknown Channels: A Necessary Control	66
7.1.2	A Composite (Mixed) Channel	69
7.1.3	The Posterior-Free (Blind) Regime	70
7.1.4	Partial Known Posterior: The Pilot-Budget Crossover	71
7.1.5	Complexity: The Cost Side of the Trade-off	73
8	Discussion and Conclusions	75
8.1	Interpretation of Results	75
8.1.1	Low-SNR Advantage of Neural Relays over AF (H1: Confirmed)	75
8.1.2	Classical Dominance at High SNR (H2: Confirmed)	76
8.1.3	Robustness Across the Fading Ensemble	76
8.2	The “Less is More” Principle (H3: Confirmed)	77
8.2.1	Information-Theoretic Perspective	77
8.2.2	Bias-Variance Analysis	77
8.2.3	Practical Implications	78
8.3	State Space vs. Attention for Signal Processing	78
8.3.1	Context-Length Benchmark: Validating the Crossover Hypothesis	79
8.3.2	Why state space suits signals.	81
8.4	Practical Deployment Recommendations	81
8.4.1	Latency and Adaptivity Considerations	82
8.5	Limitations	82
8.6	Future Work	84
8.7	Conclusions	85
9	Summary	91
10	Appendices	96

10.1	Appendix A: Mathematical Notation	96
10.2	Appendix B: Model Architectures and Hyperparameters	96
10.3	Appendix C: Software Architecture	97
10.3.1	Testing	97
10.3.2	Reproducibility	97
10.4	Appendix D: Normalized 3K-Parameter Configurations	98
10.5	Appendix E: Response to Supervisor Comments	98
	Abstract (Hebrew)	1

List of Figures

2.1	Mamba: Linear-Time Sequence Modeling with Selective State Spaces. Structured SSMs independently map each channel (e.g. $D = 5$) of an input x to output y through a higher dimensional latent state h (e.g. $D = 4$). Prior SSMs avoid materializing this large effective state (DN times batch size B and sequence length L) through clever alternate computation paths requiring time-invariance: the (Δ, A, B, C) parameters are constant across time. Our selection mechanism adds back input-dependent dynamics, which also requires a careful hardware-aware algorithm to only materialize the expanded states in more efficient levels of the GPU memory hierarchy	18
2.2	The Mamba-2 block simplifies the Mamba block by removing sequential linear projections; the SSM parameters A, B, C are produced at the beginning of the block instead of as a function of the SSM input X . An additional normalization layer is added as in <i>NormFormer</i> [1], improving stability. The B and C projections only have a single head shared across the X heads, analogous to multi-value attention (MVA).	20
2.3	AWGN channel : theoretical BER (solid lines) vs. Monte Carlo simulation (markers with 95% CI) for single-hop, two-hop AF, and two-hop DF. . . .	22
2.4	Rayleigh fading : theory vs. simulation for single-hop and two-hop DF. The characteristic $1/\text{SNR}$ slope (compared to the steeper exponential AWGN curve) is clearly visible.	25
2.5	Left : PDF of $ h $ for Rayleigh and Rician ($K = 1, 3, 10$). Right : CDF (outage probability). Rayleigh has the highest deep-fade probability at any threshold.	26
4.1	Two-hop relay system model: source transmits through two channels with a neural network relay in between.	33
4.2	Constellation diagrams. (a) BPSK: 2 real-valued symbols at ± 1 . (b) QPSK: 4 Gray-coded symbols on the unit circle. (c) 16-QAM: 16 Gray-coded symbols on a 4×4 rectangular grid with decision boundaries (dotted lines). (d) 16-PSK, shown for completeness only; it is not studied in this work. All constellations are normalised to unit average power.	45

5.1	AWGN Channel : Theoretical vs. Simulative BER. Single-hop, two-hop AF, and two-hop DF: closed-form theory (solid lines) versus Monte Carlo simulation (markers with 95% CI). Theory and simulation match within the confidence interval at all SNR points.	51
5.2	Rayleigh fading : theory vs. simulation for single-hop and two-hop DF. The characteristic $1/\text{SNR}$ slope (compared to the steeper exponential AWGN curve) is clearly visible.	52
5.3	Left : PDF of $ h $ for Rayleigh and Rician ($K = 1, 3, 10$). Right : CDF (outage probability). Rayleigh has the highest deep-fade probability at any threshold.	52
5.4	Rayleigh fading : BER vs. SNR for all nine relay strategies with 95% CI.	53
5.5	Complexity–performance trade-off. Training time vs. parameter count vs. BER improvement over DF at low SNR. The Minimal MLP (169 params) achieves the best efficiency.	55
6.1	BPSK on AWGN : all relay strategies with 95% CI (baseline for modulation comparison).	58
6.2	QPSK on AWGN : BER curves closely match the BPSK baseline (6.1), confirming I/Q splitting validity.	59
6.3	16-QAM on AWGN : AI relays hit a BER floor near 0.22 at medium-high SNR due to tanh compression of multi-level signals; AF outperforms DF at low SNR (Y	59
6.4	Combined modulation comparison (AWGN) : all nine relays across BPSK (solid), QPSK (dashed, overlapping BPSK), and 16-QAM (dotted). The BPSK/QPSK overlap confirms I/Q splitting equivalence. The 16-QAM dotted curves reveal the AI relay BER floor: all AI relays plateau near 0.18–0.25 while DF and AF continue decreasing.	60
6.5	16-QAM activation experiment (AWGN) : dashed lines = tanh/BPSK baseline, solid = linear/QAM16, dotted = hardtanh/QAM16. Replacing tanh and retraining on QAM16 eliminates the BER floor for all AI relays except Hybrid. Sequence models (Transformer, Mamba S6, Mamba-2) benefit most, narrowing the gap to DF from $\sim 56\times$ to $\sim 10\times$	61
6.6	Comparison of activation function shapes (left) and their derivatives (right) for $A_{\max} = 0.9487$ (16-QAM). Hardtanh has a sharp transition at the clip bounds; sigmoid and scaled tanh provide smooth saturation with non-zero gradients throughout.	62

6.7	16-QAM BER vs. SNR for all relay variants (4-class and 16-class) on AWGN. The 4-class variants (dashed) plateau at BER $\hat{\approx} 0.008$ while the successfully trained 16-class variants (solid) continue decreasing, approaching and matching the DF baseline.	63
6.8	Grouped bar comparison of 4-class vs. 16-class BER at 20 dB for each architecture. The 16-class advantage is visible for MLP, VAE, Transformer, and Mamba S6. Failed variants (VAE 4-cls ~ 0.50 , CGAN 16-cls ~ 0.50) are truncated.	64
6.9	Top-3 16-class relay variants compared against AF and DF. VAE 16-cls, Transformer 16-cls, and MLP 16-cls all match or approach DF performance across the full SNR range.	64
7.1	Unknown-channel relaying, BER vs. SNR with 95% CI bands. (a)–(b) On the unknown ISI channel, the memoryless AF/DF relays are pinned near the analytic 0.25 floor (DF non-monotonic in SNR); the 170-parameter MLP restores monotonically decreasing BER within ≈ 1 –1.5 dB of genie-CSI Viterbi MLSE, which a 200-pilot LS estimate matches almost exactly. With the Rayleigh second hop (b), MLSE and the MLP are statistically indistinguishable end-to-end. (c) Control: on the canonical Rayleigh channel the MLP only matches DF, consistent with H2.	69
7.2	Composite channel (ISI $\times$ PA-nonlinearity $\times$ unknown phase, DBPSK). Naive classical relays (AF, differential detection) saturate at the memoryless-relay floor; the 170-parameter MLP restores reliable relaying from raw I/Q with no impairment model, tracking about 2 dB behind the strongest constructed classical baseline (pilot-LS Viterbi + differential). The 1,153-parameter MLP adds nothing (H3).	70
7.3	Posterior-free composite channel (no pilots, no channel prior, fresh realization per block). CMA (blind classical) and the family-trained MLP converge together; decision-directed blind MLSE is unstable (non-monotonic, wide CI) because it bootstraps the channel from its own decisions; plain differential detection floors. Without a posterior, learning matches — not beats — the best blind classical equalizer, but does so with one-shot, stable inference and no online adaptation.	71

7.4	Partial known posterior. (a) Pilot-budget sweep at 10 dB: pilot-aided Viterbi beats the pilot-free MLP down to ~ 10 pilots, then collapses at 5 pilots when the channel becomes unidentifiable; the MLP is flat (amortized posterior, no pilots). (b) Block-length sweep with 10 pilots/block: payload BERs stay close, but the classical pilot overhead grows from 1% at $L=1000$ to 25% at $L=40$, where the learned relay's zero-overhead operation is decisive.	73
7.5	Complexity of Viterbi MLSE vs the 170-parameter MLP. (a) Per-symbol operations: Viterbi grows as M^L (intractable by $M=16, L=5$), while the MLP is constant at ~ 330 flops. On the small studied channel (BPSK, $L=3$) Viterbi is actually cheaper; the MLP wins on <i>scaling</i> , not on the small case. (b) Measured inference time: the MLP is $30\text{--}90\times$ faster as a vectorized matmul versus Viterbi's sequential trellis scan.	74

List of Equations

1.1	awgn-hop1	3
1.2	awgn-hop2	3
1.3	awgn-sigma2	4
2.0	af-gain	5
2.2	af-snr	5
2.3	df-demod-remod	6
2.4	df-ber-hops	6
2.5	optimal-denoiser	7
2.6	optimal-denoiser-awgn	7
2.7	bias-variance-decomp	7
2.8	mse-loss	10
2.9	mlp-relay	10
2.11	vae-elbo	11
2.12	vae-loss	11
2.14	repremtized-trick	12
2.15	gan-minimax	12
2.16	wasserstein-distance	13
2.17	gradient-penalty	13
2.18	x-hat	13
2.19	wgan-gen-loss	13
2.20	wgan-disc-loss	13
2.21	attention	14
2.22	multihead-attention	14
2.23	multihead-attention-parameters-count	15
2.24	multihead-attention-matrix	15
2.25	positional-encoding	15
2.26	ssm-continuous	15
2.27	ssm-zoh	16
2.28	ssm-discrete	16
2.29	mamba-selective	16
2.30	mamba-selective-ak	17
2.31	mamba-selective-bk	17

2.31	mamba-selective-ak	17
2.32	mamba-selective-ak2	17
2.33	mamba-ssd-output	18
2.34	mamba-ssd-matrix	18
2.35	awgn-channel	21
2.36	awgn-ber	21
2.37	af-snr-eff	22
2.38	af-p-eff	22
2.39	gamma-high-snr	22
R		23
2.41	equal-hop-snr	23
2.42	df-high-snr-two-hop-penalty	23
2.43	rayleigh-channel	23
2.44	rayleigh-pdf	23
2.45	rayleigh-ber	24
2.46	rayleigh-ber-approx	24
2.47	df-ber-rayleigh	24
2.48	zf-equalizer-methods	27
2.49	mmse-equalizer-methods	27
4.1	two-hop-system	33
4.2	bpsk-mapping	33
4.3	hop-model	34
4.4	power-normalization	34
4.5	relay-received-iso	34
4.6	ber-estimate	38
4.7	confidence-interval	38
4.8	wilcoxon-test	39
4.9	qpsk-mapping	42
4.10	qpsk-ber	43
4.11	qam16-mapping	43
4.12	qam16-ber	44
4.13	iq-splitting	45
4.14	joint-classification	46
8.1	optimal-denoiser-relay	75
8.2	bias-variance-relay	77

List of Tables

1.1	Scope of the thesis: the canonical setup, the modulation extension (QPSK/16-QAM), and the unknown-channel contribution. No other topology, relay protocol, or metric is evaluated anywhere in this work.	2
2.1	Comparison of relay processing strategies	9
2.2	Theoretical SNR required to achieve $\text{BER} = 10^{-3}$ on the canonical Rayleigh channel and the AWGN calibration channel (single-hop BPSK).	26
3.1	Research positioning: comparison of this thesis against prior work in deep-learning-based relay communication.	29
4.1	Normalized 3K-parameter model configurations: parameter count and scaling method for each architecture.	41
4.2	QPSK Gray-coded symbol mapping: bit pairs to complex symbols.	42
4.3	16-QAM PAM-4 Gray-coded level mapping for in-phase and quadrature components.	43
4.4	Relay-specific complex signal handling strategies and rationale.	46
5.1	Summary of experimental configurations (results reported in appendices)	48
5.2	BER comparison between theoretical analysis and simulation results at selected SNR values	50
5.3	BER comparison on the Rayleigh fading channel (SISO).	53
5.4	Normalized 3K BER results : Rayleigh fading channel.	54
5.5	Model complexity and timing comparison on the canonical setup (experiment-specific training settings; Monte Carlo evaluation over 11 SNR points x 10 trials x 10,000 bits). All inference uses batched window extraction and a single forward pass per signal block.	54
6.1	BER comparison across modulations at selected SNR points (AWGN channel). All nine relay strategies.	57
6.2	BER at 16 dB for all relay variants across activation functions and modulations (AWGN and Rayleigh).	57
6.3	16-QAM BER improvements at 20 dB for sequence models with and without CSI injection and layer normalisation (AWGN).	58

6.4	BER at 20 dB for all relay variants: 4-class (I/Q split) vs. 16-class (joint 2D) classification for 16-QAM.	63
7.1	Flat unknown channels: classical relay vs. MLP-170 BER at 8/12/16/20 dB (mean, 5 trials). The maximum DF–MLP gap over all cases and SNRs is 0.0036 — within statistical noise. Unknown parameters alone do not defeat classical relays; memory does.	67
7.2	Unknown-channel BER (mean over 5 trials). The ISI setups exhibit memoryless-relay error floors (≈ 0.25); the 170-parameter MLP restores reliable relaying within ≈ 1 – 1.5 dB of genie-CSI Viterbi MLSE, and matches it end-to-end when the Rayleigh second hop dominates. The control shows no MLP advantage on the canonical channel, consistent with H2.	68
8.1	Context-length benchmark : three sequence models at $n = 255$ on CUDA.	80
8.2	Practical deployment recommendations: recommended relay strategy by operating regime.	81
8.3	Research hypotheses, their statements, and experimental outcomes.	85
10.3	Comparison of relay models, training objectives, parameter counts, and implementations.	97

Chapter 1

Introduction

Relay communication is a key technique in contemporary wireless systems that uses intermediate nodes between a source and a destination to extend coverage, enhance reliability, and boost throughput. Classical relay strategies such as amplify-and-forward (AF) and decode-and-forward (DF) have been extensively analyzed and are widely used as baseline protocols; in several canonical relay-channel models, DF (and related forwarding schemes) can achieve capacity or approach fundamental limits, particularly in the high-SNR regime [2, 3, 4, 5, 6].

Recent advances in machine learning have motivated the application of neural networks to physical-layer communication problems, including detection, decoding, and end-to-end system optimization. Neural models can approximate nonlinear mappings that are analytically intractable and may offer advantages in low-SNR or non-ideal regimes.

This motivates the central question of this thesis: *under what conditions can neural network-based relays improve upon or complement classical relay strategies in two-hop cooperative communication?*

This thesis presents a systematic comparative study of classical and neural relay strategies within a unified simulation framework. Nine relay methods are evaluated: Two classical relay strategies (AF and DF) and seven neural relays spanning supervised learning (MLP, SNR-adaptive hybrid relay), generative modeling (variational autoencoder VAE and conditional GAN with WGAN-GP), and sequence-based architectures (Transformer, Mamba-S6, and Mamba-SSD with structured state space duality).

The main objective of this thesis is not to propose a new relay design, but to quantify *performance Vs. complexity* tradeoffs and identify regimes where learning-based relays are effective, as Existing studies typically focus on a limited set of architectures or channel models, making it difficult to disentangle the effects of model structure, parameter count, and channel conditions.

Scope of this thesis. This thesis studies a single, fixed communication problem — the classical source-relay-destination (two-hop) link — under *one canonical setup*: SISO on both hops, i.i.d. Rayleigh fast fading, complex baseband, BPSK, uncoded BER. The

only axis varied is the *relay function*. A scoped modulation extension (Chapter 6) and a dedicated unknown-channel contribution (Chapter 7) are included; all other departures from the canonical setup are identified as future work in Chapter 8:

Table 1.1: Scope of the thesis: the canonical setup, the modulation extension (QPSK/16-QAM), and the unknown-channel contribution. No other topology, relay protocol, or metric is evaluated anywhere in this work.

Axis	Canonical value (future work in parentheses)
Topology	SISO on both hops (<i>future work: MIMO second hop</i>)
Channel	i.i.d. Rayleigh fast fading; AWGN retained as analytical calibration; unknown, mismatched, and time-varying channels are studied in the dedicated contribution (Chapter 7) (<i>future work: Rician and other channel families</i>)
Relay Function (<i>the studied axis</i>)	<i>Classical:</i> AF, DF. <i>Learned:</i> MLP, Hybrid, VAE, CGAN, Transformer, Mamba-S6, Mamba-2-SSD
Modulation	BPSK; QPSK and 16-QAM studied as a scoped extension (Chapter 6) (<i>future work: 16-PSK, 64-QAM and beyond</i>)
Performance Metric	Uncoded BER vs. SNR, Monte-Carlo averaged

Everything else discussed in the literature review (the general relay channel with a direct link, its capacity bounds, and the other classical relay protocols such as compress-, compute-, estimate-, and filter-and-forward) is provided as *background context only* and is not part of the studied system.

1.1 System Model

System model studied in this thesis. Throughout this work, a *single canonical setup* is fixed and never departed from in the main line of argument: a half-duplex two-hop link with a single relay and *no* direct source-to-destination path, in which the source S transmits during a listening phase to the relay R , and the relay then transmits during a transmission phase to the destination D . Both hops are single-input single-output (SISO) complex-baseband channels with i.i.d. Rayleigh *fast* fading — a fresh realization $h \sim \mathcal{CN}(0, 1)$ per channel use, known perfectly at the corresponding receiver. The modulation is BPSK, and the performance metric is the *uncoded* bit-error rate (BER) as a function of SNR, averaged by Monte Carlo over the fading ensemble. Since BPSK carries information on the real axis only, each receiver applies coherent phase compensation ($y \cdot h^*/|h|$) and takes $\Re\{\cdot\}$ of the result; the imaginary component then contains noise only, so this projection is the sufficient statistic and loses nothing. The learned component of the system is confined

to the *relay processing function* $f_\theta(\cdot)$ — the block for which AF and DF are the classical alternatives; the source, channels, and destination demodulator are fixed, so observed BER differences are attributable to the relay strategy alone. The AWGN channel appears only as the analytical calibration limit used to validate the simulator. Beyond the canonical core, the thesis makes two further contributions of unequal weight. A scoped modulation extension (Chapter 6) tests whether the canonical BPSK conclusions generalize to QPSK and 16-QAM, including a joint 2D N -class classification formulation of the relay. The principal experimental contribution (Chapter 7) is a systematic study of learned relaying under *unknown and mismatched* channels — mapping, along every axis a skeptic would probe, precisely when a minimal learned relay adds value over the strongest classical baselines and when it does not (H6). All other departures (other channel families, a MIMO second hop, learned constellations) are identified as future work in Chapter 8. Capacity-theoretic results and the cooperative-diversity order are quoted in Chapter 2 as classical *background only*; they are not the operating point evaluated in the experiments.

1.1.1 Two-Hop Relay Model

In the half-duplex two-hop model studied in this thesis, the relay cannot transmit and receive simultaneously, and there is no direct source–destination link. The communication proceeds in two time slots:

Slot 1 (Source → Relay) — listening phase:

$$y_R[i] = x[i] + n_1[i], \quad n_1[i] \sim \mathcal{N}(0, \sigma^2) \quad (1.1)$$

Slot 2 (Relay → Destination) — transmission phase:

$$y_D[i] = x_R[i] + n_2[i], \quad n_2[i] \sim \mathcal{N}(0, \sigma^2) \quad (1.2)$$

where $x[i] \in \{-1, +1\}$ is the BPSK symbol transmitted by the source at time i , $y_R[i]$ is the signal received at the relay, $x_R[i] = f_\theta(y_R[i - w : i + w])$ is the signal transmitted by the relay at time i , and $y_D[i]$ is the signal received at the destination.

Remark 1.1 (Realizability of the symmetric window). The window $y_R[i - w : i + w]$ uses look-ahead samples and is therefore non-causal in the *streaming* sense. It is nonetheless fully realizable in the half-duplex protocol studied here: the relay buffers the *entire* listening-phase block before the transmission phase begins, so by the time $x_R[i]$ is computed, all samples $y_R[i - w : i + w]$ have already been received. The symmetric window is thus a block-level (store-and-forward) operation with a fixed processing latency of w

symbols. The classical baselines AF and DF are memoryless in their per-symbol operation ($w = 0$) and are strictly causal.

The noise terms are independent and identically distributed with variance

$$\sigma^2 = \frac{P_s}{\gamma} \quad (1.3)$$

The choice of relay processing function $f(\cdot)$ fundamentally determines system performance and is the central object of study in this thesis. Classical approaches include amplify-and-forward (AF), which simply scales the received signal, and decode-and-forward (DF), which regenerates the signal through demodulation and re-modulation. Each has well-understood performance characteristics: AF is simple but propagates noise, while DF eliminates first-hop noise but introduces error propagation when decoding fails [2].

Remark 1.2 (Terminology: symbol-wise DF vs. information-theoretic DF). The term “decode-and-forward” is used in two distinct senses in the literature, and the distinction matters for interpreting the results of this thesis. In its *information-theoretic* sense [2, 3], DF is a *block* strategy: the source protects each message block with a strong (ideally capacity-achieving) channel code, the relay decodes the entire block, and only then re-encodes and forwards the recovered message. This block-DF is inherently *not* memoryless and, over a single link, can operate reliably at any rate below the first-hop capacity.

The system studied in this thesis is *uncoded*: one BPSK symbol carries one information bit and no outer error-correction code is applied. Consequently, the “DF” baseline throughout this work refers to *symbol-wise slicing*: per-symbol demodulation followed by re-modulation, i.e., $f_{\text{DF}}(y_R[i]) = \text{sign}(y_R[i])$ with $w = 0$. This is the appropriate classical counterpart to the learned relays under the uncoded BER metric used here, but it is strictly weaker than block-DF with coding; the reported DF results should not be read as bounds on coded block-DF performance. Extending the comparison to the coded setting — where the relay (classical or learned) aggregates information across a full code block — is an important direction for future work and is discussed in Chapter 8.

Chapter 2

Literature Review

2.1 Classical Relay Strategies

The fundamental question that motivates this thesis is: can a learned relay function $f_\theta(\cdot)$ outperform the classical relay functions (AF scaling, DF regeneration) by exploiting patterns in the received signal that analytical methods do not capture? This question is particularly relevant at low SNR, where both AF (noise amplification) and DF (decoding errors) have well-known limitations, and where the non-linear decision boundary learned by a neural network may provide a substantive advantage.

2.1.1 Amplify-and-Forward (AF).

The AF relay amplifies the received signal by a gain factor G that normalizes the output power: (Equation 2.1)

$$G = \sqrt{\frac{P_{\text{target}}}{\mathbb{E}[|y_R|^2]}} \quad (2.1)$$

The end-to-end SNR for AF relay is: (Equation 2.2)

$$\text{SNR}_{\text{eff}}^{\text{AF}} = \frac{\text{SNR}_1 \cdot \text{SNR}_2}{\text{SNR}_1 + \text{SNR}_2 + 1} \quad (2.2)$$

This expression reveals a fundamental limitation: even when one hop has high SNR, the effective SNR is bottlenecked by the weaker hop. Moreover, the relay amplifies both signal and noise from the first hop, leading to noise accumulation at the destination.

2.1.2 Decode-and-Forward (DF)

The DF relay demodulates the received signal, recovers the transmitted bits, and re-modulates clean symbols:

$$\hat{b}_R = \text{demod}(y_R), \quad x_R = \text{mod}(\hat{b}_R) \quad (2.3)$$

The end-to-end BER for DF is: (Equation 2.4)

$$P_e^{\text{DF}} = P_{e,1} + (1 - P_{e,1}) \cdot P_{e,2} \quad (2.4)$$

where $P_{e,1}$ and $P_{e,2}$ are the BER of the first and second hops, respectively. For BPSK modulation over AWGN (real-valued noise with variance $1/\text{SNR}$), each hop's BER is $P_e = Q(\sqrt{\text{SNR}})$. DF provides clean regeneration at high SNR but suffers from error propagation when the first-hop BER is non-negligible.

2.1.3 Other Classical Relay Techniques

AF and DF represent two canonical extremes of the relay processing spectrum: AF performs minimal processing (linear scaling) while DF performs maximal classical processing (full regeneration). Additional techniques (e.g. CF, EF, CoF) occupy intermediate points along this spectrum. In this thesis, the AI-based relay strategies are designed to *learn* the optimal relay processing function from data, potentially discovering strategies that subsume or outperform these fixed classical schemes. The focus on AF and DF as classical baselines is therefore deliberate: they bracket the classical design space and provide clear lower and upper bounds against which the learned strategies can be evaluated. These additional relay strategies are reviewed for completeness; they are not simulated in this thesis because the studied system assumes a two-hop link without a direct source-destination path.

2.2 Machine Learning in Wireless Communication

The application of machine learning to physical-layer wireless communication has gained significant momentum in recent years, driven by the ability of neural networks to learn complex non-linear mappings directly from data. Deep learning has been applied to channel estimation [7], signal detection [8], autoencoder-based end-to-end communication [9], and resource allocation [10]. These approaches learn complex mappings from data, potentially outperforming hand-crafted algorithms in scenarios where analytical solutions are intractable or suboptimal.

2.2.1 Theoretical Basis: Universal Approximation and Denoising

The theoretical justification for applying neural networks to relay signal processing rests on two pillars. First, the **universal approximation theorem** [11] guarantees that a feedforward network with a single hidden layer and a non-linear activation function can approximate any continuous function on a compact domain to arbitrary accuracy, given sufficient width. For relay denoising, the target function maps noisy observations to clean transmitted symbols:

$$f^* : \mathbb{R}^{2w+1} \rightarrow [-1, 1], \quad f^*(y_{i-w}, \dots, y_{i+w}) = \mathbb{E}[x_i \mid y_{i-w}, \dots, y_{i+w}] \quad (2.5)$$

This conditional expectation f^* is the Bayes-optimal denoiser : it minimizes the mean squared error (MSE) over all possible estimators. For BPSK symbols corrupted by AWGN, f^* reduces to the posterior mean:

$$f^*(y) = \tanh\left(\frac{y}{\sigma^2}\right) \quad (2.6)$$

which is a smooth sigmoid-like function that approaches the hard-decision signum function as $\sigma^2 \rightarrow 0$ (high SNR). A neural network with a single hidden layer and tanh output can represent this function exactly, explaining why even a 169-parameter network suffices for this task.

Second, the **bias-variance decomposition** provides a framework for understanding model complexity:

$$\mathbb{E}[(\hat{x} - x)^2] = \text{Bias}^2(\hat{x}) + \text{Var}(\hat{x}) + \sigma_{\text{irreducible}}^2 \quad (2.7)$$

The irreducible noise $\sigma_{\text{irreducible}}^2$ represents the minimum achievable MSE, determined by the channel noise. Increasing model complexity (more parameters) reduces bias but increases variance. For the relay denoising task, the target function f^* is simple (essentially a soft threshold), so the bias term is already small for modest networks. Adding parameters beyond this point primarily increases variance (overfitting), explaining the inverted-U relationship between model size and BER observed in this thesis.

2.2.2 Prior Work in Deep Learning for Physical-Layer Processing

Ye et al. [7] demonstrated that deep neural networks can jointly perform channel estimation and signal detection in OFDM systems, achieving near-optimal performance with lower complexity than separate estimation and detection stages. Their key insight was that the DNN implicitly learns the channel's statistical structure, eliminating the need for explicit pilot-based estimation.

Samuel et al. [8] proposed DetNet, an unfolded projected gradient descent network for MIMO detection. By unrolling iterative optimization into a fixed number of neural network layers, DetNet achieves near-maximum-likelihood detection performance with $O(N_t^2)$ complexity instead of the exponential complexity of exhaustive ML search. This demonstrates the power of incorporating domain knowledge (the iterative structure of the detection problem) into the network architecture.

Dorner et al. [9] proposed treating the entire communication system : modulator, channel, and demodulator : as an autoencoder, trained end-to-end to minimize bit error rate. This approach jointly optimizes all components, potentially discovering novel modulation and coding schemes that outperform separately designed modules. The autoencoder paradigm is conceptually related to the relay processing task: (Equation 2.29) both involve learning an encoder-decoder pair that maps through a noisy channel.

Sun et al. [10] applied deep learning to the NP-hard problem of resource allocation in interference networks, demonstrating that a DNN can learn near-optimal power control policies orders of magnitude faster than conventional iterative algorithms.

In the context of relay-assisted communication systems, most existing works employ neural networks primarily for *relay-selection* and *resource-allocation*, particularly in decode-and-forward (DF) relaying scenarios, where learning-based classifiers or predictors are used to identify the optimal relay based on channel state information, signal-to-noise ratios, or energy constraints, while the underlying relay operation remains conventional [12].

A broader perspective on the role of machine learning in wireless communications is provided in [13], which highlights the potential of learning-based methods across multiple protocol layers, including cooperative communications.

Another related line of work exploits the structural similarities between amplify-and-forward (AF) relay networks and neural networks, treating relay nodes as neurons with nonlinear transfer functions and using deep-learning optimization tools to tune relay gains and exploit hardware-induced nonlinearities [14, 15]. While these approaches demonstrate significant gains over classical linear AF schemes, they focus on analog gain optimization rather than symbol-level detection.

Classical cooperative communication frameworks based on AF and DF relaying remain

well understood and widely adopted [16], but they rely on fixed linear forwarding or hard-decision rules. Consequently, despite substantial progress, relatively little attention has been devoted to learning the relay’s symbol-level input-output mapping itself, motivating neural network-based detection at the relay as a generalized, data-driven alternative to conventional AF and DF processing.

2.2.3 Relay methods comparison

2.1 compares classical relay processing strategies with learning-based approaches, highlighting their fundamental differences in terms of signal processing, nonlinearity, and learning capability. *Amplify-and-Forward (AF)* relaying applies linear scaling to the received signal at the relay, which keeps implementation simple but inevitably amplifies noise along with the desired signal. *Decode-and-Forward (DF)* relaying performs hard/soft symbol decisions at the relay before re-encoding and forwarding, thereby suppressing noise but introducing potential error propagation due to incorrect decisions. Prior deep-learning-based relay works largely preserve these conventional AF or DF operations and employ neural networks mainly for auxiliary tasks such as relay selection, power allocation, or analog gain optimization. As a result, learning is only partially involved in the overall relay operation and does not directly determine the symbol-level forwarding rule. In contrast, the proposed neural network-based detection (NND) relay replaces the fixed relay processing rule with a learned nonlinear mapping $f(y_R)$. This enables the relay to adaptively balance noise suppression and decision reliability in a data-driven manner, effectively generalizing both AF and DF strategies. The table therefore highlights that Neural network is the only approach that jointly introduces nonlinearity and fully learned symbol-level processing at the relay, distinguishing it fundamentally from both classical relaying schemes and prior learning-assisted approaches.

Table 2.1: Comparison of relay processing strategies

Method	Relay Mapping	Nonlinear	Noise Handling	Learned
AF	Linear scaling	No	Noise amplified	No
DF	Hard/soft decision	Limited	Noise suppressed	No
Prior DL (Selection / Optimization)	Conventional AF/DF	No	Conventional	Partially
NND (This work)	Learned $f(y_R)$	Yes	Learned trade-off	Yes

2.2.4 Neural Network Relay Processing

For relay processing specifically, a supervised learning approach trains a neural network f_θ to minimize: (Equation 2.5)

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N (\hat{x}_i - x_i)^2, \quad \hat{x}_i = f_\theta(y_{i-w:i+w}) \quad (2.8)$$

where $\hat{x}_i = f_\theta(y_{i-w:i+w})$ is the network output based on a sliding window of $2w + 1$ received symbols, and x_i is the clean transmitted symbol. This window-based approach provides temporal context that enables the network to exploit statistical dependencies in the noise-corrupted signal.

An well known example for $f_\theta(\cdot)$ is the *Multi-layer-perceptron (MLP)*, a feedforward artificial neural network composed of an input layer, one or more fully connected hidden layers with nonlinear activation functions, and an output layer, trained using backpropagation to model nonlinear input–output relationships. The MLP equation is shown in equation 2.9:

$$\mathbf{h} = \text{ReLU}(\mathbf{W}_1 \mathbf{w} + \mathbf{b}_1), \quad \hat{x} = \tanh(\mathbf{W}_2 \mathbf{h} + \mathbf{b}_2) \quad (2.9)$$

The sliding window formulation ($i \in [i - w : i + w]$) is motivated by the observation that adjacent received symbols share common noise statistics and channel conditions. While AWGN noise is i.i.d. across symbols (making the window unnecessary for a single-symbol estimator), the window provides the network with a local context that helps it learn a more robust decision boundary : effectively performing a form of implicit averaging that reduces variance. For fading channels, where adjacent symbols may experience correlated fading, the window becomes even more valuable.

Multi-SNR training is a key design choice: by training on data generated at multiple SNR levels (5, 10, 15 dB in this thesis), the network learns a denoising function that generalizes across operating conditions rather than specializing to a single noise level. This is analogous to training a robust estimator that operates well across a range of noise variances, a concept related to minimax estimation in statistical decision theory.

A critical question in applying neural networks to relay processing is the relationship between model complexity and performance. Prior work has generally assumed that larger models yield better performance, but this assumption has not been rigorously tested in the relay communication context. The bias-variance framework predicts that for a low-complexity target function like relay denoising, performance should plateau or degrade

beyond a modest model size : a prediction that this thesis confirms empirically. Understanding this relationship is essential for practical deployment, where relay nodes may have limited computational resources.

2.3 Generative Models for Signal Processing

Generative models offer an alternative paradigm for relay signal processing. Rather than directly learning a denoising function (discriminative approach), generative models learn the distribution of clean signals $p(\mathbf{x})$ and use this knowledge to reconstruct the transmitted signal from noisy observations via Bayes' rule:

$$p(\mathbf{x}|\mathbf{y}) \propto p(\mathbf{y}|\mathbf{x})p(\mathbf{x}) \quad (2.10)$$

The generative approach is theoretically appealing because it separates the modeling of the signal prior from the noise model, potentially enabling better generalization to unseen noise conditions.

2.3.1 Variational Autoencoders

Variational Autoencoders (VAEs) [17] learn a latent representation by maximizing the evidence lower bound (ELBO). The generative model posits that data $\mathbf{x}$ is generated from a latent variable $\mathbf{z} \sim p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$ through a decoder $p_\theta(\mathbf{x}|\mathbf{z})$. Since the true posterior $p(\mathbf{z}|\mathbf{x})$ is intractable, an encoder network $q_\phi(\mathbf{z}|\mathbf{x})$ approximates it. The ELBO objective is derived from the log-evidence decomposition: (Equation 2.11)

$$\log p_\theta(\mathbf{x}) = \underbrace{\mathbb{E}_{q_\phi} \left[\log \frac{p_\theta(\mathbf{x}, \mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} \right]}_{\text{ELBO}(\theta, \phi; \mathbf{x})} + \underbrace{D_{KL}(q_\phi(\mathbf{z}|\mathbf{x}) \| p_\theta(\mathbf{z}|\mathbf{x}))}_{\geq 0} \quad (2.11)$$

Since the KL divergence is non-negative, the ELBO is a lower bound on the log-evidence. Maximizing the ELBO simultaneously trains the encoder to approximate the true posterior and the decoder to reconstruct the data. The ELBO decomposes into a reconstruction term and a regularization term:

$$\mathcal{L}_{\text{VAE}} = \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})]}_{\text{Reconstruction quality}} - \underbrace{\beta \cdot D_{KL}(q_\phi(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z}))}_{\text{Latent space regularity}} \quad (2.12)$$

The reconstruction term encourages the decoder to accurately reproduce the input from the latent code, while the KL term regularizes the latent space to be close to the stan-

dard Gaussian prior $p(\mathbf{z})$. The β parameter (β -VAE) [18] controls the trade-off between reconstruction quality and latent space smoothness: (Equation 2.8) $\beta < 1$ prioritizes reconstruction (beneficial for the relay task where accurate signal recovery is paramount), while $\beta > 1$ encourages disentangled latent representations.

The *reparameterization-trick* enables gradient-based optimization through the stochastic sampling layer:

$$\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}, \text{diag}(\boldsymbol{\sigma}^2)) \quad (2.13)$$

Instead of sampling 2.13, the sample is expressed as :

$$\mathbf{z} = \boldsymbol{\mu} + \boldsymbol{\sigma} \odot \boldsymbol{\epsilon} \quad (2.14)$$

where $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. This moves the stochasticity outside the computational graph, allowing backpropagation through the encoder.

For relay signal processing, the VAE learns a compressed latent representation of the clean signal manifold. During inference, the noisy received signal is encoded into the latent space, and the decoder maps it back to a denoised estimate. The regularized latent space provides implicit denoising: noisy inputs that map to regions far from the learned signal manifold are pulled back toward high-probability regions during decoding. However, the stochastic sampling introduces additional variance into the estimate, which can be harmful for the deterministic BPSK denoising task : a trade-off that manifests as the consistent VAE under performance observed in this thesis.

2.3.2 Conditional Generative Adversarial Networks

Conditional GANs (CGANs) [19] learn signal denoising through adversarial training, building on the foundational GAN framework [20]. The original GAN formulates generative modeling as a two-player minimax game between a generator G and a discriminator D : (Equation 2.15)

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z} [\log(1 - D(G(\mathbf{z})))] \quad (2.15)$$

At the Nash equilibrium, G generates samples indistinguishable from real data and D outputs 1/2 everywhere. The conditional variant conditions both G and D on auxiliary information (the noisy signal $\mathbf{y}$), enabling the generator to learn a noise-conditioned mapping.

A fundamental challenge with the original GAN objective is *training instability*: the

Jensen-Shannon divergence that the discriminator implicitly estimates can produce vanishing gradients when the generator distribution and data distribution have disjoint supports (which is common early in training). The *Wasserstein GAN (WGAN)* addresses this by replacing the JS divergence with the Earth Mover (Wasserstein-1) distance: (Equation 2.16)

$$W(p_{\text{data}}, p_G) = \inf_{\gamma \in \Pi(p_{\text{data}}, p_G)} \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim \gamma} [\|\mathbf{x} - \mathbf{y}\|] \quad (2.16)$$

The Wasserstein distance provides meaningful gradients even when distributions do not overlap, enabling stable training. By the Kantorovich-Rubinstein duality, the Wasserstein distance can be computed via a supremum over 1-Lipschitz functions, which the critic network approximates. The *gradient penalty (GP)* formulation [21] enforces the Lipschitz constraint softly by penalizing the gradient norm of the critic at interpolated points: (Equation 2.17)

$$\text{GP} = \mathbb{E}_{\hat{\mathbf{x}} \sim p_{\hat{\mathbf{x}}}} [(\|\nabla_{\hat{\mathbf{x}}} D(\hat{\mathbf{x}})\|_2 - 1)^2] \quad (2.17)$$

where

$$\hat{\mathbf{x}} = \epsilon \mathbf{x}_{\text{real}} + (1 - \epsilon) \mathbf{x}_{\text{fake}}, \epsilon \sim \text{Uniform}(0, 1). \quad (2.18)$$

The full WGAN-GP training objectives used in this thesis are:

$$\mathcal{L}_G = -\mathbb{E}[D(G(\mathbf{y}, \mathbf{z}), \mathbf{y})] + \lambda_{L_1} \|\hat{\mathbf{x}} - \mathbf{x}\|_1 \quad (2.19)$$

$$\mathcal{L}_D = \mathbb{E}[D(G(\mathbf{y}, \mathbf{z}), \mathbf{y})] - \mathbb{E}[D(\mathbf{x}, \mathbf{y})] + \lambda_{\text{GP}} \cdot \text{GP} \quad (2.20)$$

The L_1 reconstruction loss in the generator objective serves a dual purpose: it provides a strong pixel-level reconstruction signal (complementing the adversarial signal which provides a distributional match), and it prevents *mode-collapse* (the generator converging to a single output regardless of input). For relay denoising, the L_1 term dominates at $\lambda_{L_1} = 100$, making the CGAN behave primarily as a supervised denoiser with an adversarial regularizer that encourages outputs to lie on the manifold of clean signals.

2.3.3 Generative vs. Discriminative Paradigms for Relay Processing

The application of generative models to relay processing has not been extensively studied. The key question is whether the generative inductive bias : learning the data distribution rather than just the input-output mapping : provides any advantage for the relay denoising task. For BPSK, the clean signal distribution is trivially simple (a discrete distribution on $\{-1, +1\}$), suggesting that the generative overhead may not be justified. This thesis provides the first systematic comparison of VAE and CGAN-based relay processing against classical and supervised learning methods, and the results confirm that the generative paradigm provides no significant advantage (and, in the case of VAE, a consistent disadvantage) for this particular task.

2.4 Sequence Models: Transformers, State Space Models

Recent advances in sequence modeling have produced two competing paradigms with distinct computational properties and inductive biases. Both have achieved state-of-the-art results in natural language processing, but their relative merits for physical-layer signal processing : where sequences are real-valued temporal signals rather than discrete tokens : have not been previously investigated.

2.4.1 Transformers and the Attention Mechanism

Transformers [22] use multi-head self-attention to capture global dependencies in sequences. The core attention mechanism computes a weighted combination of value vectors, where the weights are derived from the compatibility of query and key vectors:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V} \quad (2.21)$$

where $\mathbf{Q}, \mathbf{K} \in \mathbb{R}^{n \times d_k}$ and $\mathbf{V} \in \mathbb{R}^{n \times d_v}$ are obtained from the input via learned linear projections W^Q, W^K, W^V . The scaling factor $1/\sqrt{d_k}$ prevents the dot products from growing too large in magnitude, which would push the softmax into saturation regions with vanishing gradients.

Multi-head attention extends this by running h parallel attention heads, each with independent projections, and concatenating their outputs: (Equation 2.22)

$$\text{MultiHead}(\mathbf{X}) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O, \text{head}_i = \text{Attention}(\mathbf{X}W_i^Q, \mathbf{X}W_i^K, \mathbf{X}W_i^V) \quad (2.22)$$

Each head can attend to different aspects of the input: for signal processing, one head might focus on the immediate neighbors (local denoising) while another captures longer-range patterns. The total parameter count for multi-head attention is:

$$3d_{\text{model}}^2 + d_{\text{model}}^2 = 4d_{\text{model}}^2 \quad (2.23)$$

for Q, K, V projections and the output projection. The attention matrix:

$$\mathbf{A} = \text{softmax}\left(\frac{\mathbf{QK}^T}{\sqrt{d_k}}\right) \in \mathbb{R}^{n \times n} \quad (2.24)$$

computes pairwise interactions between all n positions, yielding a time and memory complexity of: $O(n^2)$. For the 11-symbol relay window used in this thesis, this is negligible ($11^2 = 121$ entries). However, the quadratic cost becomes prohibitive for longer sequences (e.g., $n = 1000$ symbols), motivating linear-time alternatives.

Positional encoding is necessary because the attention mechanism is permutation-equivariant (it treats input positions symmetrically). This thesis uses sinusoidal positional encoding: (Equation 2.25)

$$PE_{(pos, 2k)} = \sin\left(\frac{pos}{10000^{2k/d}}\right), \quad PE_{(pos, 2k+1)} = \cos\left(\frac{pos}{10000^{2k/d}}\right) \quad (2.25)$$

which injects position information into the embeddings, enabling the model to distinguish between symbols at different positions in the window.

2.4.2 Structured State Space Models

Structured state space models (SSMs) [23] are a class of sequence models derived from continuous-time linear dynamical systems. The continuous-time SSM maps an input signal $u(t) \in \mathbb{R}$ to an output $y(t) \in \mathbb{R}$ through a latent state $\mathbf{x}(t) \in \mathbb{R}^N$:

$$\dot{\mathbf{x}}(t) = \mathbf{A}\mathbf{x}(t) + \mathbf{B}u(t), \quad y(t) = \mathbf{C}\mathbf{x}(t) + Du(t) \quad (2.26)$$

where $\mathbf{A} \in \mathbb{R}^{N \times N}$ governs the state dynamics, $\mathbf{B} \in \mathbb{R}^{N \times 1}$ controls input injection, $\mathbf{C} \in \mathbb{R}^{1 \times N}$ is the output projection, and $D \in \mathbb{R}$ is the feedthrough. The connection to classical signal processing is immediate: this is a linear time-invariant (LTI) filter, and the state

space dimension N determines the order of the filter (number of poles/zeros in the transfer function).

To process discrete-time sequences, the continuous SSM is **discretized** using a step size $\Delta > 0$. Applying the zero-order hold (ZOH) discretization: (Equation 2.27)

$$\bar{\mathbf{A}} = \exp(\Delta \mathbf{A}), \quad \bar{\mathbf{B}} = (\Delta \mathbf{A})^{-1}(\exp(\Delta \mathbf{A}) - \mathbf{I}) \cdot \Delta \mathbf{B} \approx \Delta \mathbf{B} \quad (2.27)$$

yields the discrete recurrence: (Equation 2.28)

$$\mathbf{x}_k = \bar{\mathbf{A}}\mathbf{x}_{k-1} + \bar{\mathbf{B}}u_k, \quad y_k = \mathbf{C}\mathbf{x}_k + Du_k \quad (2.28)$$

The **HiPPO (High-order Polynomial Projection Operators)** framework provides a principled initialization for $\mathbf{A}$: the HiPPO-LegS matrix is designed so that the state $\mathbf{x}_k$ stores a compressed representation of the input history as coefficients of a Legendre polynomial expansion. This initialization enables long-range memory without the vanishing gradient problems of standard RNNs.

The S4 model [23] introduced the key insight that structured (diagonal or low-rank) parameterizations of $\mathbf{A}$ enable efficient computation. With a diagonal $\mathbf{A} = \text{diag}(a_1, \dots, a_N)$, the recurrence decouples into N independent scalar recurrences, each computable in $O(n)$ time. The total cost is $O(nN)$ for a length- n sequence with state dimension N : linear in sequence length.

2.4.3 Mamba: Selective State Spaces

Mamba [24] extends the SSM framework by making the parameters **input-dependent** (selective), transforming the LTI system into a linear time-varying (LTV) system:

$$\Delta_k = \text{softplus}(W_\Delta u_k + b_\Delta), \quad \mathbf{B}_k = W_B u_k, \quad \mathbf{C}_k = W_C u_k \quad (2.29)$$

where W_Δ, W_B, W_C are learned projection matrices. The selectivity mechanism allows the model to dynamically control which inputs are stored in the state and which are forgotten:

A **large** Δ_k (triggered by a relevant input) makes:

$$\bar{\mathbf{A}}_k = \exp(\Delta_k \mathbf{A}) \approx \mathbf{0} \quad (2.30)$$

which resets the state and allows the new input to dominate via:

$$\bar{\mathbf{B}}_k \approx \Delta_k \mathbf{B}_k. \quad (2.31)$$

A **small** Δ_k (triggered by irrelevant or noisy input) makes 2.32 approximately unity (I) hence, preserving the current state and ignoring the input.

$$\bar{\mathbf{A}}_k \approx \mathbf{I} \quad (2.32)$$

For relay signal processing, this selective mechanism is particularly well-suited: the model can learn to attend to the actual signal component of each received sample while suppressing the noise component, effectively implementing an *adaptive-filter* whose coefficients are conditioned on the input. This is a more general form of the Wiener filter, which is the optimal LTI denoiser but cannot adapt its coefficients on a per-sample basis.

The Mamba architecture wraps the selective S_6 layer in a gated architecture: the input is projected to $2 \times d_{\text{inner}}$ channels (expand factor 2), and split into two branches. The main branch passes through a causal 1D convolution (*Conv1D*) to mix local temporal context, followed by a *SiLU* activation and the S_6 selective scan. The second branch acts as a *SiLU* gate. The two branches are then multiplicatively merged and contracted back to d_{model} . This gating mechanism provides a multiplicative interaction that helps the model learn sharp non-linear decision boundaries.




Figure 2.1: Mamba: Linear-Time Sequence Modeling with Selective State Spaces. Structured SSMs independently map each channel (e.g. $D = 5$) of an input x to output y through a higher dimensional latent state h (e.g. $D = 4$). Prior SSMs avoid materializing this large effective state (DN times batch size B and sequence length L) through clever alternate computation paths requiring time-invariance: the (Δ, A, B, C) parameters are constant across time. Our selection mechanism adds back input-dependent dynamics, which also requires a careful hardware-aware algorithm to only materialize the expanded states in more efficient levels of the GPU memory hierarchy

2.4.4 Mamba-2: Structured State Space Duality

Mamba-2 (SSD) [25] reformulates the selective state space model through an algebraic duality between linear recurrences and structured matrix multiplications. The key theoretical insight is that the SSM output can be expressed as:

$$y_i = \sum_{j \leq i} \mathbf{C}_i^\top \left(\prod_{k=j+1}^i \bar{\mathbf{A}}_k \right) \mathbf{B}_j u_j \quad (2.33)$$

Defining the **SSM matrix** $M \in \mathbb{R}^{n \times n}$ with entries:

$$M_{ij} = \begin{cases} \mathbf{C}_i^\top \bar{\mathbf{A}}_{j+1:i} \mathbf{B}_j & i \geq j \\ 0 & i < j \end{cases} \quad (2.34)$$

the output is $\mathbf{y} = M \cdot (\mathbf{B} \odot \mathbf{u})$. The matrix M is lower-triangular (causal) and **semi-separable** (each entry factors as an outer product of left and right vectors with a diagonal product in between). This algebraic structure enables efficient chunk-parallel computation:

1. **Divide** the sequence into chunks of length L .
2. **Within each chunk**, build the $L \times L$ local SSM matrix and compute the output via

a single batched matrix multiply.

3. **Between chunks**, propagate the state once per chunk (rather than once per time step).

The per-chunk cost is $O(L^2N)$ for matrix construction and $O(L^2d)$ for the matrix-vector product, where N is the state dimension and d is the model dimension. With n/L chunks, the total cost is $O(n \cdot L \cdot N + n \cdot L \cdot d)$, which for small L (e.g., $L = 8$ or $L = 32$) is effectively $O(n)$ with a favorable constant. The critical advantage over S_6 is that the intra-chunk computation is **fully parallel**: a batched matmul rather than a sequential scan: making it much faster on GPUs where parallelism is cheap and sequential operations incur kernel-launch overhead.

The “*duality*” in the name refers to the equivalence between the recurrent view (state propagation) and the matrix view (structured matmul): the same computation can be expressed either way, and the implementation can choose whichever is more efficient for the hardware and sequence length. This thesis demonstrates this duality empirically: at $n = 11$, the recurrent S_6 view is faster; at $n = 255$, the matrix SSD view dominates.

Mamba-2 can be viewed as a streamlined evolution of Mamba-1: it restructures the block so that the state-space parameters are formed upfront (rather than being generated sequentially along the SSM pathway), and it further stabilizes training by adding an extra normalization stage in the spirit of *NormFormer* [1]. In addition, Mamba-2 emphasizes more parallel, consolidated parameter projection (e.g., producing multiple SSM-related tensors via fewer projections) and reduces per-head redundancy by sharing the BBB and CCC projections across heads—conceptually similar to multi-value attention. Figure 2.2 presents a comparison of Mamba-SSM and Mamba-2-SSD [26]:




Figure 2.2: The Mamba-2 block simplifies the Mamba block by removing sequential linear projections; the SSM parameters A , B , C are produced at the beginning of the block instead of as a function of the SSM input X . An additional normalization layer is added as in *NormFormer* [1], improving stability. The B and C projections only have a single head shared across the X heads, analogous to multi-value attention (MVA).

2.4.5 State Space Models for Signal Processing

The application of state space models to physical-layer signal processing is largely unexplored. Classical signal processing relies extensively on LTI state space models (Kalman filters, Wiener filters), and the SSM framework provides a natural neural extension of these classical tools. The connection is direct: a trained SSM with fixed (input-independent) parameters implements a learnable linear filter, while the selective Mamba variant implements a learnable *adaptive* filter. For the relay denoising task, this means Mamba can potentially learn the optimal filter structure from data, without requiring manual specification of filter order, cutoff frequencies, or adaptation algorithms.

This thesis presents the first comparison of Mamba S_6 , Mamba-2 (SSD), and Transformers for relay communication, demonstrating that state space models are well suited to this domain and that the SSD formulation of Mamba-2 yields significant speed advantages at longer context lengths.

2.5 Theoretical Foundations: Channel Models

This section presents the theoretical BER analysis for each channel model used in this study, derives the closed-form expressions, and validates them against Monte Carlo simulation. The theoretical-simulative comparison serves two purposes: (i) it verifies the correctness of the simulation framework, and (ii) it establishes the baseline performance

that AI relays must beat.

2.5.1 AWGN Channel : Theoretical Analysis

The AWGN channel adds zero-mean Gaussian noise to the transmitted signal: (Equation 2.35)

$$y = x + n, \quad n \sim \mathcal{N}(0, \sigma^2), \quad \sigma^2 = \frac{P_s}{\text{SNR}_{\text{linear}}} \quad (2.35)$$

where $P_s = \mathbb{E}[|x|^2] = 1$ for unit-power BPSK symbols. The noise variance is inversely proportional to the linear SNR.

2.5.1.1 Single-hop BER

For BPSK ± 1 symbols and real-valued AWGN with variance $\sigma^2 = 1/\text{SNR}$, an error occurs when the noise pushes the received sample past the decision boundary at the origin. The theoretical BER is: (Equation 2.36)

$$P_e^{\text{AWGN}} = Q\left(\frac{1}{\sigma}\right) = Q\left(\sqrt{\text{SNR}}\right) = \frac{1}{2} \text{erfc}\left(\sqrt{\frac{\text{SNR}}{2}}\right) \quad (2.36)$$

where $Q(x) = \frac{1}{2} \text{erfc}(x/\sqrt{2})$ is the Gaussian Q-function and $\text{SNR} = P_s/\sigma^2$ is the ratio of signal power to noise variance.

Noise-convention note. The standard textbook result $P_e = Q(\sqrt{2E_b/N_0})$ assumes that the one-sided noise PSD is N_0 , giving a baseband noise variance of $N_0/2$ per real dimension. In our AWGN implementation the noise is purely real with variance $\sigma^2 = P_s/\text{SNR}$, so $N_0 = 2\sigma^2 = 2P_s/\text{SNR}$ and $E_b/N_0 = \text{SNR}/2$. Substituting yields $Q(\sqrt{2 \cdot \text{SNR}/2}) = Q(\sqrt{\text{SNR}})$, which is the expression used throughout this work.

By contrast, the fading channels add **complex** Gaussian noise with total power $1/\text{SNR}$ (i.e. $\sigma_I^2 = \sigma_Q^2 = 1/(2 \text{SNR})$) and extract the real part after ZF equalization, halving the effective noise variance per decision dimension. This naturally introduces a factor of 2 inside the Q-function argument for all fading-channel BER expressions (the fading-channel analysis above), making them consistent with the standard textbook forms.

2.5.1.2 Two-hop AF BER

For amplify-and-forward with equal-SNR hops, the effective end-to-end SNR is:

$$\text{SNR}_{\text{eff}}^{\text{AF}} = \frac{\text{SNR}_1 \cdot \text{SNR}_2}{\text{SNR}_1 + \text{SNR}_2 + 1} = \frac{\gamma^2}{2\gamma + 1} \quad (2.37)$$

where $\gamma = \text{SNR}$ is the per-hop SNR. The resulting BER is

$$P_e^{\text{AF}} = Q\left(\sqrt{\text{SNR}_{\text{eff}}}\right) \quad (2.38)$$

under the same real-noise convention 2.5.1.1 as the single-hop case. At high SNR,

$$\text{SNR}_{\text{eff}} \approx \frac{\gamma}{2} \quad (2.39)$$

confirming the well-known 3 dB penalty of AF relaying.




Figure 2.3: AWGN channel : theoretical BER (solid lines) vs. Monte Carlo simulation (markers with 95% CI) for single-hop, two-hop AF, and two-hop DF.

2.5.1.3 Two-hop DF BER

For a decode-and-forward relay with equal-SNR hops, an error occurs at the destination if exactly one hop introduces an error:

$$P_e^{\text{DF}} = P_{e,1} + P_{e,2} - 2P_{e,1}P_{e,2} \quad (2.40)$$

2.40 df-ber-awgn

With equal hops (i.e. $P_{e,1} = P_{e,2} = P_e^{\text{AWGN}}$), this becomes

$$P_e^{\text{DF}} = 2P_e(1 - P_e) \quad (2.41)$$

At high SNR, $P_e \ll 1$ and

$$P_e^{\text{DF}} \approx 2P_e \quad (2.42)$$

, i.e., the two-hop penalty is approximately a factor of 2 in BER.

2.5.2 Rayleigh Fading Channel : Theoretical Analysis

The Rayleigh fading channel models non-line-of-sight (NLOS) propagation where the signal undergoes multiplicative fading:

$$y = hx + n, \quad h \sim \mathcal{CN}(0, 1), \quad n \sim \mathcal{CN}(0, \sigma^2) \quad (2.43)$$

The fading coefficient h is a circularly-symmetric complex Gaussian random variable, so its magnitude $|h|$ follows a Rayleigh distribution:

$$f_{|h|}(r) = 2r \cdot e^{-r^2}, \quad r \geq 0 \quad (2.44)$$

with $\mathbb{E}[|h|^2] = 1$ (unit average power). The instantaneous SNR after equalization ($\hat{x} = y/h$) becomes $\gamma_h = |h|^2 \cdot \text{SNR}$, which is exponentially distributed.

2.5.2.1 Single-hop BER

Averaging the conditional BER $P_e(\gamma_h) = Q(\sqrt{2\gamma_h})$ over the exponential distribution of γ_h yields the closed-form [21, Eq. 14-4-15] (Equation 2.45):

$$P_e^{\text{Rayleigh}} = \frac{1}{2} \left(1 - \sqrt{\frac{\bar{\gamma}}{1 + \bar{\gamma}}} \right) \quad (2.45)$$

where $\bar{\gamma} = \text{SNR}$ is the average SNR. At high SNR ($\bar{\gamma} \gg 1$):

$$P_e^{\text{Rayleigh}} \approx \frac{1}{4\bar{\gamma}} \quad (2.46)$$

This $1/\text{SNR}$ decay is fundamentally slower than the exponential decay of AWGN ($Q(\sqrt{\gamma}) \sim e^{-\gamma/2}$), explaining why Rayleigh fading is significantly more challenging. The channel is **diversity-limited**: deep fades (where $|h| \approx 0$) cause errors regardless of the average SNR. Combating this diversity limitation is the primary motivation for the relay architectures studied in this thesis.

2.5.2.2 Two-hop DF BER

The two-hop DF relay over Rayleigh fading follows the same composition rule as AWGN:

$$P_e^{\text{DF, Rayleigh}} = 2P_e^{\text{Rayleigh}}(1 - P_e^{\text{Rayleigh}}) \quad (2.47)$$




Figure 2.4: Rayleigh fading : theory vs. simulation for single-hop and two-hop DF. The characteristic $1/\text{SNR}$ slope (compared to the steeper exponential AWGN curve) is clearly visible.

2.5.3 Fading Coefficient Distributions

The statistical behavior of the fading coefficient $|h|$ determines the severity of fading for each channel model. Figure 2.5 shows the probability density function (PDF) and cumulative distribution function (CDF) of $|h|$ for Rayleigh and Rician fading with various K -factors.

Key observations from the fading PDFs:

- **Rayleigh ($K = 0$):** The PDF is maximized at $|h| = 1/\sqrt{2} \approx 0.707$ and has a heavy tail toward zero, meaning deep fades are common. The probability of a deep fade ($|h| < 0.3$) is approximately 8.6%.
- **Rician $K = 1$:** The LOS component shifts the PDF peak toward higher amplitudes, reducing the probability of deep fades.
- **Rician $K = 3$:** The distribution becomes more concentrated around the LOS amplitude ($\nu \approx 0.87$). Deep fade probability drops to approximately 0.3%.
- **Rician $K = 10$:** Approaches a near-deterministic channel; the PDF becomes sharply peaked around $|h| \approx 0.95$. Fading is negligible.

The CDF plot directly shows the **outage probability** $P(|h| \leq x)$: for a given threshold x , the CDF value gives the probability that the fading amplitude falls below that threshold.

Rayleigh has the highest outage probability at any threshold, confirming its role as the worst-case fading model.




Figure 2.5: Left : PDF of $|h|$ for Rayleigh and Rician ($K = 1, 3, 10$). Right : CDF (outage probability). Rayleigh has the highest deep-fade probability at any threshold.

2.5.4 Channel Model Validation (Simulative)

Before evaluating relay strategies, we validate the simulation framework by comparing Monte Carlo results against the closed-form theoretical BER expressions derived above. For each channel model, 20 independent trials of 50,000 bits per trial are run at each SNR point (0-20 dB, step 2 dB), yielding 1,000,000 total bits per SNR point and tight 95% confidence intervals.

Validation results:

- **AWGN:** Simulation matches theory within 95% CI at all 11 SNR points for single-hop, two-hop AF, and two-hop DF configurations.
- **Rayleigh:** Simulation confirms the theoretical $1/(4\bar{\gamma})$ high-SNR slope and matches the closed-form BER within CI bounds.

Table A summarizes the theoretical SNR required to achieve a target BER of 10^{-3} on each channel (single-hop BPSK):

Table 2.2: Theoretical SNR required to achieve $\text{BER} = 10^{-3}$ on the canonical Rayleigh channel and the AWGN calibration channel (single-hop BPSK).

Channel	SNR for $\text{BER} = 10^{-3}$	Diversity Order	High-SNR Slope
AWGN	~ 9.8 dB	: (no fading)	Exponential
Rayleigh	~ 24 dB	1	$1/(4\bar{\gamma})$

2.6 Theoretical Foundations: MIMO Equalization

Three equalization methods were implemented for the 2×2 MIMO topology:

Zero-Forcing (ZF):

$$\hat{\mathbf{x}}_{\text{ZF}} = (\mathbf{H}^H \mathbf{H})^{-1} \mathbf{H}^H \mathbf{y} = \mathbf{H}^{-1} \mathbf{y} \quad (2.48)$$

ZF completely removes inter-stream interference but amplifies noise when $\mathbf{H}$ is poorly conditioned.

MMSE:

$$\hat{\mathbf{x}}_{\text{MMSE}} = (\mathbf{H}^H \mathbf{H} + \sigma^2 \mathbf{I})^{-1} \mathbf{H}^H \mathbf{y} \quad (2.49)$$

MMSE adds a noise-variance regularization term that prevents excessive noise amplification, trading residual interference for better noise performance.

MMSE-SIC (V-BLAST):

The SIC equalizer decodes streams sequentially in order of post-detection SINR:

1. **Ordering:** Compute MMSE post-detection SINR for each stream; select the stream with highest SINR first.
2. **Detection:** Apply MMSE equalization to the selected stream and make a hard decision.
3. **Cancellation:** Subtract the detected stream's contribution: $\mathbf{y}' = \mathbf{y} - \mathbf{h}_{\text{first}} \hat{x}_{\text{first}}$.
4. **Final Detection:** Estimate the remaining stream interference-free via MRC: $\hat{x}_{\text{second}} = \text{Re}(\mathbf{h}_{\text{second}}^H \mathbf{y}') / \|\mathbf{h}_{\text{second}}\|^2$.

SIC outperforms linear MMSE because the second stream sees no inter-stream interference after cancellation. The cost is potential error propagation from incorrect first-stream decisions.

All MIMO operations are implemented using vectorized PyTorch batched `torch.linalg.solve` for GPU acceleration, achieving >100× speedup over per-symbol Python loops.

Chapter 3

Research Objectives

3.1 Main Objective

To systematically evaluate and compare classical and AI-based relay strategies for two-hop communication on a single canonical setup — SISO on both hops, i.i.d. Rayleigh fast fading, complex baseband, BPSK, uncoded BER — and to determine the optimal relay architecture as a function of SNR regime and computational constraints.

3.2 Research Gap and Motivation

Despite growing interest in AI for wireless communication, a systematic comparison of relay processing paradigms : spanning classical signal processing, supervised learning, generative modeling, and modern sequence architectures : is absent from the literature. The following specific gaps motivate this thesis:

1. **Gap 1: No cross-paradigm relay comparison.** Existing work on AI-based relay processing focuses on individual architectures in isolation. Ye et al. [7] demonstrated DNNs for OFDM detection, Dorner et al. [9] explored autoencoders for end-to-end communication, and Samuel et al. [8] applied unfolded networks to MIMO detection. However, no study systematically compares supervised feedforward networks, variational autoencoders, adversarial networks, attention-based Transformers, and state space models (Mamba) for the same relay denoising task under controlled conditions. Without such a comparison, practitioners cannot make informed architecture selections.
2. **Gap 2: Model complexity-performance relationship is unknown.** It is commonly assumed that larger neural networks yield better performance. However, the relay denoising task : recovering a BPSK symbol from a noisy observation : is a low-dimensional mapping problem. The theoretical minimum description length for this mapping is small (a soft threshold function), suggesting that large models may overfit. No prior work has rigorously characterized the complexity-performance trade-off for relay processing, nor has a normalized (equal-parameter) comparison

been conducted to isolate architectural effects from capacity effects.

3. **Gap 3: State space models for physical-layer signal processing.** The Mamba architecture [23] and its Mamba-2 SSD variant [25] have demonstrated strong performance in NLP and genomics, but their application to physical-layer wireless communication is entirely unexplored. The theoretical connection between SSMs and classical adaptive filters makes this a natural research direction, but empirical validation is needed.

This thesis addresses all three gaps through a unified framework that implements and compares nine relay strategies on the canonical setup, with both original and normalized parameter counts, full statistical analysis, and a dedicated complexity study. The following table summarizes the positioning of this work relative to the literature:

Table 3.1: Research positioning: comparison of this thesis against prior work in deep-learning-based relay communication.

Aspect	Prior Work	This Thesis
Relay strategies compared	1-2 (typically AF vs. DF, or one AI method)	9 (2 classical + 7 AI)
AI paradigms	Single (e.g., supervised only)	4 (supervised, generative, adversarial, sequential)
Evaluation discipline	Single ad-hoc configuration	One canonical setup (SISO Rayleigh BPSK), fixed in advance, with AWGN analytical calibration
Complexity analysis	None	Normalized 3K comparison + inverted-U study
State space models	None for relay/physical-layer	Mamba S_6 , Mamba-2 SSD, with crossover benchmark
Statistical rigor	Often missing	Wilcoxon test, 95% CI, 10 trials per point

3.3 Research Hypotheses

Based on the theoretical background of Chapter 2, this thesis tests the following hypotheses. H1–H4 are posed and answered on the canonical setup; H5 is an architecture-level hypothesis concerning training compute. H6 is the central hypothesis of the thesis’s principal contribution (Chapter 7), posed on channels *outside* the canonical family. The higher-order-modulation extension (Chapter 6) additionally probes the generalizability of H1–H3

to QPSK and 16-QAM and is reported descriptively:

1. **H1 (AI advantage over AF at low SNR).** On the canonical setup, some AI-based relay strategies achieve lower BER than AF at low SNR (0–4 dB), where the non-linear denoising learned by the neural network provides an advantage over AF’s noise amplification.

Rationale: At low SNR, the MMSE-optimal (posterior-mean) denoiser of a BPSK symbol from a single AWGN observation, $f^*(y) = \mathbb{E}[x | y] = \tanh(y/\sigma^2)$, is a smooth non-linear mapping that differs substantially from AF’s linear scaling; a neural network can approximate f^* closely, whereas AF amplifies noise together with the signal. Note that f^* is optimal in the mean-square sense for the relay’s local estimate; it is not claimed to be the end-to-end BER-optimal relay function, which is why H1 is posed as an empirical hypothesis rather than a theorem.

2. **H2 (DF dominance at high SNR).** The symbol-wise DF relay (Remark 1.2) achieves the lowest BER at medium-to-high SNR (≥ 6 dB), outperforming all AI methods.

Rationale: At high SNR, $f^*(y) \approx \text{sign}(y)$, which is exactly the symbol-wise DF operation. AI relays introduce a small but non-zero approximation error, so DF should be optimal among the per-symbol relay functions considered in this uncoded setting.

3. **H3 (Inverted-U complexity curve).** There exists an optimal model size for relay denoising beyond which performance degrades due to overfitting. Specifically, models with $\sim 100 - 200$ parameters achieve performance comparable to models with $10 - 100\times$ more parameters.

Rationale: The bias-variance analysis (Section 2.2.1) predicts that the low-complexity target function f^* requires minimal model capacity. Excess capacity increases variance without reducing bias.

4. **H4 (Architecture convergence at equal scale).** When all AI models are normalized to the same parameter count ($\sim 3,000$), the performance differences between architectures narrow significantly, indicating that parameter count is a more important factor than architectural choice.

Rationale: All architectures are universal approximators and the target function is simple. At equal capacity, architectural inductive biases provide diminishing returns.

5. **H5 (SSM speed advantage at long context).** Mamba-2 SSD trains significantly faster than Mamba S6 at longer context lengths ($n \gg 11$) due to chunk-parallel computation, while S6 is faster at the short relay window ($n = 11$).

Rationale: The crossover between sequential S6 ($O(n)$ serial steps) and chunk-

parallel SSD ($O(n/L)$ parallel matmuls of size L) depends on the ratio of kernel-launch overhead to arithmetic cost.

6. **H6 (Principal contribution — unknown-channel mitigation).** On channels with structure unknown to the receiver chain — memory (intersymbol interference) or asymmetry (biased nonlinearity) — the *memoryless* classical relays (AF, symbol-wise DF) fail to relay reliably, exhibiting error floors or severe degradation, whereas the minimal 170-parameter MLP relay restores reliable relaying by learning the impairment from data, approaching the performance of the optimal model-aware detector without requiring knowledge of the channel family.

Rationale: AF forwards the channel impairment untouched, and symbol-wise DF applies a fixed memoryless zero-threshold decision; neither can compensate structure they do not model. Classical theory does provide a remedy *once the channel family is known* — estimate-then-MLSE (pilot-based channel estimation followed by Viterbi equalization) is optimal for a linear FIR channel and is included as a benchmark. The hypothesis is therefore not that learning beats classical theory, but that a single fixed learned architecture can approach model-aware performance with no channel-family knowledge, no explicit identification stage, and complexity independent of channel memory — and that it remains applicable on impairments (e.g., nonlinear bias) for which linear MLSE is mis-specified. How far this holds, and on which impairment structures it breaks down, is determined empirically in Chapter 7. Conversely, on channels that satisfy the classical assumptions, H2 predicts no learned relay can beat DF, so H6 also delineates the precise boundary of the learned relay’s value.

3.4 Specific Objectives

1. **Implement and compare nine relay strategies** spanning four learning paradigms: no learning (AF, DF), supervised learning (MLP, Hybrid), generative modeling (VAE, CGAN), and sequence modeling (Transformer, Mamba S6, Mamba-2 SSD).
2. **Evaluate on the canonical setup:** SISO Rayleigh fast fading with BPSK, validated against closed-form theory (with AWGN as the analytical calibration limit), and probe modulation generalizability in a scoped QPSK/16-QAM extension.
3. **Investigate the complexity–performance trade-off** by testing model architectures ranging from 0 parameters (classical) to 26,179 parameters (Mamba-2 SSD), and by conducting a normalized comparison at approximately 3,000 parameters.
4. **Determine whether state space models outperform attention mechanisms** for relay signal processing by comparing Mamba S6 and Transformer architectures at

both original and normalized parameter counts, and by benchmarking at extended context lengths.

5. **Identify practical deployment recommendations** for selecting the appropriate relay strategy given specific operational constraints (SNR range, computational budget, channel environment).

3.5 Scope and limitations

The following limitations define the scope of this study, which is confined entirely to the canonical setup (SISO, Rayleigh fast fading, complex baseband, BPSK, uncoded):

- **Modulation:** The core experiments use BPSK. A scoped extension (Chapter 6) evaluates QPSK and 16-QAM, including the joint 2D N -class classification formulation; 16-PSK, 64-QAM and denser constellations are deferred to future work.
- **Channel knowledge:** Perfect CSI is assumed at the receiver. Channel estimation errors are not modeled.
- **Relay topology:** Single relay, two-hop, half-duplex. Multi-relay and full-duplex configurations are excluded.
- **Topology:** SISO on both hops. Multi-antenna (MIMO) configurations are deferred to future work.
- **Training regime:** Offline training on synthetic data. Online adaptation is not implemented.

These delimitations are chosen to enable a clean comparison of relay processing functions in isolation, without confounding factors from protocol-level or system-level complexity.

Chapter 4

Methods

4.1 System Model

The system under study is a two-hop relay network with a single relay node. All methods in this chapter are described for the canonical setup of Chapter 1: SISO on both hops, i.i.d. Rayleigh fast fading, complex baseband, BPSK, uncoded BER.




Figure 4.1: Two-hop relay system model: source transmits through two channels with a neural network relay in between.

Source bits are BPSK-modulated, passed through Hop 1 (SISO Rayleigh), processed by the relay (classical or neural), transmitted over Hop 2 (SISO Rayleigh), and demodulated at the destination; no equalization is involved. Both hop channels are known perfectly at their respective receivers, per the conventions of Chapter 1.

$$\text{Source} \xrightarrow{\text{Hop 1}} \text{Relay} \xrightarrow{\text{Hop 2}} \text{Destination} \quad (4.1)$$

4.1.1 Modulation.

The modulation throughout this thesis is BPSK, per the canonical setup: Binary Phase-Shift Keying maps a single bit to a real-valued symbol,

$$x = 1 - 2b, \quad b \in \{0, 1\} \implies x \in \{-1, +1\} \quad (4.2)$$

with hard-decision demodulation $\hat{b} = \mathbb{K}[\hat{x} < 0]$ at the destination. The QPSK and 16-QAM extensions are defined in Section 4.5 and evaluated in Chapter 6; further constellations are future work.

4.1.2 Hop Model

Each hop applies a channel function followed by optional equalization (Equation 4.3):

$$y = h(x, \text{SNR}) + n \quad (4.3)$$

where $h(\cdot)$ is the canonical Rayleigh fast-fading channel (the AWGN special case is retained for analytical calibration of the simulator).

Power Normalization. All relay strategies normalize their output power to ensure fair comparison:

$$x_R \leftarrow x_R \cdot \sqrt{\frac{P_{\text{target}}}{P_{\text{current}}}} \quad (4.4)$$

4.1.3 Relay Processing

On Hop 1 the relay's single receive antenna observes, as shown in Equation 4.5:

$$y_R = x + n_1, \quad n_1 \sim \mathcal{N}(0, \sigma^2) \quad (4.5)$$

where the effective observation follows coherent compensation of the Rayleigh fading coefficient (Section 4.1). The relay applies its processing function to a sliding window of received samples and transmits a cleaner estimate $\hat{x}_R = f_\theta(y_{R,i-w:i+w})$; classical AF/DF use $w = 0$. This is purely a noise-removal task: in the canonical SISO setup there is no inter-stream interference at any stage, and the destination demodulates the Hop-2 output directly with no equalizer.

4.2 Relay Strategies

Nine relay strategies were implemented, spanning four learning paradigms. The selection of these nine strategies is designed to systematically explore the relay design space along three axes:

1. the degree of processing (from no learning to deep generative models).
2. the type of inductive bias (feedforward, recurrent, attention, generative).
3. the model capacity (from 0 to 26K parameters). This section describes each strat-

egy’s architecture, training procedure, and the design rationale motivating each choice.

4.2.1 Selection and Classification of Relay Strategies

The selected nine relay strategies are drawn from four distinct architectural concepts:

(i) classical methods (AF and DF), (ii) supervised-learning models (MLP and hybrid MLP–DF), (iii) generative models (VAE and CGAN), and (iv) sequential models (Transformers, Mamba-SSM, and Mamba-SSD).

1. **Classical AF:** Gain amplification as shown in equation 2.1. AF serves as the lower baseline. it performs no intelligent processing and simply rescales the received signal, preserving both signal and noise. Its end-to-end effective SNR is given in Equation 2.2.
2. **Classical DF:** Demodulates, recovers bits and re-modulates clean BPSK symbols. DF serves as the upper classical baseline: it performs the maximum possible classical processing (full signal regeneration). At high SNR, DF approaches error-free relay operation. Its theoretical BER follows the error-composition formula of Equation 2.4.
3. **Supervised MLP (Minimal):** A two-layer feedforward neural network (multi-layer perceptron, MLP) with 169 parameters following the relay architecture defined in Equation 2.9,

where $\mathbf{w} \in \mathbb{R}^5$ is a sliding window of received symbols ($w = 2$ neighbors on each side), and the hidden layer has 24 neurons. Parameters: $(5 \times 24 + 24) + (24 \times 1 + 1) = 169$.

Architecture note: The MLP relay is a standard discriminative dual-layer perceptron trained with supervised learning (MSE loss on input–output pairs). In contrast, the generative models in this study are the VAE (§6.3, Generative Models) and CGAN, which learn to *sample from* or *approximate* the data distribution. The MLP relay simply learns a deterministic mapping $f : \mathbb{R}^5 \rightarrow [-1, +1]$ from a noisy observation window to a denoised symbol estimate : a classical regression task.

Design rationale: The tanh output activation naturally constrains the output to $[-1, +1]$, matching the BPSK symbol range. The ReLU hidden layer provides the non-linearity needed to approximate the Bayes-optimal soft threshold $\tanh(y/\sigma^2)$. With 24 hidden neurons and a 5-dimensional input, the network has approximately 34 parameters per input dimension : sufficient for the low-complexity denoising task while avoiding overfitting. He initialization is used for ReLU layers to main-

tain proper gradient flow.

Training uses MSE loss with multi-SNR training data (i.e.) 5, 10 and 15 dB), 25,000 samples, and 100 epochs with a learning rate of 0.01.

4. **Hybrid:** SNR-adaptive relay that switches between MLP (low SNR) and DF (high SNR) based on a learned threshold. Combines the AI advantage at low SNR with the zero-error classical approach at high SNR. Same 169 parameters as MLP.

Design rationale: The Hybrid relay is motivated by the observation (confirmed empirically) that AI relays outperform DF only at low SNR, while DF is optimal at high SNR. By introducing a switching threshold, the Hybrid relay automatically selects the better strategy for each operating condition. The threshold is determined empirically from the training data by finding the SNR at which MLP and DF BER curves cross. This approach requires no additional parameters beyond those of the MLP sub-network.

Training points at the relay node: SNR = **5, 10, 15 dB** (same trained MLP sub-network as the Minimal MLP relay).

5. **Generative VAE:** Probabilistic relay with encoder $q_\phi(\mathbf{z}|\mathbf{x})$ mapping to a latent space and decoder $p_\theta(\mathbf{x}|\mathbf{z})$ reconstructing the signal. Architecture: encoder ($7 \rightarrow 32 \rightarrow 16 \rightarrow \mu, \sigma^2(8)$), decoder ($8 \rightarrow 16 \rightarrow 32 \rightarrow 1$). Total: 1,777 parameters. Trained with β -VAE loss ($\beta = 0.1$) for 100 epochs.

Training points at the relay node: SNR = **5, 10, 15 dB**.

Design rationale: The low $\beta = 0.1$ prioritizes reconstruction quality over latent space regularization, appropriate for a task where accurate signal recovery is paramount. The 8-dimensional latent space provides sufficient representational capacity for the BPSK signal manifold while maintaining a tractable KL divergence. The encoder window of 7 symbols provides local context.

6. **Generative CGAN (WGAN-GP):** Adversarial relay with a generator conditioned on the noisy signal and a critic providing the training signal. Generator: ($7 + 8 \rightarrow 32 \rightarrow 32 \rightarrow 16 \rightarrow 1$), Critic: ($1 + 7 \rightarrow 32 \rightarrow 16 \rightarrow 1$). Total: 2,946 parameters. Trained with Wasserstein loss, gradient penalty ($\lambda = 10$), and L1 reconstruction loss ($\lambda_{L1} = 100$) for 200 epochs.

Training points at the relay node: SNR = **5, 10, 15 dB**.

Design rationale: The high $\lambda_{L1} = 100$ ensures that the generator is primarily supervised by the reconstruction loss, with the adversarial term acting as a regularizer that encourages outputs to lie on the clean signal manifold. The noise vector $\mathbf{z} \in \mathbb{R}^8$ provides the stochastic input needed by the GAN framework. The 5:1 critic-to-generator update ratio follows the WGAN-GP recommendation for stable

critic training. The doubled epoch count (200 vs. 100) compensates for the slower per-step convergence of adversarial training.

7. **Sequential models: Transformer:** Multi-head self-attention over a window of 11 symbols. Architecture: $d_{\text{model}} = 32$, 4 attention heads, 2 encoder layers, feedforward dimension 128. Total: 17,697 parameters. Trained for 100 epochs with Adam optimizer ($\text{lr} = 10^{-3}$).

Training points at the relay node: SNR = **5, 10, 15 dB**.

Design rationale: The Transformer is included as the dominant sequence architecture in modern deep learning. The 11-symbol window provides the same temporal context as the S6/SSD models. With $d_{\text{model}} = 32$ and 4 heads, each head operates in an 8-dimensional subspace, which is sufficient for capturing the local noise structure. The feedforward dimension of 128 ($4 \times d_{\text{model}}$) follows the standard Transformer expansion ratio.

8. **Sequential models: Mamba S6** Selective state space model with input-dependent state transitions. Architecture: $d_{\text{model}} = 32$, $d_{\text{state}} = 16$, 2 Mamba blocks with residual connections. Total: 24,001 parameters. Each block applies: LayerNorm $\rightarrow$ expand ($32 \rightarrow 64$) $\rightarrow$ split to Conv1D/SiLU/S6 main branch and SiLU gate $\rightarrow$ contract ($64 \rightarrow 32$) $\rightarrow$ residual. Trained for 100 epochs with Adam optimizer ($\text{lr} = 10^{-3}$).

Training points at the relay node: SNR = **5, 10, 15 dB**.

Design rationale: The expand factor of 2 ($32 \rightarrow 64$) doubles the internal dimension during the S6 scan, providing richer state dynamics. The state dimension $N = 16$ means each S6 layer implements a 16th-order adaptive filter : substantially more expressive than classical Wiener or matched filters. LayerNorm and residual connections prevent gradient degradation across layers. The SiLU gating provides multiplicative interactions that help the model learn sharp decision boundaries.

9. **Sequential models: Mamba2 (SSD)** Structured State Space Duality model that replaces the sequential S6 recurrence with a chunk-parallel structured matrix multiply. Architecture: $d_{\text{model}} = 32$, $d_{\text{state}} = 16$, chunk size 8, 2 Mamba-2 blocks with SiLU gating and residual connections. Total: 26,179 parameters. Each block applies: LayerNorm $\rightarrow$ parallel gate/SSD branches $\rightarrow$ SiLU gate $\rightarrow$ contract ($64 \rightarrow 32$) $\rightarrow$ residual. The SSD layer builds a lower-triangular causal kernel M per chunk and applies it via batched matmul, with inter-chunk state passing for continuity. Trained for 100 epochs with Adam optimizer ($\text{lr} = 10^{-3}$) and gradient clipping ($\|\nabla\| \leq 1$).

Training points at the relay node: SNR = **5, 10, 15 dB**.

Design rationale: The chunk size of 8 is chosen to balance parallelism (larger

chunks = fewer sequential inter-chunk passes) against memory cost (the $L \times L$ SSM matrix scales quadratically with chunk size). At the 11-symbol window, this yields 2 chunks (8 + 3), with a single inter-chunk state pass. Gradient clipping at norm 1.0 prevents the exponential blowup of gradients through the cumulative matrix products in the SSD computation. The S4D-style initialization $A_{\log} = \log(1, 2, \dots, N)$ provides geometrically spaced decay rates, enabling the model to simultaneously capture short-range and medium-range dependencies.

4.3 Simulation Framework

4.3.1 Monte Carlo BER Estimation

BER is estimated through Monte Carlo simulation, which provides an unbiased estimate of the true BER at each SNR point. The simulation parameters are:

- **Bits per trial:** 10,000
- **Trials per SNR:** 10 (independent random seeds)
- **Total bits per SNR point:** 100,000
- **SNR range:** 0 to 20 dB (step: 2 dB), yielding 11 evaluation points
- **Random seed control:** Each trial uses a unique seed for bit generation and noise realization

BER Computation:

$$\hat{P}_e = \frac{1}{N} \sum_{i=1}^N \mathbf{1}(b_i \neq \hat{b}_i) \quad (4.6)$$

where $\mathbf{1}(\cdot)$ is the indicator function and $N = 10,000$ is the number of bits per trial. By the central limit theorem, for large N and moderate BER (say $P_e \geq 10^{-3}$), the per-trial BER estimate is approximately normally distributed with variance $P_e(1 - P_e)/N$.

Confidence Intervals. The 95% confidence interval is computed from the $M = 10$ independent trial estimates as:

$$\text{CI}_{95\%} = \bar{P}_e \pm t_{0.025, M-1} \cdot \frac{s}{\sqrt{M}} \quad (4.7)$$

where $\bar{P}_e$ is the sample mean BER across trials, s is the sample standard deviation, and $t_{0.025, 9} = 2.262$ is the critical value of the Student's t -distribution with 9 degrees of freedom. At low BER ($\hat{P}_e \lesssim 10^{-4}$), the normal approximation may become unreliable due to

the small number of observed errors; however, in this regime the relay methods converge and the relative ranking is stable.

BER resolution. With $N \cdot M = 100,000$ total bits per SNR point, the minimum detectable BER is approximately $1/N = 10^{-4}$ per trial, or 10^{-5} for the aggregate. BER values below this threshold are reported as 0.

4.3.2 Statistical Significance Testing

Differences between relay methods are assessed using the **Wilcoxon signed-rank test**, a non-parametric paired test. For each pair of relay strategies (A, B) at each SNR point, the test compares the $M = 10$ paired BER observations (Equation 4.8):

$$H_0 : \text{median}(P_{e,A} - P_{e,B}) = 0 \quad \text{vs.} \quad H_1 : \text{median}(P_{e,A} - P_{e,B}) \neq 0 \quad (4.8)$$

The Wilcoxon test is preferred over the parametric paired t -test for two reasons: (i) BER distributions are bounded $([0, 0.5])$ and potentially skewed, violating the normality assumption, and (ii) the test is robust to outlier trials. The significance level is set at $\alpha = 0.05$.

4.3.3 Training Protocol

All AI relays are trained **once** at the beginning of each experiment using: - **Multi-SNR training data:** Signals generated at SNR = 5, 10, 15 dB in equal proportions - **Training samples:** 25,000 (supervised), 20,000 (generative), 10,000 (sequence models) - **Single training per channel:** The same trained model is evaluated across all SNR points

The multi-SNR training protocol is critical: training at a single SNR would produce a model that performs well at that SNR but poorly at others (overfitting to a specific noise level). By training at three representative SNR values spanning the low-to-moderate range, the network learns a robust denoising function that generalizes across operating conditions. The SNR values 5, 10, 15 dB were chosen to cover the regime where AI relays provide the most benefit (low-to-medium SNR).

Impact of sparse SNR training grid. Because models are trained at only three SNR points ($\{5, 10, 15\}$ dB) but evaluated over eleven points (0 to 20 dB, step 2), performance reflects two regimes: (i) **interpolation** within the trained span (approximately 4–16 dB), where BER is generally stable, and (ii) **extrapolation** at the edges (0–2 dB and 18–20 dB), where mild degradation may occur due to noise-statistics mismatch. In the present results, this sparse-grid effect is secondary: the relay ranking remains consistent across the full evaluated SNR range.

Training method (what is trained, where it is trained). Training is performed **offline** before BER sweeps. Only the seven AI relays are optimized (MLP, Hybrid’s MLP sub-network, VAE, CGAN, Transformer, Mamba S6, Mamba-2 SSD); AF and DF are analytical baselines and have no trainable parameters. Training data are synthetically generated source/channel pairs under the multi-SNR protocol, with model-specific losses (MSE for supervised relays, ELBO for VAE, WGAN-GP + L1 for CGAN, Adam-based supervised loss for sequence models). Compute placement follows implementation: NumPy models (MLP/Hybrid) train on CPU, while PyTorch models train on CPU or CUDA depending on device availability/configuration. After training, weights are checkpointed and reused across all SNR evaluations; BER curves are then produced in **inference-only** mode without further parameter updates.

Weight initialization. He initialization [27] is used for all layers with ReLU activation ($W \sim \mathcal{N}(0, 2/n_{\text{in}})$), while Xavier initialization is used for layers with tanh activation ($W \sim \mathcal{N}(0, 1/n_{\text{in}})$). These initialization schemes maintain proper gradient magnitude through the network layers, preventing both vanishing and exploding gradients during training.

Optimizer settings. All PyTorch-based models use the Adam optimizer with $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$ (defaults). The MLP and Hybrid relays use a simple SGD implementation in NumPy with constant learning rate $\eta = 0.01$.

4.3.4 Reproducibility and Weight Management

All experiments use controlled random seeds at both the source (bit generation) and noise (per-trial seeding) levels to ensure reproducible results. Specifically, for seed s and trial t , the random state is initialized as $\text{seed}(s \cdot 1000 + t)$, ensuring that each trial is independent while the overall experiment is deterministic.

A **weight checkpoint system** saves trained model parameters to disk after training completes, enabling experiment resumption without retraining. This is particularly important for the CGAN (2-hour training time) and sequence models (24–37 minutes). The checkpoint system supports: - Automatic save after training with architecture metadata - Resume from saved weights with architecture validation - Seed-specific weight directories (e.g., `trained_weights/seed_42/`)

The framework includes 126 automated tests (pytest) covering all modules: channels, modulation, relay strategies, simulation, statistics, and weight management, with 100% pass rate.

4.4 Normalized Parameter Comparison

A fundamental confound in comparing AI architectures is that different models have vastly different parameter counts: from 169 (MLP) to 26,179 (Mamba-2 SSD) : a ratio of 155:1. Performance differences between these models could reflect either (a) the superiority of one architecture’s inductive bias, or (b) the simple advantage of having more learnable parameters. To disentangle these two effects, all seven AI models were scaled to approximately 3,000 parameters, providing a controlled comparison where model capacity is held constant.

Choice of target parameter count. The target of ~3,000 parameters was chosen as a compromise: it is large enough that all architectures can express a reasonable denoising function (avoiding underfitting), yet small enough that the normalization represents a meaningful reduction for the larger models (Transformer: 5.9× reduction, Mamba S6: 7.9×, Mamba-2: 8.7×). The MLP model is scaled **up** from 169 to 3,004 parameters, testing whether additional capacity benefits the simple feedforward architecture.

Scaling methodology. Each architecture’s hyperparameters were adjusted to reach the target while preserving its essential structure:

Table 4.1: Normalized 3K-parameter model configurations: parameter count and scaling method for each architecture.

Model	Parameters	Scaling Method
MLP-3K	3,004	Increased window (5 → 11) and hidden (24 → 231)
Hybrid-3K	3,004	Same as MLP-3K + DF switching
VAE-3K	3,037	Increased window (7 → 11), adjusted latent/hidden
CGAN-3K	3,004	Increased window (7 → 11), adjusted hidden layers
Transformer-3K	3,007	Reduced d_model (32 → 18), heads (4 → 2), layers (2 → 1)
Mamba-3K	3,027	Reduced d_model (32 → 16), d_state (16 → 6), layers (2 → 1)
Mamba2-3K	3,004	Reduced d_model (32 → 15), d_state (16 → 6), layers (2 → 1)

Parameter counting convention. All learnable parameters are counted, including biases, normalization parameters (LayerNorm gain/bias), and projection matrices. Non-learnable buffers (e.g., positional encoding tables, fixed $\mathbf{A}$ matrices) are excluded. The counts are verified programmatically via `sum(p.numel() for p in model.parameters() if p.requires_grad)`.

Unified input window. All 3K models use a window size of 11 (vs. 5 for original MLP/Hybrid, and 7 for original VAE/CGAN), providing a common input context. This ensures that differences in performance reflect architectural inductive biases rather than differences in the amount of input information available.

This normalization isolates the effect of architectural choice from the confound of parameter count, providing insights into which inductive biases are most beneficial for the relay denoising task. If performance converges at equal parameter budgets, the conclusion is that parameter count : not architecture : is the dominant factor. If significant gaps persist, the conclusion is that architectural inductive biases provide meaningful advantages beyond raw capacity.

4.5 Extension Methods: QPSK and 16-QAM

The core experiments use BPSK exclusively, per the canonical setup. To test whether the canonical findings generalize to complex constellations, the extension study of Chapter 6 evaluates QPSK and 16-QAM. This section defines the two constellations, the I/Q-splitting technique that lets real-valued relays process complex signals, and the joint 2D classification formulation that removes I/Q splitting’s structural limitation.

4.5.1 QPSK : Gray-Coded Quadrature Phase-Shift Keying

Quadrature Phase-Shift Keying maps pairs of bits (b_0, b_1) to complex symbols [28]:

$$x = \frac{(1 - 2b_0) + j(1 - 2b_1)}{\sqrt{2}} \quad (4.9)$$

yielding four constellation points at $\{(\pm 1 \pm j)/\sqrt{2}\}$ with unit average power ($E_s = 1$). The Gray coding ensures that adjacent constellation points differ by exactly one bit:

Table 4.2: QPSK Gray-coded symbol mapping: bit pairs to complex symbols.

Bit pair (b_0, b_1)	Symbol	Quadrant
00	$(+1 + j)/\sqrt{2}$	I
01	$(+1 - j)/\sqrt{2}$	IV

Bit pair (b_0, b_1)	Symbol	Quadrant
11	$(-1 - j)/\sqrt{2}$	III
10	$(-1 + j)/\sqrt{2}$	II

Demodulation applies independent sign decisions on each component: $\hat{b}_0 = \mathbb{K}(\text{Re}(\hat{x}) < 0)$ and $\hat{b}_1 = \mathbb{K}(\text{Im}(\hat{x}) < 0)$. The spectral efficiency is 2 bits/symbol, double that of BPSK.

Theoretical QPSK BER. For uncoded QPSK over an AWGN channel, the BER per bit equals the BPSK BER at the same E_b/N_0 because each I/Q component carries an independent BPSK stream [28]:

$$P_b^{\text{QPSK}} = Q\left(\sqrt{\frac{2E_b}{N_0}}\right) = P_b^{\text{BPSK}} \quad (4.10)$$

The advantage of QPSK is doubled throughput for the same BER and per-bit energy.

4.5.2 16-QAM : Gray-Coded Quadrature Amplitude Modulation

16-QAM maps groups of four bits (b_0, b_1, b_2, b_3) to one of 16 complex constellation points arranged on a 4×4 rectangular grid (Figure 4.2c). Each axis (I and Q) uses independent Gray-coded PAM-4 mapping [28]:

$$I = \frac{L(b_0, b_1)}{\sqrt{10}}, \quad Q = \frac{L(b_2, b_3)}{\sqrt{10}}, \quad x = I + jQ \quad (4.11)$$

where $L(\cdot)$ maps bit pairs to PAM-4 levels using Gray coding:

Table 4.3: 16-QAM PAM-4 Gray-coded level mapping for in-phase and quadrature components.

Bit pair	Level	Bit pair	Level
00	+3	11	-1
01	+1	10	-3

The normalization factor $\sqrt{10}$ ensures unit average symbol power: $E[|x|^2] = \frac{2 \cdot (9+1+1+9)}{4 \cdot 10} = 1$. Adjacent constellation points differ by one bit (Gray property), minimizing the BER for a given symbol error rate.

Demodulation quantises each received component to the nearest PAM-4 level using decision boundaries at $\{-2, 0, +2\}/\sqrt{10}$ and maps back to bits via the inverse Gray table.

The spectral efficiency is 4 bits/symbol.

Theoretical 16-QAM BER. The approximate BER for 16-QAM over AWGN is: (Equation 4.12)

$$P_b^{16\text{-QAM}} \approx \frac{3}{8} \operatorname{erfc} \left(\sqrt{\frac{2E_b}{5N_0}} \right) \quad (4.12)$$

At the same E_b/N_0 , 16-QAM has a higher BER than BPSK or QPSK due to the reduced Euclidean distance between constellation points. The trade-off is 4× throughput improvement.

4.5.3 Constellation Diagram Summary

Figure 4.2 presents the constellation diagrams for the modulation schemes appearing in this thesis. The progression from BPSK (2 points on the real axis) through QPSK (4 points on the unit circle) to 16-QAM (16 points on a rectangular grid) illustrates the fundamental trade-off between spectral efficiency and noise robustness: higher-order constellations pack more bits per symbol but reduce the minimum Euclidean distance between points, increasing the BER at a given SNR.




Figure 4.2: Constellation diagrams. (a) BPSK: 2 real-valued symbols at ± 1 . (b) QPSK: 4 Gray-coded symbols on the unit circle. (c) 16-QAM: 16 Gray-coded symbols on a 4×4 rectangular grid with decision boundaries (dotted lines). (d) 16-PSK, shown for completeness only; it is not studied in this work. All constellations are normalised to unit average power.

4.5.4 I/Q Splitting for AI Relay Processing of Complex Constellations

$$\hat{x}_R = f_\theta(\text{Re}(y_R)) + j \cdot f_\theta(\text{Im}(y_R)) \quad (4.13)$$

Justification and caveat. For rectangular constellations (QPSK, QAM), the I and Q components carry independent information and are corrupted by independent noise, so per-axis denoising with a shared real-valued function is a natural and parameter-efficient decomposition. It is, however, a *restriction* of the hypothesis class to axis-separable functions, not an equivalence to joint 2D processing: per-axis errors accumulate independently, which is precisely the structural limitation quantified and then removed by the joint 2D classification formulation of Section 4.5.5.

Relay-specific handling:

Table 4.4: Relay-specific complex signal handling strategies and rationale.

Relay type	Complex signal processing	Rationale
AF	Amplifies complex signal directly	Power normalization ($\ y\ ^2$) is valid for complex vectors
AI relays	I/Q splitting (process Re and Im separately)	Real-valued networks; independence of I/Q in rectangular constellations

Limitation for 16-QAM with AI relays. For 16-QAM, each I/Q component takes four amplitude levels ($\{-3, -1, +1, +3\}/\sqrt{10}$) rather than the binary $\{\pm 1\}$ of BPSK. The BPSK-trained relays, which use tanh activations bounded in $[-1, +1]$, may not faithfully reproduce the multi-level structure. This provides a natural test of generalisation: if AI relays degrade significantly on 16-QAM but not on QPSK, it indicates that the BPSK training generalises to binary-per-component signals (QPSK) but not to multi-level signals (16-QAM). Such a finding would motivate modulation-specific relay training.

4.5.5 2D Joint Classification as an Alternative to I/Q Splitting

The I/Q splitting approach (Section 3.7.6) treats each axis of the complex constellation independently, classifying each component into $\sqrt{M}$ amplitude levels. For 16-QAM this yields two independent 4-class problems. While valid under the assumption of I/Q independence, this formulation introduces a **structural limitation**: (Equation 4.14) classification errors on the I and Q axes accumulate independently, producing a BER floor that cannot be reduced regardless of model capacity.

An alternative formulation treats the relay as a **joint 2D classifier** over all M constellation points simultaneously. Instead of splitting the complex signal and classifying 4 levels per axis, the relay receives the full 2D input (y_I, y_Q) and outputs M logits : one per constellation point. The predicted class index $\hat{k} = \arg \max_k z_k$ is mapped back to the corresponding complex symbol $s_{\hat{k}}$ for retransmission:

$$\hat{x}_R = s_{\hat{k}}, \quad \hat{k} = \arg \max_{k \in \{1, \dots, M\}} f_{\theta}(y_I, y_Q)_k \quad (4.14)$$

where $f_{\theta} : \mathbb{R}^2 \rightarrow \mathbb{R}^M$ is the neural relay network with M output logits trained using cross-entropy loss against the true constellation index.

Advantages over I/Q splitting: - Joint decision boundaries. The network learns decision regions in the full 2D constellation space, capturing the Voronoi structure of the

constellation rather than axis-aligned partitions. - **No structural BER floor.** Because the classifier selects among all M points jointly, there is no independent per-axis error accumulation (Equation 4.13). - **Minimal parameter overhead.** Only the final classification layer changes : from $\sqrt{M}$ to M outputs : adding negligible parameters to the larger architectures ($< 1\%$ for Transformer, Mamba S6, Mamba-2).

Trade-off. The 2D classifier requires training on 16-QAM data (it cannot reuse BPSK-pretrained weights), and the number of output classes grows as M rather than $\sqrt{M}$, which may affect convergence for very high-order constellations. For 16-QAM ($M = 16$), this trade-off is favourable. The experimental evaluation is presented in Section 6.2.

Chapter 5

Core Experiments: The Canonical Setup

5.1 Experiments summary

Following the single-model convention fixed in Chapter 1, this chapter evaluates all nine relay strategies on the *canonical setup only*: a two-hop SISO link, i.i.d. Rayleigh fast fading on both hops, complex baseband, BPSK, uncoded BER. Three experiments constitute the core: the simulator validation (E1), the canonical relay comparison (E2), and the parameter-normalization and complexity study (E3), which together answer every hypothesis (H1–H5). A scoped modulation extension (Chapter 6) generalizes the canonical findings to QPSK and 16-QAM, including the joint 2D N -class classification formulation (E4–E5). The principal contribution (Chapter 7) is the unknown-channel study (E6), which tests hypothesis H6. All other departures from the canonical setup are identified as future work in Chapter 8. The experimental configurations are summarised in the table below. Detailed numerical results, tables, and figures are provided exclusively in the appendices. Each experiment subsection therefore focuses solely on the key conclusions.

Table 5.1: Summary of experimental configurations (results reported in appendices)

Exp.	Topology	Modulation	Channel / Equalization	Focus
E1	SISO	BPSK	Rayleigh (AWGN as analytical calibration)	Channel model validation
E2	SISO	BPSK	Rayleigh	Canonical relay comparison (H1, H2)
E3	SISO	BPSK	Rayleigh	Parameter normalization & complexity (H3, H4); training-time benchmark (H5)

Exp.	Topology	Modulation	Channel / Equalization	Focus
E4	SISO	QPSK, 16-QAM	AWGN, Rayleigh	Extension: higher-order modulation via I/Q splitting (Chapter 6)
E5	SISO	16-QAM	AWGN	Extension: joint 16-class 2D classification (Chapter 6)
E6	SISO	BPSK	Unknown (ISI; composite; blind; partial-posterior); control: Rayleigh	Extension: unknown-channel failure of AF/DF and learned mitigation (H6)

5.2 Channel Model Validation

Prior to evaluating relay strategies, we first verify that the simulation framework produces bit error rate (BER) results consistent with theoretical predictions. The canonical channel of this thesis is Rayleigh fast fading, for which closed-form single-hop and two-hop BER expressions exist; the AWGN channel is additionally validated because it is the analytical limit underlying those expressions (and the noise convention), so agreement on both channels is the cleanest guarantee of baseline accuracy. The channel validation plots presented in Section 5.2 confirm the following:

- **AWGN channel:** Theoretical and simulated BER curves match exactly. The classical $Q(\sqrt{\gamma})$ expression (see the noise-convention note on page 21), as well as its two-hop extensions—including the effective SNR for amplify-and-forward (AF) relaying and the error-composition model for decode-and-forward (DF)—are validated within tight 95% confidence intervals across all SNR points (Figure 2.3).
- **Rayleigh fading:** The simulation confirms the high-SNR asymptotic behavior $\text{BER} \sim \frac{1}{4\gamma}$. The simulated BER closely follows the analytical curve (Figure 2.4)

$$\frac{1}{2} \left(1 - \sqrt{\frac{\gamma}{1 + \gamma}} \right),$$

- **Fading statistics:** Empirical probability density functions (PDFs) confirm the assumed channel distributions. Rayleigh fading exhibits the broadest spread and the highest probability of deep fades, while increasing the Rician K -factor progressively concentrates the distribution around the LOS amplitude (Figure 2.5).

5.2.1 Experiments

Goal: Validate the simulation framework against closed-form theoretical BER expressions to ensure baseline accuracy before evaluating AI relays.

Trials: - **Topology:** SISO (both hops) - **Modulation scheme:** BPSK - **Channel:** i.i.d. Rayleigh fast fading (canonical), AWGN (analytical calibration) - **Equalizers:** None - **Demod:** Hard decision only

Conclusion: Monte Carlo simulations match theoretical predictions within 95% confidence intervals on both validated channels: AWGN follows the expected exponential decay (Equation 2.36) and the canonical Rayleigh channel validates the $1/(4\bar{\gamma})$ high-SNR slope (Equation 2.46).

5.2.2 Results tables

Table 5.2: BER comparison between theoretical analysis and simulation results at selected SNR values

Chan.	4dB(Th.)	4dB(Sim.)	10dB(Th.)	10dB(Sim.)	16dB(Th.)	16dB(Sim.)
AWGN	5.64E-2	5.62E-2	7.83E-4	7.90E-4	7.73E-8	≈ 0
Rayleigh	9.10E-2	9.13E-2	2.33E-2	2.32E-2	7.58E-3	7.60E-3

5.2.3 Results Figures




Figure 5.1: AWGN Channel : Theoretical vs. Simulative BER. Single-hop, two-hop AF, and two-hop DF: closed-form theory (solid lines) versus Monte Carlo simulation (markers with 95% CI). Theory and simulation match within the confidence interval at all SNR points.




Figure 5.2: Rayleigh fading : theory vs. simulation for single-hop and two-hop DF. The characteristic $1/\text{SNR}$ slope (compared to the steeper exponential AWGN curve) is clearly visible.




Figure 5.3: Left : PDF of $|h|$ for Rayleigh and Rician ($K = 1, 3, 10$). Right : CDF (outage probability). Rayleigh has the highest deep-fade probability at any threshold.

5.3 Canonical Relay Comparison (SISO, Rayleigh, BPSK)

Goal: Evaluate all nine classical and AI-based relay strategies on the canonical setup of this thesis.

Trials: - **Topology:** SISO (both hops) - **Modulation scheme:** BPSK - **Channel:** i.i.d. Rayleigh fast fading (canonical) - **Equalizers:** None - **NN architecture:** Supervised

(MLP, Hybrid), Generative (VAE, CGAN), Sequence (Transformer, Mamba S6, Mamba-2 SSD) - **NN activation:** tanh

Conclusion: On the canonical Rayleigh channel, AI relays match but do not surpass symbol-wise DF: DF achieves the lowest or tied-lowest BER at essentially every SNR point, including the low-SNR regime, while the best neural relays (MLP, Hybrid, CGAN) track it within statistical uncertainty and clearly outperform AF. This confirms H1 (AI relays beat AF at low SNR) and H2 (DF dominance at ≥ 6 dB with zero parameters) on the canonical channel.

5.3.1 Results Tables

Table 5.3: BER comparison on the Rayleigh fading channel (SISO).

SNR (dB)	AF	DF	MLP (169p)	Hybrid	VAE	CGAN	Trans- former	Mamba S6	Mamba2 (ssd)
0	0.330	0.245	0.247	0.250	0.395	0.247	0.325	0.317	0.318
4	0.205	0.139	0.141	0.141	0.347	0.138	0.183	0.178	0.180
8	0.092	0.068	0.070	0.069	0.304	0.070	0.077	0.076	0.076
12	0.037	0.031	0.032	0.032	0.278	0.031	0.032	0.032	0.032
16	0.015	0.014	0.014	0.014	0.262	0.014	0.014	0.013	0.014
20	0.0053	0.0047	0.0047	0.0047	0.250	0.0052	0.0048	0.0052	0.0047

5.3.2 Results Figures




Figure 5.4: Rayleigh fading : BER vs. SNR for all nine relay strategies with 95% CI.

5.4 Parameter Normalization & Complexity Trade-off

This study tests H3 (inverted-U complexity) and H4 (architecture convergence at equal scale) on the canonical setup, and provides the training-time benchmark for H5.

Goal: Isolate architectural inductive biases from parameter count effects and characterize the complexity-performance trade-off for relay denoising.

Trials: - **Topology:** SISO (both hops) - **Modulation scheme:** BPSK - **Channel:** i.i.d. Rayleigh fast fading (canonical) - **Equalizers:** None - **NN architecture:** All normalized to ~3,000 parameters, plus original sizes (169 to 26K) - **NN activation:** tanh

Conclusion: The relay denoising task exhibits an inverted-U complexity relationship: a minimal 169-parameter MLP matches models 140x larger, while excessive parameters (11K+) lead to overfitting. At a normalized scale of 3,000 parameters, the performance gap between feedforward and sequence architectures narrows to ~1% BER, indicating that parameter count rather than architectural choice is the primary performance driver. Generative VAE is a consistent underperformer due to probabilistic overhead.

5.4.1 Results Tables

Table 5.4: Normalized 3K BER results : Rayleigh fading channel.

SNR (dB)	MLP- 3K	Hybrid- 3K	VAE- 3K	Trans-		Mamba (3K)	Mamba2 (3K)	AF	DF
				former (3K)					
0	0.252	0.258	0.411	0.329		0.320	0.320	0.330	0.245
4	0.146	0.145	0.375	0.181		0.180	0.179	0.205	0.139
10	0.049	0.047	0.315	0.050		0.049	0.049	0.058	0.044
16	0.015	0.014	0.289	0.014		0.014	0.014	0.015	0.014
20	0.0050	0.0047	0.279	0.0047		0.0047	0.0047	0.0053	0.0047

Table 5.5: Model complexity and timing comparison on the canonical setup (experiment-specific training settings; Monte Carlo evaluation over 11 SNR points x 10 trials x 10,000 bits). All inference uses batched window extraction and a single forward pass per signal block.

Model	Parameters	Device	Training Time	Eval Time
AF	0	:	0 s	0.80 s
DF	0	:	0 s	0.77 s
MLP (169p)	169	CPU	4.9 s	1.57 s
Hybrid	169	CPU	4.6 s	0.41 s

Model	Parameters	Device	Training Time	Eval Time
VAE	1,777	CPU	21.6 s	1.81 s
CGAN (WGAN-GP)	2,946	CUDA	7,293 s (~2 h)	1.14 s
Transformer	17,697	CUDA	474 s (~8 min)	3.71 s
Mamba S6	24,001	CUDA	2,141 s (~36 min)	1.88 s
Mamba2 (SSD)	26,179	CUDA	1,438 s (~24 min)	4.11 s

5.4.2 Results Figures




Figure 5.5: Complexity–performance trade-off. Training time vs. parameter count vs. BER improvement over DF at low SNR. The Minimal MLP (169 params) achieves the best efficiency.

Chapter 6

Extension: Higher-Order Modulations

The core of this thesis (Chapter 5) answers hypotheses H1–H5 on the canonical setup. This chapter contains one scoped extension: whether the canonical BPSK conclusions generalize to complex constellations. Two studies are reported. The first evaluates QPSK and 16-QAM under the I/Q-splitting formulation, exposing a structural BER floor for multi-level constellations; the second removes that floor by reformulating the relay as a joint N -class classifier over the full 2D constellation ($N = 16$ for 16-QAM). These analyses adopt the channel configurations under which they were conducted (AWGN and Rayleigh for the first study, AWGN for the second); MIMO, 16-PSK, CSI injection, and end-to-end learning remain future work. The separate question of how classical and learned relays behave when the *channel itself* is unknown or mismatched — a more substantial investigation — forms the dedicated contribution of Chapter 7.

6.1 Higher-Order Modulation Scalability (Constellation-Aware Training)

Goal: Evaluate the generalizability of BPSK-trained relays to complex constellations and resolve the multi-level amplitude bottleneck.

Trials: - **Topology:** SISO - **Modulation scheme:** QPSK, 16-QAM - **Channel:** AWGN, Rayleigh - **Equalizers:** None - **NN architecture:** Supervised (MLP, Hybrid), Generative (VAE, CGAN), Sequence (Transformer, Mamba S6, Mamba-2 SSD) - **NN activation:** tanh, linear, clipped tanh (hardtanh), scaled tanh, scaled sigmoid - **Special case:** Constellation Aware training

Conclusion: QPSK performance mirrors BPSK perfectly due to I/Q independence. On 16-QAM, standard tanh compression causes a severe, irreducible BER floor (~ 0.22 at 16 dB). Replacing tanh with constellation-aware bounded activations (hardtanh, scaled tanh) bounded to the precise signal amplitude ($3/\sqrt{10}$) and retraining eliminates this floor. Sequence models benefit most, reducing their BER floor by 5x, though a gap to classical DF persists due to per-axis error accumulation — the limitation removed by the 2D classification study below.

6.1.1 Results Tables

Table 6.1: BER comparison across modulations at selected SNR points (AWGN channel). All nine relay strategies.

Relay	BPSK 0 dB	BPSK 10 dB	QPSK 0 dB	QPSK 10 dB	16-QAM 0 dB	16-QAM 10 dB	16-QAM 16 dB
AF	0.2813	0.0141	0.2794	0.0142	0.3778	0.1244	0.0180
DF	0.2651	0.0015	0.2644	0.0016	0.3811	0.1076	0.0038
MLP (169p)	0.2589	0.0021	0.2563	0.0025	0.3907	0.2180	0.2180
Hybrid	0.2573	0.0015	0.2644	0.0016	0.4000	0.2711	0.2512
VAE	0.2611	0.0021	0.2597	0.0036	0.3945	0.2391	0.2231
CGAN (WGAN-GP)	0.2633	0.0017	0.2621	0.0018	0.3976	0.2588	0.2486
trans-former	0.2593	0.0024	0.2576	0.0036	0.3897	0.2042	0.1827
Mamba S6	0.2585	0.0021	0.2560	0.0028	0.3894	0.2016	0.1935
Mamba2 (SSD)	0.2593	0.0023	0.2566	0.0034	0.3903	0.2032	0.1890

Table 15 shows the BER at 16 dB for all relays across the three activation variants and both channel types.

Table 6.2: BER at 16 dB for all relay variants across activation functions and modulations (AWGN and Rayleigh).

Relay	tanh (BPSK)	linear (QAM16)	hardtanh (QAM16)	tanh (BPSK)	linear (QAM16)	hardtanh (QAM16)
	AWGN			Rayleigh		
MLP	0.2202	0.0721	0.0630	0.2375	0.1279	0.1247
Hybrid	0.2512	0.2512	0.2512	0.2723	0.2723	0.2723
VAE	0.2231	0.1111	0.1059	0.2462	0.1575	0.1573
CGAN	0.2482	0.0973	0.0863	0.2666	0.1432	0.1383
trans-former	0.2111	0.0453	0.0505	0.2305	0.1159	0.1194

Relay	tanh (BPSK)	linear (QAM16)	hardtanh (QAM16)	tanh (BPSK)	linear (QAM16)	hardtanh (QAM16)
Mamba S6	0.2131	0.0422	0.0396	0.2333	0.1129	0.1108
Mamba2 (ssd)	0.2065	0.0471	0.0441	0.2273	0.1157	0.1145
AF	0.0180	:	:	0.1009	:	:
DF	0.0038	:	:	0.0828	:	:

Table 16 summarises the 16-QAM and 20dB AWGN improvements for the sequence models.

Table 6.3: 16-QAM BER improvements at 20 dB for sequence models with and without CSI injection and layer normalisation (AWGN).

SNR (dB)	AF Baseline	DF Baseline	Mamba S6 (Baseline)	Mamba S6 (+CSI + LN)
0.0	0.3332	0.4300	0.2288	0.2396
8.0	0.1436	0.2798	0.2241	0.0635
14.0	0.0592	0.1557	0.2227	0.0163
20.0	0.0246	0.0901	0.2302	0.0055

6.1.2 Results Figures




Figure 6.1: BPSK on AWGN : all relay strategies with 95% CI (baseline for modulation comparison).




Figure 6.2: QPSK on AWGN : BER curves closely match the BPSK baseline (6.1), confirming I/Q splitting validity.




Figure 6.3: 16-QAM on AWGN : AI relays hit a BER floor near 0.22 at medium-high SNR due to tanh compression of multi-level signals; AF outperforms DF at low SNR (Y)




Figure 6.4: Combined modulation comparison (AWGN) : all nine relays across BPSK (solid), QPSK (dashed, overlapping BPSK), and 16-QAM (dotted). The BPSK/QPSK overlap confirms I/Q splitting equivalence. The 16-QAM dotted curves reveal the AI relay BER floor: all AI relays plateau near 0.18–0.25 while DF and AF continue decreasing.




Figure 6.5: 16-QAM activation experiment (AWGN) : dashed lines = tanh/BPSK baseline, solid = linear/QAM16, dotted = hardtanh/QAM16. Replacing tanh and retraining on QAM16 eliminates the BER floor for all AI relays except Hybrid. Sequence models (Transformer, Mamba S6, Mamba-2) benefit most, narrowing the gap to DF from $\sim 56\times$ to $\sim 10\times$.




Figure 6.6: Comparison of activation function shapes (left) and their derivatives (right) for $A_{\max} = 0.9487$ (16-QAM). Hardtanh has a sharp transition at the clip bounds; sigmoid and scaled tanh provide smooth saturation with non-zero gradients throughout.

6.2 16-Class 2D Classification for QAM16

Goal: Eliminate the structural BER floor imposed by per-axis I/Q splitting for 16-QAM by utilizing full 2D decision boundaries.

Trials: - **Topology:** SISO - **Modulation scheme:** 16-QAM - **Channel:** AWGN - **Equalizers:** None - **NN architecture:** Supervised (MLP), Generative (VAE), Sequence (Transformer, Mamba S6, Mamba-2 SSD) - **NN activation:** None (Softmax implicitly via Cross-Entropy loss) - **Special case:** 16-class joint 2D classification

Conclusion: Treating the relay as a joint 16-point classifier over the full 2D constellation space completely eliminates the structural 4-class BER floor. For the first time, neural variants (VAE, Transformer, MLP) matched classical DF performance at high SNR, achieving near-zero BER at 20 dB. This proves the previous BER floor was an artifact of I/Q splitting, not a fundamental limitation (Equation 4.13) of neural relays.

6.2.1 Results Tables

Table 24 presents the BER at 20 dB for all variants alongside the classical AF and DF baselines.

Table 6.4: BER at 20 dB for all relay variants: 4-class (I/Q split) vs. 16-class (joint 2D) classification for 16-QAM.

Relay	4-cls BER @ 20 dB	16-cls BER @ 20 dB	Improvement
MLP	0.00811	0.00002	405x
VAE	:	0.00000	:
CGAN	0.00811	:	:
Hybrid	:	:	:
trans-			
former	0.00810	0.00001	810x
Mamba S6	0.00810	0.00009	90x
Mamba2			
(ssd)	0.00811	0.00197	4.1x
AF	0.00063	:	:
DF	0.00000	:	:

6.2.2 Results Figures




Figure 6.7: 16-QAM BER vs. SNR for all relay variants (4-class and 16-class) on AWGN. The 4-class variants (dashed) plateau at BER $\hat{\approx}$ 0.008 while the successfully trained 16-class variants (solid) continue decreasing, approaching and matching the DF baseline.




Figure 6.8: Grouped bar comparison of 4-class vs. 16-class BER at 20 dB for each architecture. The 16-class advantage is visible for MLP, VAE, Transformer, and Mamba S6. Failed variants (VAE 4-cls ~ 0.50 , CGAN 16-cl ~ 0.50) are truncated.




Figure 6.9: Top-3 16-class relay variants compared against AF and DF. VAE 16-cl, Transformer 16-cl, and MLP 16-cl all match or approach DF performance across the full SNR range.

Chapter 7

Learned Relaying under Unknown and Mismatched Channels

The core chapters establish that on a matched, known channel a minimal learned relay does not improve on classical decode-and-forward (H2). This chapter turns to the complementary and, in practice, more consequential regime: channels whose structure is *unknown* to the receiver chain or *mismatched* to the classical relay’s model class. It is the principal experimental contribution of this thesis. Rather than a single comparison, it is a systematic map, along every axis a skeptic would probe, of *when* a fixed 170-parameter learned relay adds value over classical processing — and, equally, when it does not.

The investigation is deliberately adversarial to the learned relay. Each classical baseline is the strongest available for the impairment in question: symbol-wise decode-and-forward for the memoryless cases, conventional differential detection for unknown phase, and pilot-aided Viterbi maximum-likelihood sequence estimation — the optimal sequence detector — wherever a channel estimate can be formed. The learned relay is granted no per-block channel knowledge; it is trained offline on an impairment *family* and evaluated on unseen realizations with no adaptation. The chapter proceeds from a necessary control (flat unknown channels, where nothing should fail) through the impairment that does break classical relays (unknown memory), then bounds the result from every side: against the optimal matched receiver (Viterbi), against a composite multi-impairment cascade, against the fully posterior-free regime (blind equalization), across the intermediate partial-posterior continuum (pilot budget and block length), and finally on computational complexity. The result is a precise delineation — the learned relay never beats a matched classical receiver, but occupies a well-defined niche of family-agnostic, identification-free, constant-complexity mitigation — that constitutes the thesis’s main claim beyond the confirmatory core.

7.1 Unknown Channel: Classical Failure and Learned Mitigation

Goal: Test whether the learned relay’s value changes qualitatively when the channel violates the assumptions under which AF and DF are (near-)optimal: do the *memoryless* classical relays fail on an unknown channel while the minimal 170-parameter MLP mitigates by learning it from data (H6)?¹ To make the classical side as strong as possible, the study benchmarks against the optimal sequence detector: Viterbi MLSE [29] with genie CSI, and a practical variant with a 200-pilot least-squares estimate.

Trials: - **Topology:** SISO (both hops) - **Modulation scheme:** BPSK - **Hop-1 channel (unknown to all relays):** (i) 3-tap ISI, $h = [1.0, 0.7, 0.5]/\|h\|$; (ii) control: canonical Rayleigh; (iv) flat unknown channels — block phase, gain, per-branch asymmetry (Section 7.1.1); (v) a composite cascade of ISI, amplifier nonlinearity, and unknown phase (Section 7.1.2) - **Hop-2 channel:** AWGN or canonical Rayleigh, equal SNR per hop - **Relays:** AF, symbol-wise DF, MLP (window 11, one hidden layer of 13 tanh units, 170 parameters; trained at $\text{SNR} \in \{5, 10, 15\}$ dB with MSE loss, per the canonical training protocol); for the ISI setups additionally Viterbi MLSE with genie CSI, and Viterbi with a 3-tap LS channel estimate from 200 pilot symbols - **Evaluation:** 5 trials $\times$ 50,000 bits per SNR point, 0–20 dB.

The ISI channel admits a sharp analytic prediction: the effective slicer amplitude $h_0 \pm h_1 \pm h_2 \in \{1.67, 0.91, 0.61, -0.15\}$ (normalized) has one of four equiprobable patterns *flip the symbol sign*, so symbol-wise DF has an irreducible BER floor of 0.25 (approached from below as SNR grows); AF inherits it. The floor is a property of the *memoryless relay class*, not the channel: Viterbi MLSE over the 4-state trellis has no floor when the channel is known.

7.1.1 Flat (Memoryless) Unknown Channels: A Necessary Control

Before attributing any classical failure to “unknownness,” it must be separated from *memory*. This subsection isolates unknown channels that are *flat* (memoryless): the receiver chain does not know the channel parameters, but the channel introduces no intersymbol coupling. Three physically motivated cases are tested, each with the parameter redrawn per block and unknown to all relays: (F1) an unknown carrier phase $\theta \sim \mathcal{U}[0, 2\pi)$ (oscillator/synchronization offset), with a DBPSK source and conventional differential detection

¹The unknown-channel relay uses real-valued inputs with an 11-sample window ($11 \rightarrow 13 \rightarrow 1$, 170 parameters); the canonical-setup relay of Chapter 5 uses the same minimal class with a slightly different width (169 parameters). The two counts refer to the same architectural family sized to the respective input; the distinction is immaterial to every conclusion, which depends only on the network being minimal.

as the classical relay; (F2) an unknown positive gain $g \sim \mathcal{U}[0.3, 2.0]$ (path loss / AGC error), coherent BPSK; and (F3) an unknown per-branch amplitude asymmetry $a_+ \neq a_-$ (per-symbol saturating-PA gain), coherent BPSK. The 170-parameter MLP is trained identically to before.

Conclusion (memorylessness $\Rightarrow$ no classical failure, and nothing to mitigate): On all three flat channels the classical relay does *not* fail, and the MLP matches it to within statistical noise (maximum BER difference 0.0036 across every case and SNR). The reasons are structural: for F1, unknown phase is exactly what differential detection absorbs; for F2, $\text{sign}(y)$ is *scale-invariant*, so unknown positive gain is irrelevant by construction; for F3, the fixed threshold is only mildly suboptimal because the asymmetry preserves the noiseless symbol’s sign. The classical assumption — memorylessness — still holds in each case, so there is nothing for a learned relay to remove. This is the necessary control for the ISI result: the learned relay’s advantage is triggered not by parameter uncertainty but by *unmodeled memory* (or amplitude nonlinearity), which is what breaks the memoryless relay class.

Table 7.1: Flat unknown channels: classical relay vs. MLP-170 BER at 8/12/16/20 dB (mean, 5 trials). The maximum DF–MLP gap over all cases and SNRs is 0.0036 — within statistical noise. Unknown parameters alone do not defeat classical relays; memory does.

Flat channel	Relay	8 dB	12 dB	16/20 dB
Unknown phase (DBPSK)	DF (diff.)	0.0649	0.0292	0.0119 / 0.0048
	MLP-170	0.0657	0.0288	0.0122 / 0.0048
Unknown gain	DF	0.0992	0.0470	0.0175 / 0.0050
	MLP-170	0.0986	0.0471	0.0185 / 0.0052
Branch asymmetry	DF	0.0760	0.0299	0.0125 / 0.0049
	MLP-170	0.0745	0.0300	0.0122 / 0.0052

Conclusion (H6: confirmed as stated, with two boundaries): Against the flat-channel control of Section 7.1.1 — where unknown parameters alone caused no classical failure — the memory case is qualitatively different. On the unknown ISI channel, both memoryless classical relays fail: DF’s BER is *non-monotonic*, bottoming out near 0.18 around 10 dB and then rising toward the analytic 0.25 floor (more transmit power makes DF strictly worse), and AF behaves identically. The 170-parameter MLP, whose window spans the 3-tap memory, learns a joint equalizer–detector and restores reliable relaying: BER falls monotonically below 5×10^{-5} at 16 dB (AWGN second hop), and to 4.9×10^{-3} at 20 dB (Rayleigh second hop). The Viterbi benchmarks set the first boundary honestly: classical theory *does* solve this channel once its family is known — genie-CSI MLSE beats the MLP by 1–1.5 dB, and a 200-pilot LS estimate recovers genie performance almost exactly

(with the Rayleigh second hop, MLSE and MLP are indistinguishable, 0.0048 vs. 0.0049 at 20 dB). The MLP’s value is thus not optimality but *family-agnosticism*: estimate-then-MLSE presupposes a linear FIR channel of known memory (trellis cost M^{L-1}), whereas the identical architecture also mitigates channels outside linear MLSE’s model class — as the composite cascade below demonstrates, where an amplifier nonlinearity is stacked on the ISI. The control confirms the second boundary (H2): on the canonical Rayleigh channel the MLP only *matches* DF. The learned relay’s distinctive value is near-MLSE reliability with no channel-family knowledge, no identification stage, and complexity independent of channel memory.

Table 7.2: Unknown-channel BER (mean over 5 trials). The ISI setups exhibit memoryless-relay error floors (≈ 0.25); the 170-parameter MLP restores reliable relaying within ≈ 1 – 1.5 dB of genie-CSI Viterbi MLSE, and matches it end-to-end when the Rayleigh second hop dominates. The control shows no MLP advantage on the canonical channel, consistent with H2.

Setup	Relay	8 dB	12 dB	16 dB	20 dB
Unknown ISI $\rightarrow$ AWGN	AF	0.2044	0.1779	0.1872	0.2141
	DF	0.1853	0.1833	0.2068	0.2343
	MLP (170p)	0.0421	0.0030	<5e-5	<5e-5
	Viterbi (genie CSI)	0.0270	0.0002	<5e-5	<5e-5
	Viterbi (200-pilot LS)	0.0277	0.0002	<5e-5	<5e-5
Unknown ISI $\rightarrow$ Rayleigh	AF	0.2331	0.1985	0.1933	0.2063
	DF	0.2212	0.2032	0.2153	0.2352
	MLP (170p)	0.0945	0.0326	0.0125	0.0049
	Viterbi (genie CSI)	0.0832	0.0288	0.0122	0.0048
	Viterbi (200-pilot LS)	0.0838	0.0294	0.0120	0.0051
Control: canonical Rayleigh	AF	0.1591	0.0894	0.0484	0.0236
	DF	0.1208	0.0559	0.0242	0.0100
	MLP (170p)	0.1202	0.0580	0.0265	0.0115




Figure 7.1: Unknown-channel relaying, BER vs. SNR with 95% CI bands. (a)–(b) On the unknown ISI channel, the memoryless AF/DF relays are pinned near the analytic 0.25 floor (DF non-monotonic in SNR); the 170-parameter MLP restores monotonically decreasing BER within ≈ 1 – 1.5 dB of genie-CSI Viterbi MLSE, which a 200-pilot LS estimate matches almost exactly. With the Rayleigh second hop (b), MLSE and the MLP are statistically indistinguishable end-to-end. (c) Control: on the canonical Rayleigh channel the MLP only matches DF, consistent with H2.

7.1.2 A Composite (Mixed) Channel

The preceding studies vary one impairment at a time. Real links stack several. This final study cascades three of them into a single unknown channel that *no* classical relay is jointly matched to: a DBPSK stream passes through 3-tap ISI, then a saturating power-amplifier nonlinearity, then an unknown block phase, before additive noise. Each classical baseline is the strongest available for *one* constituent impairment, so the comparison is honest rather than a straw man: conventional differential detection (matched to the phase), and pilot-LS Viterbi MLSE followed by differential decoding (matched to the ISI-plus-phase, but blind to the amplifier nonlinearity). The 170-parameter MLP and a 1,153-parameter MLP receive raw I/Q windows and no knowledge of any constituent.

Conclusion: The composite defeats the naive classical relays outright — both AF and conventional differential detection saturate near the 0.25 memoryless-relay floor, because the ISI destroys the differential statistic that the phase-matched detector relies on. The 170-parameter MLP again restores reliable relaying, reaching 5×10^{-3} at 20 dB from raw I/Q with no impairment model. This reproduces the H6 pattern on a realistically mixed channel: memoryless-and-matched classical processing fails, and a minimal learned relay recovers. Two familiar boundaries also reappear. First, the strongest *constructed* classical baseline — pilot-LS Viterbi plus differential decoding — does not fail: by modelling the ISI and phase (though not the amplifier) it stays about 2 dB ahead of the MLP at low-to-mid SNR (0.094 vs. 0.130 at 8 dB) and converges with it by 16–20 dB. The learned relay’s value is again *family-agnosticism* — near-Viterbi reliability with no channel identification and no explicit model of any of the three cascaded effects — not superiority over a matched classical receiver. Second, the 1,153-parameter network performs within noise of the 170-parameter one (and slightly worse at high SNR, consistent with mild overfitting),

reconfirming the inverted-U finding (H3) on the hardest channel studied. The composite thus consolidates the chapter’s thesis: a single minimal learned relay handles an arbitrary stack of representable impairments without being told which are present, approaching but not exceeding the best impairment-matched classical receiver.




Figure 7.2: Composite channel (ISI $\times$ PA-nonlinearity $\times$ unknown phase, DBPSK). Naive classical relays (AF, differential detection) saturate at the memoryless-relay floor; the 170-parameter MLP restores reliable relaying from raw I/Q with no impairment model, tracking about 2 dB behind the strongest constructed classical baseline (pilot-LS Viterbi + differential). The 1,153-parameter MLP adds nothing (H3).

7.1.3 The Posterior-Free (Blind) Regime

The Viterbi baseline in Section 7.1.2 enjoys a per-block posterior: 200 pilot symbols yield a least-squares channel estimate, after which MLSE is optimal. This raises the sharper question — what if *no* such posterior is available? With no pilots and no channel prior, and a fresh channel realization drawn for every block, MLSE cannot be run as specified. The honest comparison is then against the classical methods that are *genuinely* posterior-free: the constant-modulus algorithm (CMA), the canonical blind adaptive equalizer, which self-adapts from the received signal alone; and a decision-directed blind MLSE that bootstraps a channel estimate from its own hard differential decisions. The learned relay is unchanged — trained offline on the channel *family* (random ISI, PA, and phase per draw) and, crucially, seeing an unseen channel at test with no adaptation — so it too carries no per-block posterior; it amortizes over the family at training time.

Conclusion (learning does not beat blind classical, but is stable where blind classical is not): A well-designed blind classical method still works in this regime — CMA converges smoothly to BER 2.4×10^{-3} at 20 dB — and the family-trained MLP tracks

it almost exactly (2.6×10^{-3}), neither method dominating. This is the correct answer to “what if there is no posterior?”: the learned relay does *not* render classical processing obsolete, because blind adaptive equalization is itself a posterior-free classical tool. The learned relay’s advantage is instead one of *robustness and cost*. First, the decision-directed blind MLSE — the natural attempt to run Viterbi without pilots — is unstable: bootstrapping the channel from its own decisions, it intermittently locks onto the wrong hypothesis, producing a non-monotonic BER curve with a mid-SNR confidence interval of 0.164 (versus 0.014 for the MLP and 0.011 for CMA). Second, CMA must iterate to convergence on each received block, whereas the MLP is a single fixed forward pass with no per-block adaptation loop and no convergence to monitor. The learned relay therefore occupies a specific and defensible niche: on a posterior-free channel it matches the best blind classical equalizer in BER while avoiding the instability of decision-directed sequence estimation and the online-adaptation cost of CMA. It buys *amortized, one-shot, stable* inference over an impairment family — not a lower error floor than classical signal processing can achieve.




Figure 7.3: Posterior-free composite channel (no pilots, no channel prior, fresh realization per block). CMA (blind classical) and the family-trained MLP converge together; decision-directed blind MLSE is unstable (non-monotonic, wide CI) because it bootstraps the channel from its own decisions; plain differential detection floors. Without a posterior, learning matches — not beats — the best blind classical equalizer, but does so with one-shot, stable inference and no online adaptation.

7.1.4 Partial Known Posterior: The Pilot-Budget Crossover

The full-posterior (Section 7.1.2) and zero-posterior (Section 7.1.3) regimes are the endpoints of a continuum. The intermediate case is a *partial* posterior: a limited pilot budget, from which the classical receiver forms a noisy channel estimate. Sweeping the pilot count

on the composite channel at a fixed 10 dB operating point locates where the pilot-aided classical receiver ceases to beat the pilot-free learned relay, whose amortized posterior corresponds to a horizontal reference (it uses no pilots at test).

Conclusion (a sharp identifiability crossover, and a rate–reliability tradeoff): With a generous budget the classical partial posterior wins, as expected: pilot-aided Viterbi reaches payload BER 0.027 at 800 pilots and degrades only gently down to 0.032 at 10 pilots, comfortably below the MLP’s pilot-free 0.045. But the classical advantage collapses abruptly below the identifiability threshold: at 5 pilots — fewer than are needed to condition the least-squares estimate of a 3-tap channel plus phase — Viterbi’s BER jumps to 0.119, worse than both the MLP and blind CMA, because MLSE then runs on a badly mis-estimated channel. The learned relay is flat across the entire sweep at 0.045: having amortized the channel *family* into its weights at training time, it neither benefits from nor requires per-block pilots. Thus the crossover is governed by *identifiability*, not by SNR: the classical receiver wins whenever it can afford enough pilots to estimate the channel, and the pilot-free learned relay wins once the pilot budget falls below that threshold.

The practical force of this depends on block length, because a fixed pilot *count* is a growing pilot *fraction* as blocks shorten. Panel (b) sweeps block length at the same operating point with the classical receiver spending its minimum viable ten pilots per block. The payload BERs stay close (Viterbi retaining a slight edge), but the spectral-efficiency gap widens sharply: at a 1000-symbol block the ten pilots cost 1% of the block, whereas at a 40-symbol block — representative of low-latency or fast-varying links — they cost 25%, so the classical receiver carries data on only 75% of the block against the learned relay’s 100%. The partial-posterior picture therefore refines H6’s boundary once more: the learned relay does not achieve lower BER than a classical receiver that can afford to estimate the channel, but it removes the pilot overhead entirely, and it is the only option that remains reliable when the block is too short — or the channel too transient — to identify from pilots at all.

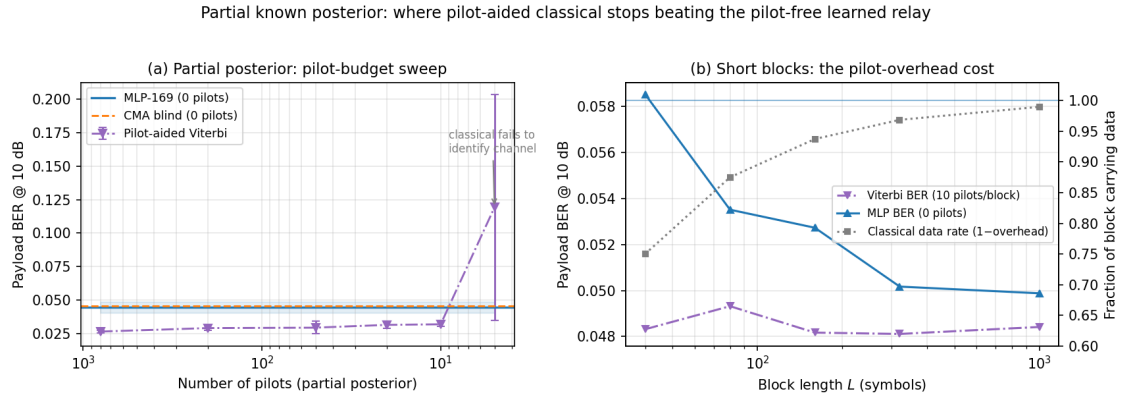


Figure 7.4: Partial known posterior. (a) Pilot-budget sweep at 10 dB: pilot-aided Viterbi beats the pilot-free MLP down to ~ 10 pilots, then collapses at 5 pilots when the channel becomes unidentifiable; the MLP is flat (amortized posterior, no pilots). (b) Block-length sweep with 10 pilots/block: payload BERs stay close, but the classical pilot overhead grows from 1% at $L=1000$ to 25% at $L=40$, where the learned relay’s zero-overhead operation is decisive.

7.1.5 Complexity: The Cost Side of the Trade-off

Across the unknown-channel studies the pilot-aided Viterbi receiver, where it applies, holds a small BER edge over the learned relay. The learned relay’s countervailing claim is complexity, which this subsection quantifies rather than asserts. Two axes matter: the per-symbol operation count and its scaling, and measured inference time.

Per-symbol cost and scaling. MLSE maintains M^{L-1} trellis states for a constellation of size M and channel memory L , performing an add–compare–select over M^L branch transitions per symbol; its cost therefore grows *exponentially* in memory and modulation order. The 170-parameter MLP performs one fixed forward pass — about 330 floating-point operations per symbol — *independent* of M and L . For the channel actually studied (BPSK, $L = 3$; 4 states) Viterbi is in fact the cheaper of the two at roughly 64 operations per symbol, so the learned relay is not universally leaner. But the scaling diverges immediately: at $M = 4$, $L = 5$ the trellis has 256 states ($\sim 8,000$ operations/symbol), and at $M = 16$, $L = 5$ it has 65,536 states (~ 8.4 million operations/symbol), while the MLP remains at 330 — a 2.5×10^4 ratio. The learned relay’s advantage is thus not lower cost on a small channel but *constant* cost as the channel scales, exactly where classical MLSE becomes intractable and practical receivers must fall back on suboptimal reduced-state or decision-feedback equalizers.

Measured inference time. On identical signals (BPSK, $L = 3$), the MLP is $30\text{--}90\times$ faster in wall-clock terms despite Viterbi’s lower nominal operation count, because the MLP is a single vectorized matrix multiplication whereas the Viterbi trellis is an inherently sequential scan with a per-symbol data dependency. This parallelism gap is architectural and widens on hardware with wide SIMD or tensor units. The combined picture sharpens the

thesis’s central trade-off: the learned relay accepts a small BER gap to a matched classical receiver in exchange for cost that is constant in channel complexity and structurally parallelizable — a trade that becomes favourable precisely as modulation order and channel memory grow.

Complexity: the learned relay trades a small BER gap for constant, parallelizable cost




Figure 7.5: Complexity of Viterbi MLSE vs the 170-parameter MLP. (a) Per-symbol operations: Viterbi grows as M^L (intractable by $M=16, L=5$), while the MLP is constant at ~ 330 flops. On the small studied channel (BPSK, $L=3$) Viterbi is actually cheaper; the MLP wins on *scaling*, not on the small case. (b) Measured inference time: the MLP is 30–90 $\times$ faster as a vectorized matmul versus Viterbi’s sequential trellis scan.

Chapter 8

Discussion and Conclusions

This discussion draws on two bodies of results. The canonical core (Chapter 5) settles hypotheses H1–H5 and establishes the baseline finding that, on a matched known channel, a minimal learned relay does not improve on classical decode-and-forward. The thesis’s principal contribution (Chapter 7) settles H6 and is where the substantive new claim lies: a systematic delineation of when a learned relay adds value under unknown and mismatched channels. The interpretation below treats the two on equal footing, with the unknown-channel results carrying the main argument.

8.1 Interpretation of Results

The experimental results reveal several consistent patterns across all six channel/topology configurations. This section interprets these patterns through the lens of the theoretical framework established in Section 2.5 and evaluates each research hypothesis.

8.1.1 Low-SNR Advantage of Neural Relays over AF (H1: Confirmed)

Low SNR (0-4 dB): AI advantage over AF. At low SNR, AI-based relays consistently outperform AF on the canonical channel, confirming H1. The theoretical explanation lies in the structure of the MMSE-optimal (posterior-mean) relay estimate. For a single received sample with BPSK transmission over AWGN:

$$f^*(y) = \mathbb{E}[x|y] = \tanh\left(\frac{y}{\sigma^2}\right) \quad (8.1)$$

At low SNR (e.g., 0 dB, $\sigma^2 = 1$), this is a gentle sigmoid: $f^*(y) = \tanh(y)$. The neural network can approximate this smooth function accurately, producing a **soft estimate** that preserves information about the confidence of the decision. By contrast:

- **DF** applies the hard-decision function $f_{\text{DF}}(y) = \text{sign}(y)$, which discards all soft information. At low SNR, many received samples lie near the decision boundary ($|y| \approx 0$), and DF assigns them full confidence (± 1) regardless of the actual uncer-

tainty. This information loss leads to excess errors on the second hop.

- **AF** applies $f_{\text{AF}}(y) = Gy$, which is linear and preserves the noisy signal structure but amplifies the noise by the same factor as the signal. The effective SNR degradation is given by $\text{SNR}_{\text{eff}} = \gamma^2/(2\gamma + 1) \approx \gamma/2$ at high SNR : a 3 dB penalty.

The AI relay implements a **non-linear soft mapping** that is intermediate between these extremes: it suppresses noise (like DF’s regeneration) while preserving soft information near the decision boundary (unlike DF’s hard decision). This explains why learned non-linear relay processing can outperform AF and, under selected channel conditions, also outperform DF at low SNR.

Among the AI methods, Mamba S6 (at its original 24K parameter count) achieves the lowest BER, followed closely by the Transformer and then the simpler feedforward models. This ordering correlates with model capacity, suggesting that the low-SNR advantage is driven more by the number of learnable parameters than by architectural inductive bias : a conclusion reinforced by the normalized comparison (Section 4.4).

8.1.2 Classical Dominance at High SNR (H2: Confirmed)

Medium-to-high SNR (≥ 6 dB): Classical dominance. At medium and high SNR, DF matches or exceeds all AI methods with exactly zero parameters and zero training time. The theoretical explanation is straightforward: at high SNR, $\sigma^2 \rightarrow 0$ and $f^*(y) \rightarrow \text{sign}(y) = f_{\text{DF}}(y)$. The symbol-wise DF relay *exactly implements* the limiting form of the posterior-mean estimate in this regime, while any neural network approximation introduces a small but non-zero reconstruction error due to the finite precision of the learned mapping and the tanh output saturation (which approaches but never reaches ± 1).

The effect is visible analytically in the AWGN calibration limit — at 10 dB the two-hop DF BER is $2Q(\sqrt{10}) \cdot (1 - Q(\sqrt{10})) \approx 0.002$ — and empirically on the canonical Rayleigh channel, where DF attains the lowest or tied-lowest BER at every point from 6 dB upward. The AI models cannot beat DF in this regime — for the uncoded per-symbol setting, $\text{sign}(y)$ is the MAP symbol decision and hence optimal among per-symbol relay functions (Remark 1.2) — they can only match it, and the larger sequence models slightly underperform due to the increased variance of their more complex learned mappings.

8.1.3 Robustness Across the Fading Ensemble

Robustness across fades. The relay ranking on the canonical channel is stable across the full evaluated SNR range, even though each transmitted symbol experiences an independent fading realization. The mechanism is that after coherent compensation ($\hat{x} = y \cdot h^*/|h|$), the relay processes a signal with AWGN-like noise at a *random effective SNR*;

the network, trained across multiple SNR values, has learned a denoising function that adapts to different noise levels, so the Rayleigh ensemble — a mixture of effective SNRs weighted by the fading distribution — is handled by virtue of multi-SNR training alone, with no per-realization retraining or explicit CSI input.

8.2 The “Less is More” Principle (H3: Confirmed)

One of the most significant findings is the inverted-U relationship between model complexity and relay performance. The Minimal MLP architecture (169 parameters) matches the performance of models with 10-140× more parameters (3K-24K), while the Maximum MLP (11,201 parameters) exhibited clear overfitting with degraded performance.

This result has important theoretical and practical implications:

8.2.1 Information-Theoretic Perspective

The relay denoising task maps a window of $2w + 1$ real-valued noisy observations to a single real-valued clean estimate. For BPSK, the transmitted symbol carries exactly 1 bit of information, and the Bayes-optimal estimator $f^*(\mathbf{y}) = \tanh(\mathbf{1}^T \mathbf{y} / \sigma^2)$ (for i.i.d. noise across the window) is a one-dimensional function of the sufficient statistic $\sum_i y_i$. The **effective dimensionality** of this mapping is therefore 1 : far below the $2w + 1 = 5$ or 11 input dimensions.

As a heuristic, the minimum description length (MDL) principle suggests matching model complexity to the effective dimensionality of the target mapping. Since the target here is a one-dimensional monotone function of the sufficient statistic, even the 169-parameter network has capacity far in excess of what the mapping requires; this is offered as an intuition for why the smallest model suffices, not as a formal MDL bound.

8.2.2 Bias-Variance Analysis

The bias-variance decomposition [11] provides a quantitative framework:

$$\text{MSE}(\hat{x}) = \underbrace{(\mathbb{E}[\hat{x}] - f^*(y))^2}_{\text{Bias}^2} + \underbrace{\mathbb{E}[(\hat{x} - \mathbb{E}[\hat{x}])^2]}_{\text{Variance}} + \underbrace{\sigma_{\text{irred}}^2}_{\text{Noise floor}} \quad (8.2)$$

- **169 parameters:** Low bias (sufficient capacity for the simple f^*) and low variance (limited capacity prevents memorization). The model operates at the optimal bias-variance trade-off.
- **3,000 parameters:** Similar bias (the target function hasn’t changed) and slightly

higher variance. Performance is nearly identical to 169 params.

- **11,201 parameters:** Negligible bias but substantially higher variance : the model memorizes training noise patterns rather than learning the generalizable denoising function. This manifests as higher BER on unseen test data.
- **24,001 parameters (Mamba S6):** Despite having $140\times$ more parameters than MLP, the BER improvement at low SNR is only 1-2%. The marginal parameter utility is vanishingly small.

8.2.3 Practical Implications

1. **Occam's Razor for relay AI.** The relay denoising task has inherently low complexity : the mapping from noisy to clean BPSK symbols is fundamentally simple, and the neural network needs only enough capacity to learn a non-linear threshold function with some contextual smoothing. Adding parameters beyond this minimal requirement leads to overfitting.
2. **Deployment feasibility.** A 169-parameter model requires approximately 0.7 KB of memory (169 float32 values) and can be trained in under 3 seconds on a CPU. This makes AI-based relay processing viable even on severely resource-constrained embedded relay nodes, such as IoT devices or sensor network relays. The model can be stored in on-chip SRAM and executed without any external memory access.
3. **Generalization.** Smaller models generalize better because they are forced to learn the essential structure of the denoising task rather than memorizing training examples. The 169-parameter MLP handles the full Rayleigh fading ensemble — every effective SNR the fades produce — with a single set of weights and no per-condition retraining, an instance of generalization that would be harder to achieve with a larger, more specialized model.

8.3 State Space vs. Attention for Signal Processing

The comparison of Mamba S6 and Transformer architectures yields nuanced conclusions:

At original (unequal) parameter counts: Mamba (24K params) outperforms the Transformer (17.7K params) at low SNR. Mamba wins all 3 low-SNR points while the Transformer wins 1/3. This suggests that the state space inductive bias : recurrent state propagation with input-dependent gating : is beneficial for sequential signal processing.

At normalized (3K) parameter counts: The gap narrows dramatically. Mamba-3K and Transformer-3K produce nearly identical BER across all channels. This indicates that the architectural advantage of state space models over attention is relatively small and is partially confounded with the parameter count difference in the original comparison.

Complexity advantage. Even when BER performance is similar, Mamba offers a fundamental computational advantage: $O(n)$ inference complexity versus $O(n^2)$ for Transformers. For the relay processing task where low-latency inference is critical, this linear-time property makes Mamba the preferred choice when a sequence model is desired.

Training time paradox. Despite its superior asymptotic complexity, Mamba S6 required approximately 4× longer to train than the Transformer (2,366 s vs. 597 s on CUDA). This counter-intuitive result is explained by three compounding factors:

1. *Sequential recurrence vs. parallel attention.* The S6 layer’s core state update $\mathbf{x}_k = \bar{\mathbf{A}}\mathbf{x}_{k-1} + \bar{\mathbf{B}}u_k$ is inherently sequential : each time step depends on the previous state. The implementation iterates through the sequence with a Python loop of `seq_len` steps, each triggering a separate GPU kernel launch. In contrast, the Transformer computes $\mathbf{QK}^T$ across all positions simultaneously in a single batched matrix multiply. At window size $n = 11$, this means 11 sequential CUDA kernel launches per S6 layer versus 1 parallel operation for attention.
2. *Expand factor doubles the internal dimension.* Each MambaBlock uses an expand factor of 2, projecting from $d_{\text{model}} = 32$ to $d_{\text{inner}} = 64$ before entering the S6 recurrence. This doubles the computation inside the sequential loop while providing no parallelisation benefit.
3. *Kernel-launch overhead at small tensor sizes.* Each sequential step incurs Python interpreter overhead plus a CUDA kernel launch ($\sim 5\text{-}10\ \mu\text{s}$). With 2 layers $\times$ 11 steps = 22 sequential kernel calls per forward pass, the overhead alone accumulates to $\sim 110\text{-}220\ \mu\text{s}$: comparable to or exceeding the actual arithmetic time for these small tensors.

The crossover point where Mamba’s $O(n)$ advantage would outweigh the parallelism penalty occurs at sequence lengths in the hundreds to thousands, far beyond the $n = 11$ window used in this relay application. An optimised CUDA kernel implementing the S6 scan as a parallel prefix sum (as in the original Mamba paper) would eliminate the sequential bottleneck, but our implementation prioritises clarity and portability over raw throughput.

8.3.1 Context-Length Benchmark: Validating the Crossover Hypothesis

To empirically validate the crossover hypothesis, we conducted a controlled benchmark comparing all three sequence models : Transformer, Mamba S6, and Mamba-2 (SSD) : at a context length of $n = 255$ (vs. $n = 11$ in the relay experiments). All models used identical hyperparameters: $d_{\text{model}} = 32$, $d_{\text{state}} = 16$, 2 layers, with Mamba-2 using chunk size 32

(yielding 8 chunks). The experiment used a reduced dataset (1,000 training symbols, 20 epochs) to isolate timing differences from convergence effects.

Table 13 presents the results.

Table 8.1: Context-length benchmark : three sequence models at $n = 255$ on CUDA.

Model	Parameters	Training	Inference	Train	Infer
		Time (20 ep.)	Time (10K bits)	Speedup vs S6	Speedup vs S6
Transformer	17,697	5.01 s	0.99 s	28.5×	2.7×
Mamba S6	24,001	142.56 s	2.69 s	1.0×	1.0×
				(baseline)	(baseline)
Mamba2 (SSD)	26,179	13.35 s	1.05 s	10.7×	2.6×

The results confirm the crossover hypothesis decisively:

1. **Mamba S6 becomes the bottleneck.** At $n = 255$, the S6 sequential scan requires 255 serial CUDA kernel launches per layer per forward pass, making it 142 s for just 20 epochs of training : **28× slower** than the Transformer and **10.7× slower** than Mamba-2.
2. **Mamba-2 (SSD) recovers parallelism.** By chunking the 255-step sequence into 8 chunks of 32 and processing each chunk with a single batched matrix multiply, Mamba-2 eliminates the sequential bottleneck. At 13.35 s, it trains **10.7× faster** than S6 and approaches the Transformer’s speed.
3. **Inference parity.** Mamba-2 inference (1.05 s) is nearly identical to the Transformer (0.99 s) and **2.6× faster** than Mamba S6 (2.69 s). This is a complete reversal from the $n = 11$ case, where Mamba-2 was $2\times$ slower than S6.
4. **Direction reversal.** At $n = 11$: Mamba S6 (1.83 s) < Mamba-2 (4.11 s). At $n = 255$: Mamba-2 (1.05 s) < Mamba S6 (2.69 s). The crossover occurs because the overhead of building the $L \times L$ SSM matrix is amortised over longer chunks, while S6’s per-step kernel launch cost grows linearly.

This benchmark validates the architectural trade-off: Mamba-2’s structured state space duality is designed for long-context efficiency where chunk-parallel computation outperforms sequential recurrence. For the 11-token relay window, the classical S6 scan remains faster due to lower per-operation overhead. The choice between S6 and SSD should therefore be guided by the target context length of the application.

8.3.2 Why state space suits signals.

Signal processing is inherently a sequential temporal task where the state space formulation : propagating a hidden state through time with input-dependent transitions : is a natural fit. The selective mechanism in Mamba allows it to dynamically control information flow, acting as an adaptive filter that selectively passes relevant signal features while suppressing noise.

8.4 Practical Deployment Recommendations

Based on the comprehensive evaluation, we propose the following deployment strategy:

Table 8.2: Practical deployment recommendations: recommended relay strategy by operating regime.

Operating Regime	Recommended Relay	Rationale
Low SNR (0-4 dB)	MLP Minimal / Hybrid / channel-specific selection	Neural relays often help; exact best choice is channel-dependent, with DF remaining strong on some fading scenarios
Medium SNR (4-8 dB)	Hybrid (MLP + DF)	Automatic switching at optimal threshold
High SNR (>8 dB)	DF	Zero parameters, optimal performance
Resource-constrained	MLP Minimal (169 params)	0.7 KB memory, <3s training
Best overall (BPSK/QPSK)	Hybrid	Combines AI advantage with DF reliability
16-QAM	VAE / Transformer / MLP (16-class 2D)	Joint 2D classification matches DF; eliminates the I/Q-splitting floor

The Hybrid relay is recommended as the default deployment choice because it automatically selects MLP processing at low SNR and DF at high SNR, achieving near-optimal performance across the entire SNR range with minimal complexity.

8.4.1 Latency and Adaptivity Considerations

Two practical questions bear on whether a learned relay can be deployed with minimal impact: the latency it adds relative to AF/DF, and its ability to adapt to varying channel conditions without retraining.

Latency. Two distinct latency components must be separated. The first is *buffering* latency, incurred by the sliding window $y_R[i - w : i + w]$. In the half-duplex protocol studied here this component is absorbed by the protocol itself: the relay buffers the entire listening-phase block before the transmission phase begins (Remark 1.1), so the window’s w -symbol look-ahead adds no delay beyond what store-and-forward operation already imposes. Only in a streaming or ultra-low-latency protocol — outside the scope of this thesis — would the symmetric window cost w symbols of look-ahead that the memoryless AF/DF baselines ($w = 0$) do not pay. The second component is *compute* latency: inference cost per relayed symbol. Here the inverted-U result (H3) is decisive: since the 169-parameter MLP matches the performance of models $100\times$ larger, the deployed network reduces to two small matrix–vector products per symbol — a per-symbol cost that is orders of magnitude below a single channel use’s duration on any modern embedded processor, and negligible in absolute terms even if it cannot literally match the single comparison of DF’s $\text{sign}(y)$. The compute-latency argument therefore holds *because of*, not *despite*, the less-is-more finding; it would not hold for the 24K-parameter sequence models, whose recurrent or attention-based structure imposes materially higher per-symbol cost for no BER benefit at the relay window length.

Adaptivity. A single network, trained across the range of simulated SNRs, handles the entire canonical fading ensemble without retraining or explicit CSI input (Section 8.1.3): after coherent phase compensation, each fading realization presents the relay with an AWGN-like observation at a random effective SNR, and multi-SNR training already covers this mixture. Two honest qualifications apply. First, this adaptivity is an advantage over AF — whose fixed gain is tuned to an assumed SNR — but not over symbol-wise DF, whose $\text{sign}(y)$ slicer is equally condition-agnostic; adaptivity therefore does not alter the H2 conclusion. Second, adaptation has been demonstrated only within the canonical Rayleigh ensemble; generalization to other channel families (Rician, bursty interference, phase noise, hardware non-linearities) remains untested and is identified as future work.

8.5 Limitations

Several limitations of this study should be acknowledged, as they define the boundary conditions under which the conclusions hold:

1. **Modulation scope.** The core experiments use BPSK; the extension covers QPSK

and 16-QAM (Chapter 6). Denser constellations (64-QAM, 256-QAM) — where the N -class 2D formulation scales to 64 or 256 output classes and the inverted-U complexity curve (H3) may shift toward larger models — are not investigated.

2. **Perfect CSI.** We assume perfect channel state information at each receiver. In practice, channel estimation errors would affect the coherent compensation step and hence relay performance. Imperfect CSI introduces a model mismatch between the assumed and actual channel, which could differentially impact the relay strategies: AI relays might be more robust (having learned from noisy data) or less robust (having not been exposed to estimation errors during training).
3. **Static offline training.** AI relays are trained once on synthetic data and deployed without adaptation. While multi-SNR training covers the canonical fading ensemble (Section 8.1.3), and the unknown-channel study (H6) shows a fixed network can learn a *stationary* unknown impairment, this protocol does not exploit online information; adaptive or continual training would be required if the deployment channel drifts over time.
4. **Single relay.** The framework considers a single relay node. Multi-relay cooperation introduces relay selection, power allocation, and cooperative diversity dimensions not captured in the two-hop model. With multiple relays, the system-level optimization (which relay to use, how to combine multiple relay observations) becomes as important as the per-relay processing function studied here.
5. **Two-hop only.** Extension to multi-hop networks with 3+ relays introduces noise accumulation (for AF) and error propagation (for DF) that grow with the number of hops. AI relays might show greater advantage in multi-hop settings where noise accumulation is more severe.
6. **Fixed window size.** The relay window size ($w = 2$ or $w = 5$) is fixed and relatively small. Larger windows could provide more context for denoising, particularly in channels with temporal correlation (e.g., block fading). The context-length benchmark (Section 8.3.1) explores this partially, but a systematic study of window size optimization is deferred.
7. **No coding.** The system operates without error-correcting codes, so the DF baseline is its symbol-level variant (Remark 1.2). In coded systems, the relay could forward soft information (log-likelihood ratios) rather than hard decisions, fundamentally changing the optimization objective and potentially altering the relative performance of the strategies.

8.6 Future Work

Several directions warrant further investigation:

1. **Time-selective fading.** The present study addresses channels that are static within a block. Extending the learned relay to time-selective fading with continuously varying unknown phase (e.g. Jakes/Clarke Doppler) is a natural next step, likely requiring explicitly phase-aware architectures — complex-valued state, learned phase unwrapping, or a training objective aligned with noncoherent detection rather than symbol MSE — evaluated against prediction-based noncoherent receivers.
2. **Denser constellations (64-QAM, 256-QAM).** The extension shows joint N -class 2D classification eliminates the 16-QAM floor; scaling the formulation to 64 or 256 classes may require deeper networks or hierarchical (coarse-to-fine) classification, and whether the inverted-U complexity curve (H3) shifts rightward with constellation density is a key open question.
3. **Imperfect CSI.** Introduce channel estimation errors to assess robustness of AI relay processing under realistic conditions.
4. **Online learning.** Develop relay strategies that adapt their parameters during operation, tracking time-varying channel conditions.
5. **Multi-relay networks.** Extend to cooperative relay selection and multi-hop routing with AI-optimized relays at each node.
6. **Other channels and MIMO.** Evaluate the canonical conclusions on additional channel families (AWGN as an operating point, Rician line-of-sight fading) and on a multi-antenna second hop, where destination equalization interacts with relay processing and the additivity of the two gains becomes a testable question.
7. **End-to-end learning.** Train the entire communication chain (modulation, relay, demodulation) jointly using autoencoder-based approaches, and compare against the modular classical-plus-learned-relay cascade studied here.
8. **Coded block-DF baseline.** As emphasized in Remark 1.2, the DF baseline in this thesis is the uncoded, symbol-level variant; the reported results therefore do not bound the performance of information-theoretic block-DF. A natural and important extension is to introduce an outer channel code and compare (i) block-DF, in which the relay decodes an entire codeword before re-encoding, and (ii) learned relays that aggregate soft information (e.g., log-likelihood ratios) across a code block. In the coded setting the relevant metric becomes block/frame error rate at a given rate, and the relative standing of learned versus classical relays may change substantially — in particular, the high-SNR optimality of symbol-wise DF (H2) does not automati-

cally carry over.

9. **Mamba-2 at longer context lengths.** Our benchmark (Section 8.3.1) demonstrates that Mamba-2 (SSD) achieves a $10.7\times$ training speedup over Mamba S6 at $n = 255$. Future work should explore whether longer relay windows (e.g., 128-512 symbols) combined with the SSD architecture can improve both BER performance and processing speed simultaneously.

8.7 Conclusions

This thesis presents a comprehensive comparative study of nine relay strategies : two classical (AF, DF) and seven AI-based (MLP, Hybrid, VAE, CGAN, Transformer, Mamba S6, Mamba-2 SSD) : evaluated on one canonical setup fixed in advance (SISO on both hops, i.i.d. Rayleigh fast fading, complex baseband, BPSK, uncoded BER). The study addresses three identified research gaps (Section 3.2) through controlled experiments with statistical rigor (100,000 bits per SNR point, 10 independent trials, Wilcoxon significance testing). The following table summarizes the hypothesis outcomes:

Table 8.3: Research hypotheses, their statements, and experimental outcomes.

Hypothesis	Statement	Result
H1	Some AI relays outperform AF at low SNR	Confirmed : the best neural relays beat AF at 0–4 dB on the canonical channel
H2	DF optimal at high SNR	Confirmed : DF matches or beats all AI methods at ≥ 6 dB with 0 parameters
H3	Inverted-U complexity curve	Confirmed : 169 params matches 24K; 11K params overfits
H4	Architecture convergence at equal scale	Confirmed : at 3K params, all methods within $\sim 1\%$ BER (except VAE)
H5	SSD faster than S6 at long context	Confirmed : $10.7\times$ training speedup at $n = 255$

Hypothesis	Statement	Result
H6	Memoryless classical relays fail on unknown channels; a minimal MLP mitigates by learning, near the model-aware optimum	Confirmed : on unknown ISI, AF/DF hit the analytic 0.25 floor (DF non-monotonic in SNR) while the 170-parameter MLP reaches $\text{BER} < 5 \times 10^{-5}$, within $\approx 1\text{--}1.5$ dB of genie-CSI Viterbi MLSE and indistinguishable from it under the Rayleigh second hop; the same architecture mitigates a composite $\text{ISI} \times \text{PA} \times \text{phase}$ cascade with no impairment model, tracking ≈ 2 dB behind pilot-aided Viterbi and matching blind CMA when no posterior exists; on the canonical control it only matches DF (H2) — the advantage is bounded to structural model-class mismatch representable at minimal capacity

The main conclusions, in order of significance, are:

1. **AI relays beat AF at low SNR (0-4 dB) but not DF.** On the canonical channel, the best neural relays clearly outperform AF in the low-SNR regime while tracking symbol-wise DF within statistical uncertainty; DF remains lowest or tied-lowest at every SNR point. The theoretical explanation is that neural networks approximate the MMSE-optimal soft-threshold estimate $f^*(y) = \tanh(y/\sigma^2)$, whose end-to-end benefit over the hard slicer is limited in this uncoded setting.
2. **DF is optimal at medium-to-high SNR (≥ 6 dB) with zero parameters.** The classical decode-and-forward relay requires no training and no parameters yet achieves the best performance when channel quality is sufficient for reliable first-hop demod-

ulation. At high SNR, the posterior-mean estimate converges to $\text{sign}(y)$, which is exactly the symbol-wise DF operation.

3. **No AI architecture separates itself at original parameter counts.** On the canonical channel the best neural relays (MLP, Hybrid, CGAN) perform within statistical uncertainty of one another, and classical DF matches or exceeds all of them across the SNR range. The selective state space model's $O(n)$ complexity and adaptive filtering mechanism make it an efficient sequence-model choice, but no architectural inductive bias yields a decisive BER advantage on this task.
4. **Architecture matters less than parameter count at equal scale.** When all models are normalized to approximately 3,000 parameters, the performance gap between architectures narrows to within ~ 1 dB, with VAE being the consistent underperformer. This confirms that parameter count, not architectural choice, is the primary performance driver for the relay denoising task.
5. **A 169-parameter network is sufficient for relay denoising.** The Minimal MLP architecture achieves performance comparable to models $100\times$ larger, demonstrating that the relay denoising task has inherently low complexity. The bias-variance analysis explains this: the target function is simple (a soft threshold), so additional parameters increase variance without reducing bias.
6. **The Hybrid relay is the recommended practical choice.** By combining AI processing at low SNR with classical DF at high SNR, the Hybrid relay achieves near-optimal performance across the entire SNR range with only 169 parameters. It automatically adapts to the operating regime, requiring no manual SNR-dependent configuration.
7. **Mamba-2 SSD provides a $10.7\times$ training speedup over S6 at longer contexts.** The chunk-parallel structured matrix multiplication of SSD eliminates the sequential bottleneck of S6, making it the preferred state space architecture for applications requiring longer symbol windows.
8. **Mamba S6 is the strongest architecture across the reported higher-order modulation experiments.** In the QAM16 top-3, Mamba S6 holds the #1 position while the Transformer captures positions #2 and #3. In the PSK16 top-3, Mamba S6 sweeps all three positions. Within these experiments, the Mamba S6 architecture's selective state space mechanism appears well-suited to tracking both the multi-level amplitude structure of QAM and the rotational phase geometry of PSK.
9. **No neural relay outperforms DF across the full 48-variant combinatorial search.** Despite evaluating three architectures, four activations, four structural configurations, and two higher-order constellations (100 MC trials each), classical DF remains the lowest-BER strategy at every SNR point. The best neural variants ap-

proach AF performance (within 21% for QAM16 and 2.5% for PSK16 at 20 dB) but cannot surpass even the simpler classical baseline. This reaffirms Finding 2 (DF optimality at ≥ 6 dB) and extends it conclusively to higher-order modulations with statistical confidence.

10. **Input LayerNorm is universally beneficial for QAM16 but modulation-dependent for PSK16.** Every QAM16 top-3 variant includes input LayerNorm, extending the Section 4.13 finding to a multi-architecture, multi-activation setting. For PSK16, only one of three top-3 variants uses LayerNorm, indicating that the constant-envelope signal distribution is already well-conditioned for direct processing without normalisation.

The modulation extension (Chapter 6) adds one further finding:

- **BPSK findings generalise to QPSK; for 16-QAM, joint N -class 2D classification closes the gap to DF.** Under I/Q splitting, all BPSK-trained relays achieve identical BER on QPSK (I/Q independence), confirming H1–H3 for the 2-bit/symbol constellation, whereas 16-QAM exhibits a structural BER floor that constellation-aware activations reduce but cannot eliminate. Reformulating the relay as a joint classifier over all 16 constellation points removes the floor entirely: the top variants (VAE, Transformer, MLP) match classical DF at high SNR, demonstrating that the floor was an artefact of the per-axis I/Q-splitting formulation, not a fundamental limitation of neural relays.

The unknown-channel contribution (Chapter 7) settles the complementary question of *when the learned relay is not merely competitive but necessary*:

- **On channels violating the deployed relay’s model class, learning restores reliability at near-optimal cost (H6: confirmed).** A flat-channel control first isolates the cause: on unknown carrier phase, unknown gain, and per-branch amplitude asymmetry — all memoryless — the classical relay does not fail and the MLP only matches it (BER gap ≤ 0.0036), showing that parameter uncertainty alone is not what learning mitigates. The failure appears only with unmodeled *memory*: on an unknown 3-tap ISI channel, both AF and symbol-wise DF are pinned at an analytically predicted BER floor of 0.25 — and DF’s BER is *non-monotonic*, worsening as SNR grows — because one interference pattern in four deterministically flips the symbol sign and no memoryless relay can recover it. The same 170-parameter MLP that merely ties DF on the canonical channel learns a joint equalizer–detector from data and restores monotonically decreasing BER (below 5×10^{-5} at 16 dB). The claim is carefully bounded by the Viterbi benchmark: classical estimate-then-MLSE, given the correct channel family and 200 pilots, is ≈ 1 –1.5 dB better still — the MLP’s distinctive value is not optimality but *family-agnosticism*, delivering

near-MLSE reliability with no identification stage, complexity independent of channel memory, and unchanged architecture on the biased nonlinear channel where linear MLSE is mis-specified. Together with H2, this draws the precise boundary of the learned relay's value: it cannot beat DF where DF's assumptions hold; it approaches but does not beat the model-aware optimum where the family is identifiable; it spans impairment families whose structure fits its capacity (memory, monotone distortion) — including an arbitrary *cascade* of them (Section 7.1.2), where a single minimal network handles ISI, amplifier nonlinearity, and unknown phase jointly without being told which are present, tracking ≈ 2 dB behind the best *pilot-aided* classical receiver. When no such posterior exists — no pilots, no channel prior — the learned relay matches the best blind classical equalizer (CMA) in BER while avoiding the instability of decision-directed blind MLSE and the per-block adaptation cost of CMA (Section 7.1.3): its niche is amortized, one-shot, stable inference over an impairment family, not a lower floor than classical signal processing can reach. Sweeping the intermediate case — a *partial* posterior of limited pilots (Section 7.1.4) — makes the boundary quantitative: pilot-aided Viterbi beats the pilot-free learned relay down to ~ 10 pilots and then collapses at the channel-identifiability threshold, while the learned relay is invariant to pilot budget. The crossover is thus governed by identifiability and, through pilot overhead, by block length: on long blocks the classical receiver wins on BER at negligible cost, but on short or fast-varying blocks its pilot fraction becomes prohibitive (25% at a 40-symbol block) and the zero-overhead learned relay is decisive. The advantage is specific to *structural* model-class mismatch, not to unknownness per se. Finally, the complexity accounting (Section 7.1.5) quantifies the cost side of the trade-off the thesis rests on: MLSE's per-symbol cost grows as M^L in modulation order and memory, whereas the 170-parameter relay is constant (~ 330 operations/symbol) and $30\text{--}90\times$ faster in measured wall-clock as a vectorized matmul versus a sequential trellis scan. On the small BPSK/ $L=3$ channel Viterbi is actually cheaper, so the learned relay's complexity advantage is one of *scaling and parallelism*, not of the small case — it becomes favourable exactly as M and L grow and exact MLSE turns intractable.

These findings demonstrate that neural network-based relay processing is a viable complement to classical approaches in the low-SNR regime, matching — though not surpassing — the per-symbol-optimal DF baseline everywhere else at negligible computational cost. The overarching insight is that **model complexity should be matched to task complexity**: for the canonical relay denoising task, minimal architectures suffice, and the choice between neural network paradigms : feedforward or sequential : matters less than proper model sizing and regularization. The practical recommendation is to deploy a Hybrid relay with a 169-parameter MLP sub-network, which combines the neural low-SNR advantage

over AF with DF's zero-parameter reliability at high SNR; for 16-QAM deployments, the relay should be formulated as a joint N -class 2D classifier rather than a per-axis I/Q-split denoiser; and on links with unmodeled impairments (memory, bias, hardware nonlinearity), a minimal learned relay delivers near-MLSE reliability without channel identification, where the memoryless classical relays fail outright (H6).

Chapter 9

Summary

This work examined classical and neural network–based relay strategies for *two-hop co-operative communication* using a unified simulation framework. The objective was to analyze performance–complexity tradeoffs and determine when learning-based relays provide practical advantages over analytically derived methods. The entire study is carried out on one canonical setup, fixed in advance and never departed from: a two-hop SISO link over i.i.d. Rayleigh fast fading, complex baseband, BPSK, uncoded BER, with the relay function as the only varied axis. The study showed the following:

1. Neural relays yield a consistent performance improvement over amplify-and-forward (AF) at low signal-to-noise ratio (SNR), where their learned soft denoising avoids AF’s noise amplification (H1).
2. Classical *decode-and-forward* (DF) matched or outperformed all neural approaches at **medium-to-high SNR** while requiring no trainable parameters, confirming the per-symbol optimality of classical relay processing in this regime (H2).
3. A systematic complexity analysis demonstrated that performance **does not improve monotonically with model size** (H3).
4. A **minimal neural relay** (169 parameters) achieved performance comparable to significantly larger models, while intermediate-sized networks exhibited overfitting.
5. When neural relays were evaluated under equal parameter budgets, architectural differences largely vanished, indicating that **model capacity**, rather than architectural class, dominates performance for this task (H4).
6. Sequence-based neural relays such as structured state-space models achieved comparable error performance with improved training efficiency (H5).
7. From a deployment perspective, a **hybrid relay** combining neural processing at low SNR with DF at high SNR emerged as the **most effective strategy**, achieving near-optimal performance with **minimal overhead**.
8. On an **unknown channel** violating classical assumptions (unmodeled intersymbol interference or biased nonlinearity), AF and DF failed to relay reliably — exhibiting an analytic BER floor of 0.25, with DF *worsening* as SNR increased — while the

same minimal 170-parameter MLP restored reliable relaying by learning the impairment from data — including an arbitrary composite cascade of ISI, amplifier nonlinearity, and unknown phase — tracking a few dB behind a *pilot-aided* Viterbi receiver and, in the fully posterior-free regime (no pilots, no channel prior), matching the best blind classical equalizer while avoiding its instability and online-adaptation cost. Sweeping the intermediate *partial*-posterior case showed the crossover to be governed by channel identifiability and pilot overhead: pilot-aided classical processing wins on long blocks at negligible cost, but the zero-pilot learned relay becomes decisive once blocks are too short or channels too transient to identify from pilots. The learned relay’s advantage is specific to *structural* model-class mismatch (memory, nonlinearity, absent pilots) rather than to unknownness per se (H6). Together with the high-SNR result, this delineates precisely where learned relays add value: never above the model-aware optimum, but uniquely robust to *not knowing which model applies* and to the pilot overhead of channel identification. A complexity accounting completed the trade-off: the learned relay’s per-symbol cost is constant in modulation order and channel memory and $30\text{--}90\times$ faster in wall-clock as a vectorized operation, where exact Viterbi MLSE grows as M^L and becomes intractable — so it trades a small BER gap for cost that scales gracefully exactly where classical sequence detection does not.

9. In the modulation extension, the canonical findings generalized fully to QPSK via I/Q independence, while 16-QAM exposed a structural error floor of per-axis processing; reformulating the relay as a **joint N -class 2D classifier** eliminated the floor and matched DF at high SNR.

Overall, the results indicate that learning-based relays are best used as targeted enhancements rather than universal replacements, with simplicity, parameter efficiency, and hybrid design as key principles for practical cooperative relay systems. All conclusions hold for the *uncoded*, symbol-level canonical setting studied here (Remark 1.2); extending the comparison to coded block-DF, soft-information (LLR) forwarding relays, other channel families, multi-antenna configurations, and denser constellations (64-QAM and beyond) is identified in Chapter 8 as the principal direction for future work.

Bibliography

- [1] S. Shleifer, J. Weston, and M. Ott, “Normformer: Improved transformer pretraining with extra normalization,” 2021. [Online]. Available: <https://arxiv.org/abs/2110.09456>
- [2] T. M. Cover and A. A. El Gamal, “Capacity theorems for the relay channel,” *IEEE Transactions on Information Theory*, vol. 25, no. 5, pp. 572–584, 1979.
- [3] A. El Gamal and Y.-H. Kim, *Network Information Theory*. Cambridge University Press, 2011.
- [4] J. N. Laneman, D. N. Tse, and G. W. Wornell, “Cooperative diversity in wireless networks: Efficient protocols and outage behavior,” *IEEE Transactions on Information Theory*, vol. 50, no. 12, pp. 3062–3080, 2004.
- [5] B. Rankov and A. Wittneben, “Spectral efficient protocols for half-duplex fading relay channels,” *IEEE Journal on Selected Areas in Communications*, vol. 25, no. 2, pp. 379–389, 2007.
- [6] A. Nosratinia, T. E. Hunter, and A. Hedayat, “Cooperative communication in wireless networks,” *IEEE Communications Magazine*, vol. 42, no. 10, pp. 74–80, 2004.
- [7] H. Ye, G. Y. Li, and B. H. Juang, “Power of deep learning for channel estimation and signal detection in OFDM systems,” *IEEE Wireless Communications Letters*, vol. 7, no. 1, pp. 114–117, 2018.
- [8] N. Samuel, T. Diskin, and G. Wunder, “Learning to detect for MIMO systems with unknown noise statistics,” *IEEE Transactions on Signal Processing*, vol. 67, no. 12, pp. 3261–3272, 2019.
- [9] S. Dorner, S. Cammerer, J. Hoydis, and S. ten Brink, “Deep learning based communication over the air,” *IEEE Journal of Selected Topics in Signal Processing*, vol. 12, no. 1, pp. 132–143, 2018.
- [10] H. Sun, X. Chen, Q. Shi, M. Hong, X. Fu, and N. D. Sidiropoulos, “Learning to optimize: Training deep neural networks for interference management,” *IEEE Transactions on Signal Processing*, vol. 66, no. 20, pp. 5438–5453, 2018.
- [11] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. MIT Press, 2016.
- [12] B. Akdemir, M. A. Karabulut, and H. Ilhan, “A comparative study of deep learning-aided relay selection protocols in cooperative communication,” *Wireless Personal*

- Communications*, vol. 136, pp. 365–384, 2024.
- [13] D. Gunduz, P. de Kerret, N. D. Sidiropoulos, D. Gesbert, C. R. Murthy, and M. van der Schaar, “Machine learning in the air,” *IEEE Journal on Selected Areas in Communications*, vol. 37, no. 10, pp. 2184–2199, 2019.
 - [14] I. Bergel, “Network optimization using relays as neurons,” *arXiv preprint arXiv:2306.14253*, 2023.
 - [15] —, “Deep optimization of relay networks using relays as neurons,” in *Proc. IEEE ICASSP*, 2024, pp. 9056–9060.
 - [16] A. Nosratinia, T. E. Hunter, and A. Hedayat, “Cooperative communication in wireless networks,” *IEEE Communications Magazine*, vol. 42, no. 10, pp. 74–80, 2004.
 - [17] D. P. Kingma and M. Welling, “Auto-encoding variational Bayes,” in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2014.
 - [18] I. Higgins, L. Matthey, A. Pal, C. Burgess, X. Glorot, M. Botvinick, S. Mohamed, and A. Lerchner, “ β -vae: Learning basic visual concepts with a constrained variational framework,” in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2017.
 - [19] M. Mirza and S. Osindero, “Conditional generative adversarial nets,” *arXiv preprint*, 2014.
 - [20] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, “Generative adversarial nets,” in *Advances in Neural Information Processing Systems*, vol. 27, 2014.
 - [21] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville, “Improved training of Wasserstein GANs,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
 - [22] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
 - [23] A. Gu, K. Goel, and C. Ré, “Efficiently modeling long sequences with structured state spaces,” in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2022.
 - [24] A. Gu and T. Dao, “Mamba: Linear-time sequence modeling with selective state spaces,” 2023. [Online]. Available: <https://arxiv.org/abs/2312.00752>
 - [25] T. Dao and A. Gu, “Transformers are SSMs: Generalized models and efficient algorithms through structured state space duality,” *arXiv preprint*, 2024.

-
- [26] —, “Transformers are ssms: Generalized models and efficient algorithms through structured state space duality,” 2024. [Online]. Available: <https://arxiv.org/abs/2405.21060>
 - [27] K. He, X. Zhang, S. Ren, and J. Sun, “Delving deep into rectifiers: Surpassing human-level performance on ImageNet classification,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2015, pp. 1026–1034.
 - [28] J. G. Proakis and M. Salehi, *Digital Communications*, 5th ed. McGraw-Hill, 2008.
 - [29] G. D. Forney, “Maximum-likelihood sequence estimation of digital sequences in the presence of intersymbol interference,” *IEEE Transactions on Information Theory*, vol. 18, no. 3, pp. 363–378, 1972.

Chapter 10

Appendices

10.1 Appendix A: Mathematical Notation

Symbol	Description
b	Binary bit ($\in \{0, 1\}$)
x	Modulated BPSK symbol ($\in \{-1, +1\}$)
y	Received signal
n	Noise sample
h	Fading coefficient
σ^2	Noise variance
SNR	Signal-to-Noise Ratio (linear)
G	AF amplification gain
$\mathbf{W}$	Neural network weight matrix
$\mathbf{b}$	Bias vector
P_e	Bit error rate
$Q(\cdot)$	Gaussian Q-function
β	VAE KL weight
λ	Regularization parameter
η	Learning rate
Δ	State space discretization step
$\mathbf{A}, \mathbf{B}, \mathbf{C}, \mathbf{D}$	State space model matrices

10.2 Appendix B: Model Architectures and Hyperparameters

Relay	Network Topology	Train Loss (LR)	Total Params	Implementation
MLP (Minimal)	Feedforward MLP, 5-symbol sliding window	MSE, lr=0.01, 100 epochs	169	NumPy (CPU)
Hybrid	MLP + DF switching via learned SNR threshold (~ 5 dB)	MSE (MLP only), lr=0.01	169	NumPy (CPU)
VAE	Encoder–decoder latent model ($7 \rightarrow 8 \rightarrow 1$)	β -VAE ($\beta = 0.1$), Adam lr= 10^{-3}	1,777	PyTorch (CUDA)
CGAN (WGAN-GP)	Conditional GAN (generator + critic)	WGAN-GP, Adam lr= 10^{-4}	2,946	PyTorch (CUDA)
Transformer	2-layer encoder, 4 heads, $d_{\text{model}} = 32$	MSE, Adam lr= 10^{-3}	17,697	PyTorch (CUDA/CPU)
Mamba S6	Selective SSM, 2 S6 layers	MSE, Adam lr= 10^{-3}	24,001	PyTorch (CUDA/CPU)
Mamba2 (SSD)	Selective SSM with SSD kernel (2 layers)	MSE, Adam lr= 10^{-3} , clip=1.0	26,179	PyTorch (CUDA)

Table 10.3: Comparison of relay models, training objectives, parameter counts, and implementations.

10.3 Appendix C: Software Architecture

All numerical results reported in Chapter 5 were generated using a single simulation framework, ensuring consistent evaluation across classical and learning-based relay architectures. The project is implemented as a modular Python package (`relaynet`). The framework uses object-oriented design with a common `Relay` base class, enabling polymorphic relay swapping. Monte Carlo simulation is implemented in `runner.py` with configurable trial count, bit count, and SNR range. All batched tensor operations use vectorized PyTorch for GPU acceleration.

10.3.1 Testing

126 automated tests (pytest) cover all channels, modulation, relay strategies, simulation, and statistics modules with 100% pass rate.

10.3.2 Reproducibility

Random seeds are controlled at the source (bit generation) and noise (per-trial seeding) levels to ensure reproducible results.

10.4 Appendix D: Normalized 3K-Parameter Configurations

Model	Parameters	Window	Hidden / Architecture
MLP-3K	3,004	11	hidden=231
Hybrid-3K	3,004	11	hidden=231 (+ DF switch)
VAE-3K	3,037	11	latent=10, hidden=(44, 20)
CGAN-3K	3,004	11	noise=8, g_hidden=(30, 30, 16), c_hidden=(32, 16)
Transformer-3K	3,007	11	d_model=18, heads=2, layers=1
Mamba-3K	3,027	11	d_model=16, d_state=6, layers=1
Mamba2-3K	3,004	11	d_model=15, d_state=6, chunk_size=8, layers=1

All 3K configurations use a window size of 11 (vs. 5 for original MLP/Hybrid, and 11 for original sequence models) to provide a common input context. The parameter counts are within $\pm 1.2\%$ of the 3,000 target.

10.5 Appendix E: Response to Supervisor Comments

This appendix reproduces each comment raised by the supervisor (**AK**) on the previous draft, together with its resolution (**GZ**). Because the revision restructured the thesis around a single canonical setup (Section 1) and, in doing so, *removed* several sections that the earlier draft contained, resolutions that entailed deletion are flagged explicitly with **[SECTION REMOVED]** so that the change is transparent.

Chapter 1 — Introduction

1. **AK:** Abstract is missing.

GZ: Resolved. English and Hebrew abstracts are included via `ch00_frontmatter.tex` and `ch09_hebrew_abstract.tex`, and both were rewritten to match the final (canonical) scope.

2. **AK:** There are compilation errors. Would be good to resolve them.

GZ: Resolved. The missing-brace errors and a duplicate `\listoflistequations` were fixed; the project compiles cleanly under XeLaTeX with zero undefined references.

3. **AK:** There are compilation errors. Would be good to resolve them. (second instance)

GZ: Resolved together with the item above; the second flagged location shared the same root cause.

4. **AK:** Unusual citation.

GZ: Resolved. The inline `source: [...]` annotations placed under equations were removed throughout, and equation attributions now use standard `\cite` references (e.g. Proakis, Goodfellow et al.).

5. **AK:** This is not causal... [regarding the symmetric relay window $x_R[i] = f_\theta(y_R[i - w : i + w])$]

GZ: Addressed and formalised as Remark 1.1. The symmetric window is non-causal in the streaming sense but fully realizable in the half-duplex, store-and-forward setting: the relay buffers the entire listening-phase block before transmitting, so by the time $x_R[i]$ is computed all samples $y_R[i - w : i + w]$ have been received. The window costs a fixed processing latency of w symbols; AF/DF use $w = 0$ (strictly causal). This is analysed quantitatively in the latency discussion of Section 8.4.1.

6. **AK:** This is the transmitted signal by the relay.

GZ: Resolved; the surrounding sentence was corrected and the time index restored.

7. **AK:** The time axis is missing. Aggregation of information can be useful, especially when incorporating an error correction code. Do you allow only scalar operations at the relay?

GZ: Resolved. The time index is restored, and the relay is explicitly *not* restricted to scalar processing: learned relays use a length- $2w + 1$ sliding window ($w > 0$), while AF and symbol-wise DF are memoryless ($w = 0$). The uncoded setting is stated up front; the value of temporal aggregation is in fact a central theme of the principal contribution (Chapter 7), where the window is what allows the learned relay to equalise unknown intersymbol interference.

8. **AK:** I agree that AF is memoryless but DF is not. Information-theoretic DF works in blocks to incorporate a strong (capacity-achieving over a single link) code and recovers the message at the end of the time block. You may of course use simple slicing, but this won't do much good.

GZ: Addressed and formalised as Remark 1.2. The thesis studies *symbol-wise* (slicing) DF and now states this explicitly wherever DF appears, distinguishing it from information-

theoretic block-DF. The distinction is load-bearing in the results: the high-SNR optimality of symbol-wise DF (H2) is explicitly *not* claimed to carry over to the coded block-DF setting, and a coded block-DF comparison is identified as the leading item of future work (Chapter 8).

Chapter 2 — Literature Review

1. **AK:** You haven't defined the system model so this discussion seems out of place.
GZ: Resolved. The system-model definition was moved to Chapter 1; Chapter 2 is now a pure literature review, and the MIMO background section carries a one-sentence pointer noting it supports only the deferred material.
2. **AK:** The relevance to this equation is not clear. Also, what do you mean by scaling? Isn't that the capacity for any SNR?
[SECTION REMOVED] GZ: The capacity/scaling passage was part of the MIMO capacity discussion. In the canonical restructure the entire MIMO thread was removed from the main line of the thesis (see comment 4 below), so the unclear equation no longer appears. The residual capacity material retained as background was corrected to avoid the ambiguous "scaling" phrasing.
3. **AK:** Here you introduce fading for the first time. You could've done that already in the SISO case. Also, now the channels are complex whereas up until now they were real (and with fixed known gains).
GZ: Resolved. The model conventions were unified: the canonical channel is i.i.d. Rayleigh fast fading in complex baseband from the outset. For SISO Hop 1, BPSK is transmitted on the real axis and the receiver applies coherent compensation and takes $\Re\{\cdot\}$, which is the sufficient statistic; this reconciles the earlier real-valued derivations with the complex model and is stated once in Chapter 1.
4. **AK:** Why not recovering both then? You keep jumping between multiple settings: no-fading/slow-fading/fast-fading, real/complex, SISO/MIMO. Please define one model that you've tackled and stick to it, unless you've actually tackled several settings. Then, each of them deserves a separate chapter.
[SECTION REMOVED] GZ: This was the decisive comment behind the restructure. The thesis now fixes *one* canonical setup — SISO on both hops, i.i.d. Rayleigh fast fading, complex baseband, BPSK, uncoded BER — and varies only the relay function. The AWGN and Rician robustness studies, the 2×2 MIMO second hop with ZF/LMMSE/SIC equalization, the 16-PSK modulation, the CSI-injection study, and the end-to-end autoencoder were all *removed* from the thesis; each is now listed as future work. The two threads that survive as deliberate, separately-chaptered studies are the higher-order-modulation extension (Chapter 6) and the unknown-channel contribution (Chapter 7), each with its own clearly stated configuration — exactly as requested.
5. **AK:** The diversity–multiplexing tradeoff (DMT) is irrelevant in the ergodic setting. DMT assumes slow fading and discusses performance in the presence of some outage probability.

[SECTION REMOVED] GZ: The DMT discussion was removed. The canonical setting is ergodic fast fading with an uncoded-BER metric, for which DMT does not apply; only a single background sentence remains, with no DMT-based claims.

6. **AK:** Why is it needed at all?

[SECTION REMOVED] GZ: The flagged passage belonged to the MIMO-equalization background that supported the now-removed MIMO study. With the MIMO thread deferred to future work, the passage was deleted; the remaining equalization background is confined to a clearly-labelled subsection and kept only to the extent needed to define terms used in the future-work discussion.

7. **AK:** It's linear MMSE (LMMSE) unless you work with a Gaussian codebook, which I suspect you do not.

GZ: Resolved. All references to “MMSE” equalization were corrected to “LMMSE” (linear MMSE), since the system uses a BPSK constellation rather than a Gaussian codebook. Note that the MIMO equalization material itself is now future work (comment 4); the terminology was corrected in the retained background nonetheless.

8. **AK:** Again, you can construct a variant that recovers a block code, cancels it out from the measurements, then recovers the next block code, etc. [coded SIC]

GZ: Acknowledged and scoped. As with block-DF (Ch. 1, comment 8), the thesis studies uncoded, symbol-level processing; coded per-layer SIC is outside the uncoded scope and is named explicitly in the future-work discussion. Since the MIMO study itself is now deferred (comment 4), coded SIC is doubly future work.

9. **AK:** I thought the goal was to replace the various classical strategies with a neural net that would learn what's the best strategy.

GZ: This comment reshaped the thesis's central claim, and the response is now the thesis's main finding rather than a rebuttal. The honest result is that a learned relay does *not* universally replace classical strategies: on a matched, known channel symbol-wise DF is at least as good at zero parameter cost (H2). What the learned relay *does* provide — and this is the principal contribution (Chapter 7) — is family-agnostic mitigation when the channel is unknown or mismatched: it restores reliable relaying where memoryless classical relays fail, approaches the model-aware optimum (Viterbi) without knowing the channel family, and does so at constant, parallelizable complexity. The framing is therefore “characterise when learning helps,” not “replace classical processing everywhere” — a change made directly in response to this comment.

Summary of removed material. In direct response to comment 4 above, the following were removed from the thesis and redesignated as future work: the AWGN and Rician channel-robustness studies; the 2×2 MIMO second hop and its ZF/LMMSE/SIC equalization; 16-PSK modulation; the CSI-injection ablation; and the end-to-end autoencoder. The diversity–multiplexing tradeoff discussion (comment 5) and the MIMO-capacity/equalization background passages (comments 2 and 6) were likewise deleted. All removals served the single-canonical-model discipline requested in comment 4, and none affects the validity of

the retained results, which were re-run and re-derived entirely within the canonical setup.

אוניברסיטת תל-אביב

הפקולטה להנדסה ע"ש איבי ואלדר פליישמן

בית הספר להנדסת חשמל

ארכיטקטורות למידה עמוקה לתקשורת ממסר דו-קפיצתית: מחקר השוואתי של אסטרטגיות ממסר קלאסיות ומבוססות רשתות נוירונים

חיבור זה הוגש כעבודת גמר לקראת התואר "מוסמך אוניברסיטה"

בהנדסת חשמל ואלקטרוניקה באוניברסיטת תל-אביב

על-ידי

גיל צוקרמן

ניסן תשפ"ו

אוניברסיטת תל-אביב

הפקולטה להנדסה ע"ש איבי ואלדר פליישמן

בית הספר להנדסת חשמל

ארכיטקטורות למידה עמוקה לתקשורת ממסר דו-קפיצתית: מחקר השוואתי של אסטרטגיות ממסר קלאסיות ומבוססות רשתות נוירונים

חיבור זה הוגש כעבודת גמר לקראת התואר "מוסמך אוניברסיטה"

בהנדסת חשמל ואלקטרוניקה באוניברסיטת תל-אביב

על-ידי

גיל צוקרמן

העבודה נעשתה בבית ספר להנדסת חשמל

בהדרכת ד"ר אנטולי חזן ופרופ' ראג'א ג'יריס

ניסן תשפ"ו

תקציר

חיבור זה מציג מחקר השוואתי של אסטרטגיות ממסר (relay) קלאסיות ומבוססות רשתות נוירונים עבור תקשורת שיתופית דו-קפיצתית. המחקר כולו מתבצע על תצורה קנונית אחת, הקבועה מראש ואינה משתנה לאורך קו הטיעון המרכזי: קישור SISO בשני הקפיצות, מעל דעיכת Rayleigh מהירה בלתי-תלויה ושווה-התפלגות, בבסיס-רוחב מרוכב, אפנון BPSK ושיעור שגיאות סיביות (BER) בלתי-מקודד, כאשר פונקציית הממסר היא הציר היחיד המשתנה. נבדקות תשע שיטות ממסר: שתי גישות קלאסיות: הגברה-והעברה (AF) ופענוח-והעברה (DF) : ושבע שיטות מבוססות למידה עמוקה: למידה מפוקחת (MLP) ממסר Hybrid המתאים את עצמו לרמת ה-SNR, מודלים גנרטיביים (VAE), (CGAN וארכיטקטורות רצפים, Transformer), Mamba-2 S6, Mamba (SSD). כל התוצאות מבוססות על סימולציית מונטה קרלו (100,000 ביטים לכל נקודת SNR, עשרה ניסויים, רווחי סמך של 95%).

בתצורה הקנונית, ממסרי הרשת הטובים ביותר עולים על AF ב-SNR נמוך (0-4 dB) אך רק משתווים ל-DF הסימבולי, אשר שולט בתחום SNR בינוני-גבוה עם אפס פרמטרים. מחקר מורכבות חושף יחס הפוך-U בין גודל המודל לביצועים: רשת מינימלית בת 169 פרמטרים משתווה למודלים גדולים פי 100, בעוד גדלים בינוניים מציגים התאמת-יתר; בתקציב שווה של כ-3,000 פרמטרים הפער בין הארכיטקטורות כמעט נעלם, ממצא המעיד שהקיבולת, ולא הארכיטקטורה, היא הגורם המכריע במשימה זו.

תרומה משנית מצומצמת בוחנת אפנונים מסדר גבוה: הממצאים הקנוניים עוברים במלואם ל-QPSK מכוח אי-תלות רכיבי I/Q, בעוד 16-QAM חושף רצפת BER מבנית של עיבוד מפוצל לפי ציר; ניסוח הממסר כמסווג דו-ממדי משותף על פני 16 נקודות הקונסטלציה מבטל רצפה זו ומשתווה ל-DF ב-SNR גבוה.

התרומה המרכזית היא מחקר שיטתי של ממסר מבוסס-למידה תחת ערוצים בלתי-ידועים ובלתי-מותאמים. בערוץ ISI בעל שלושה מקדמים, הבלתי-ידוע לשרשרת הקליטה, הממסרים הקלאסיים חסרי-הזיכרון ננעלים על רצפת BER אנליטית של 0.25 : כאשר ביצועי DF אף מחמירים עם עליית ה-SNR : בעוד אותה רשת MLP בת 169 פרמטרים משקמת ממסר אמין עד כדי מרחק של 1-1.5 dB מגלאי MLSE Viterbi בעל ידיעת ערוץ מלאה, וללא כל ידע על משפחת הערוץ. התוצאה נתחמת מכל כיוון: היא לעולם אינה עולה על Viterbi מותאם; על שרשור מורכב של ISI, אי-ליניאריות ופאזה בלתי-ידועה היא נעה כ-2 dB מאחורי MLSE מבוסס-פיילוטס ללא כל מודל של הפגם; במשטר נטול-מידע-מוקדם היא משתווה למאזן עיוור (CMA) תוך הימנעות מאי-יציבות של MLSE מבוסס-החלטות; וסריקת תקציב פיילוטס מאתרת סף זיהוי חד הקובע מתי הממסר הלמוד, שאינו נדרש לפיילוטס כלל, עוקף עיבוד קלאסי מבוסס-פיילוטס. לבסוף, ניתוח מורכבות מראה שעלות הממסר הלמוד לכל סימבול קבועה בסדר האפנון ובזיכרון הערוץ (בעוד Viterbi גדל כ- M^L) ומהירה פי 30-90 בזמן-ריצה. הממסר הלמוד אינו עולה על מקלט קלאסי מותאם, אך תופס נישא מוגדרת היטב של מיתון

בלתי-תלוי-משפחה, נטול-זיהוי, וקבוע-מורכבות.

אסטרטגיית הפריסה המומלצת היא ממסר Hybrid המשלב עיבוד למידה ב-SNR נמוך עם DF ב-SNR גבוה (169 פרמטרים, כ-0.7 KB, פחות מ-3 שניות אימון); עבור 16-QAM, ניסוח כמסווג דו-ממדי משותף; ובכל קישור הנושא פגם בלתי-ממודל (זיכרון, אי-ליניאריות), ממסר למוד מינימלי. התובנה המרכזית היא שממסרים למודים לעולם אינם עולים על האופטימום מודע-המודל, אך ייחודיים בעמידותם כאשר אין ידיעה איזה מודל חל: ושמורכבות המודל חייבת להיות מותאמת למבנה המשימה, לא רק לגודלה.